

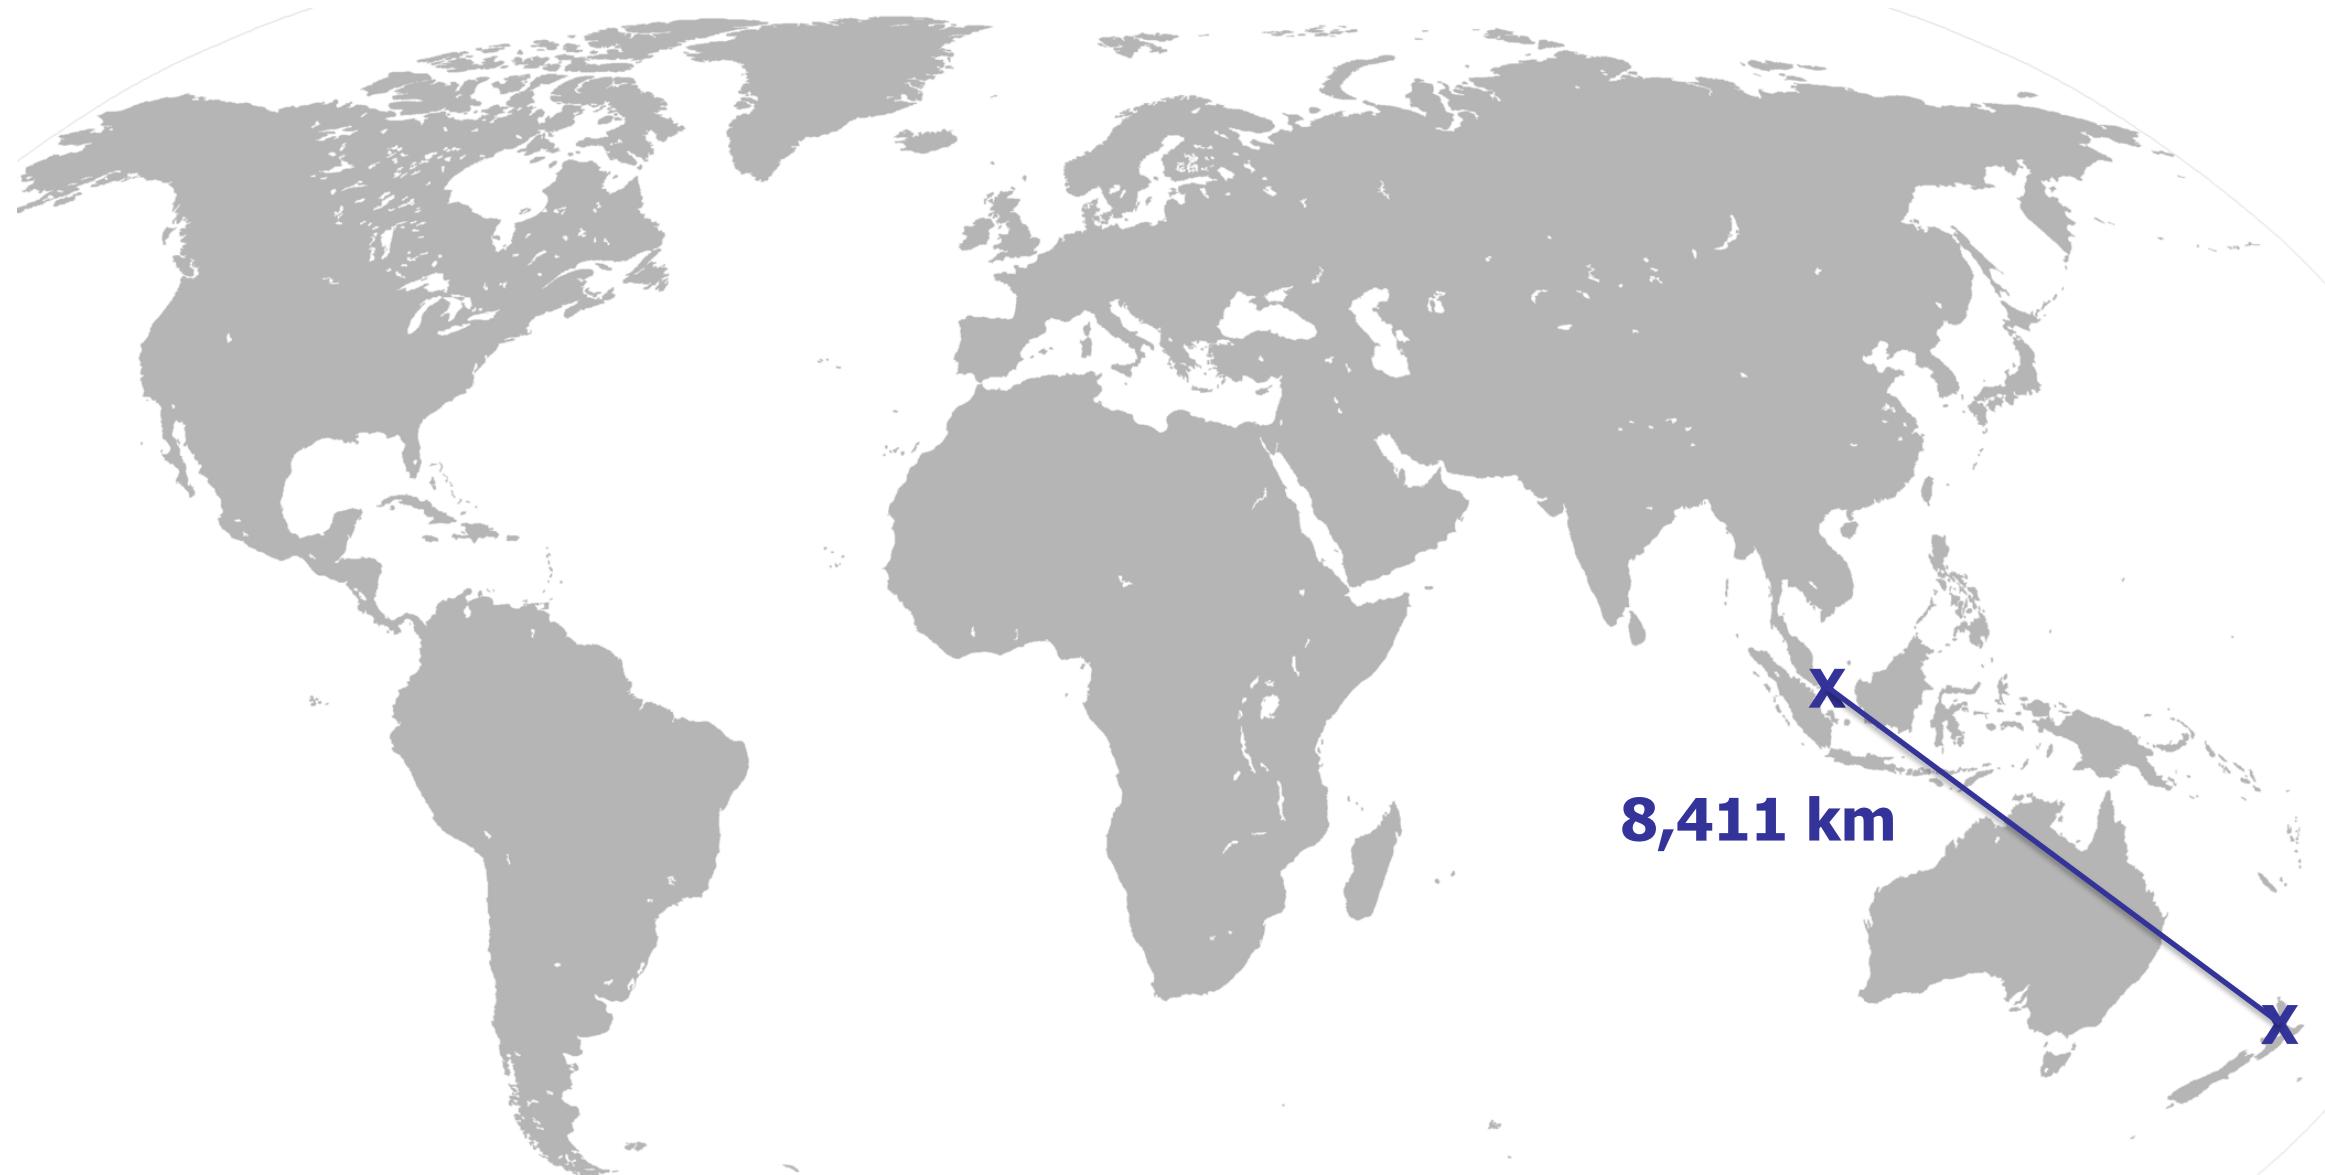


Code, Critique, Cure: Advancing LLM Reasoning for AI-Augmented Software Maintenance

David Lo, FACM, FIEEE

ICSME 2025, Auckland, New Zealand, September 2025

Self-Introduction



Self-Introduction



Self-Introduction



Self-Introduction



Singapore Management University



- Third university in Singapore
- Number of students:
 - 8000+ (UG)
 - 1800+ (PG)
- Schools:
 - Business
 - Computing
 - Economics
 - Accountancy
 - Law
 - Social Science



Center for Research on Intelligent Software Engineering (RISE)



Elsevier JSS'21, Bibliometric Study

Table 3

Most active institutions in software engineering

Rank	Name
1	University of California
2	Carnegie Mellon University
3	Nanjing University
4	Microsoft Research
5	Singapore Management University

CSRankings, SE, Aug 2025

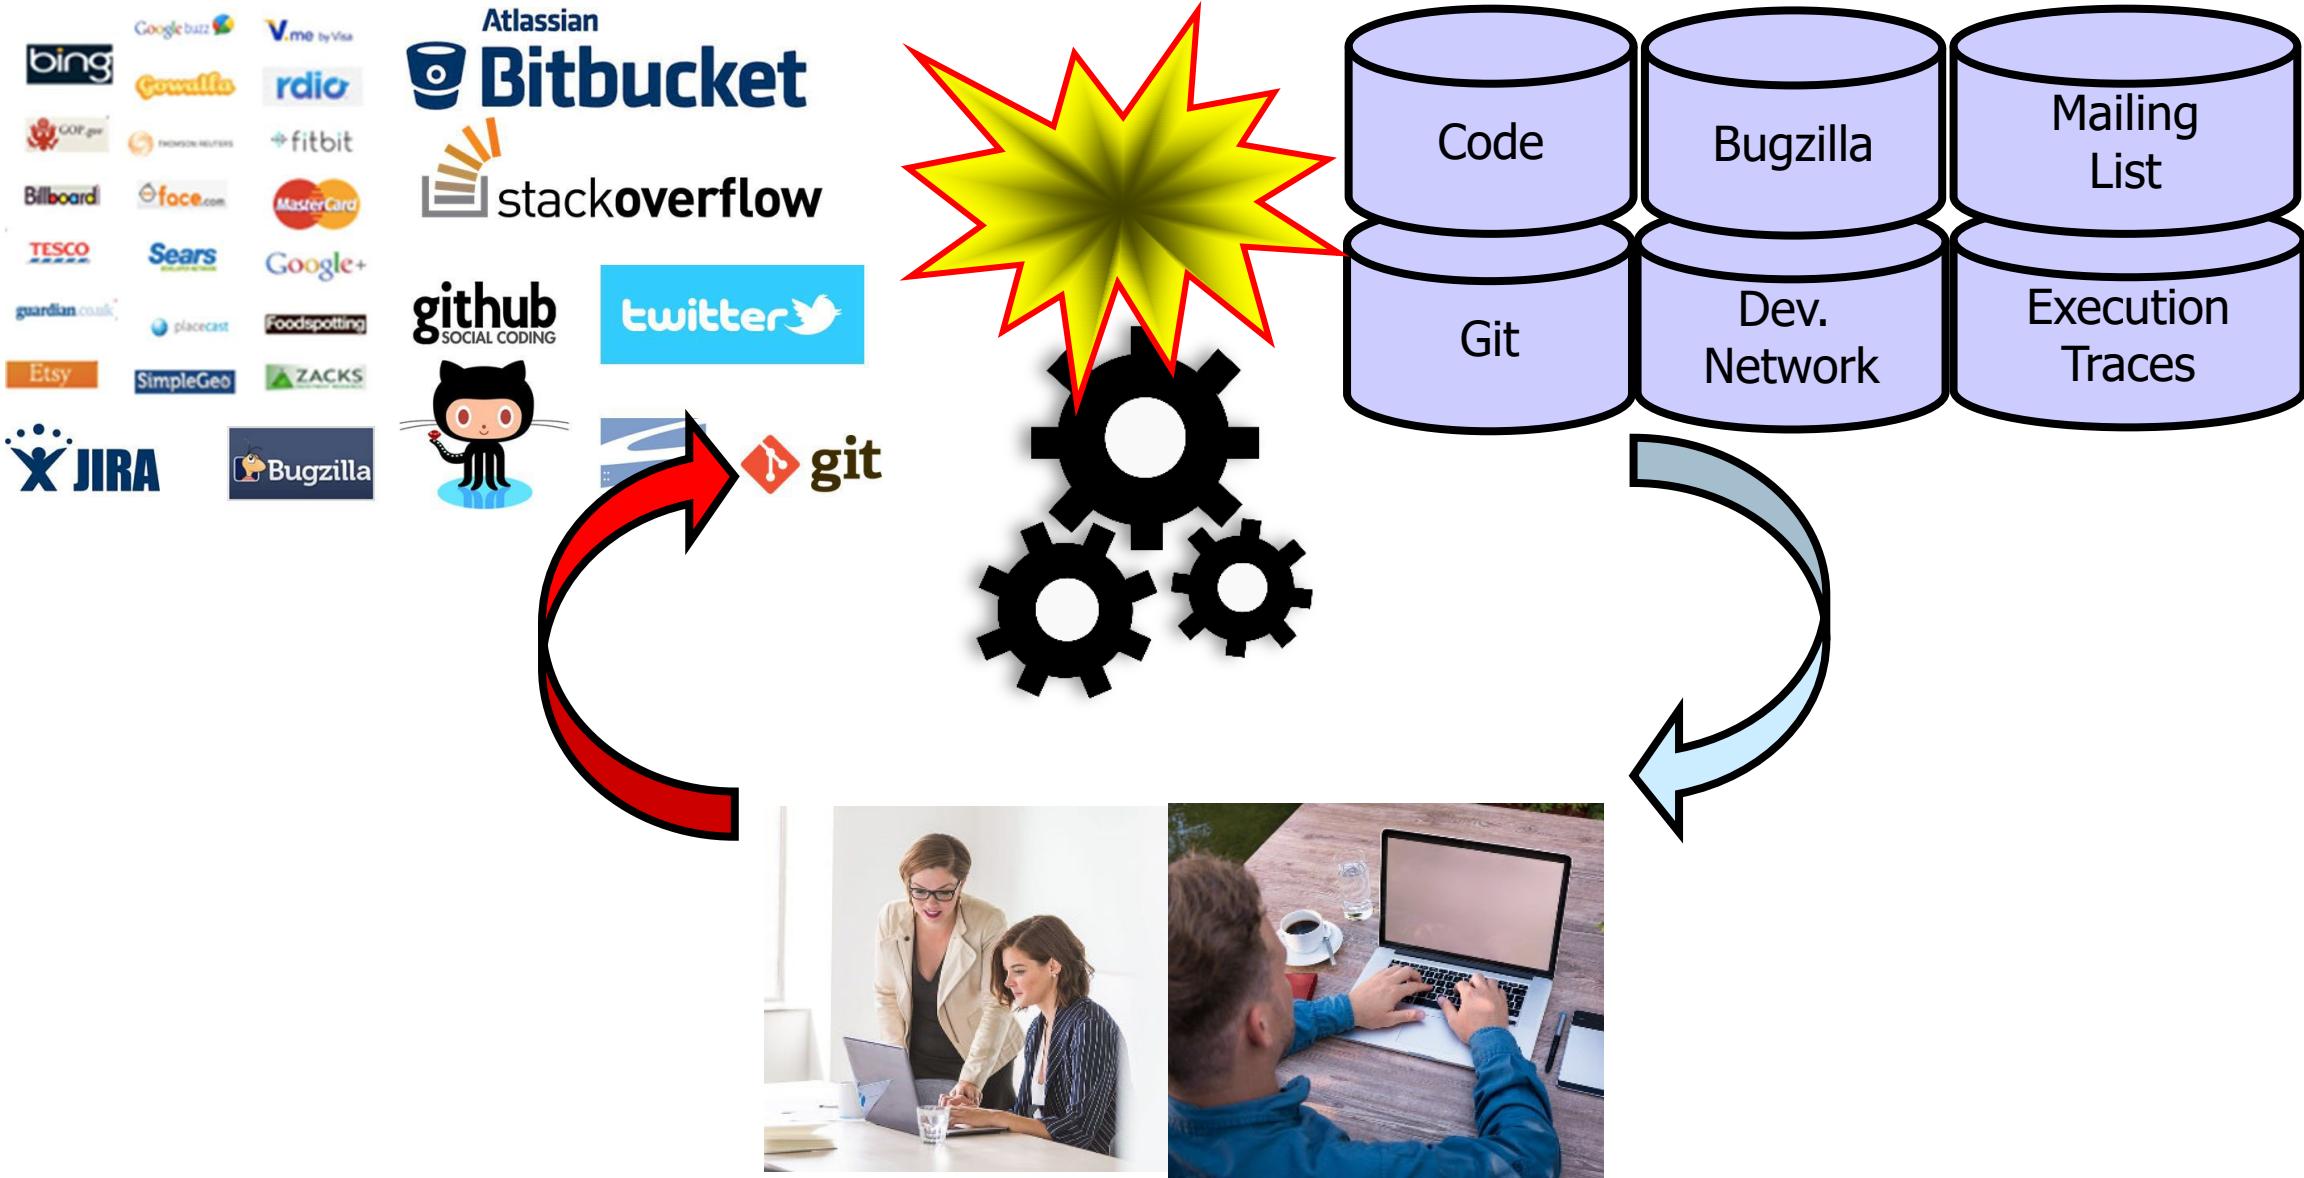
#	Institution	Count	Faculty
1	► Nanjing University 🇨🇳📊	45.6	40
2	► Peking University 🇨🇳📊	31.3	22
3	► Carnegie Mellon University 🇺🇸📊	31.2	16
4	► Singapore Management University 🇸🇬📊	25.0	9



School of
Computing and
Information Systems

Centre for Research on
Intelligent Software
Engineering

AI for Software Engineering



Experience with AI4SE

SMArTIC: Towards Building an Accurate, Robust and Scalable Specification Miner

FSE'06

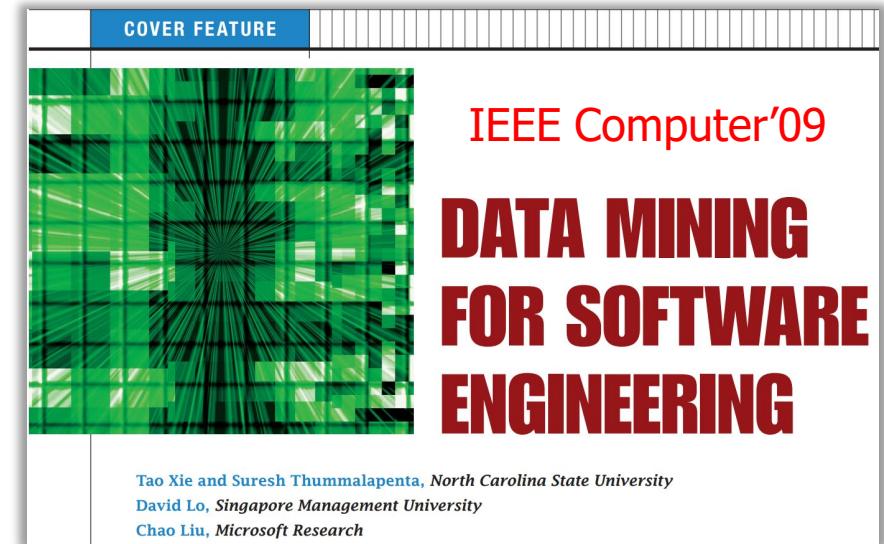
David Lo and Siau-Cheng Khoo
 Department of Computer Science, National University of Singapore
 {dlo,khoosc}@comp.nus.edu.sg

Efficient Mining of Iterative Patterns for Software Specification Discovery

KDD'07

David Lo and Siau-Cheng Khoo
 Department of Computer Science
 National University of Singapore
 {dlo,khoosc}@comp.nus.edu.sg

Chao Liu
 Department of Computer Science
 University of Illinois-UC
 chaoliu@cs.uiuc.edu



Experience with AI4SE

Classification of Software Behaviors for Failure Detection: A Discriminative Pattern Mining Approach KDD'09

David Lo
Singapore Management University
davidlo@smu.edu.sg

Jiawei Han[†]
University of Illinois at Urbana-Champaign
hanj@cs.uiuc.edu

Hong Cheng^{*}
Chinese University of Hong Kong
hcheng@se.cuhk.edu.hk

Siau-Cheng Khoo and Chengnian Sun
National University of Singapore
{khoosc,suncn}@comp.nus.edu.sg

Test oracle generation

Tag Recommendation in Software Information Sites MSR'13

Xin Xia*,‡, David Lo†, Xinyu Wang*, and Bo Zhou*§

*College of Computer Science and Technology, Zhejiang University
†School of Information Systems, Singapore Management University

Intelligent crowdsourced SE

A Discriminative Model Approach for Accurate Duplicate Bug Report Retrieval ICSE'10

Chengnian Sun¹, David Lo², Xiaoyin Wang³, Jing Jiang², Siau-Cheng Khoo¹

¹School of Computing, National University of Singapore

²School of Information Systems, Singapore Management University

³Key laboratory of High Confidence Software Technologies (Peking University), Ministry of Education
suncn@comp.nus.edu.sg, davidlo@smu.edu.sg, wangxy06@sei.pku.edu.cn,
jingjiang@smu.edu.sg, khoosc@comp.nus.edu.sg

Intelligent issue trackers

History Driven Program Repair SANER'16

Xuan-Bach D. Le, David Lo
School of Information Systems
Singapore Management University
{dxb.le.2013,davidlo}@smu.edu.sg

Claire Le Goues
School of Computer Science
Carnegie Mellon University
clegoues@cs.cmu.edu

Intelligent program repair

*"History-driven
program repair
influence
our work, the overall
pipeline is similar"*

- Facebook
Engineers

Our Research Agenda in AI4SE

Trustworthy and Synergistic Artificial Intelligence for Software Engineering: Vision and Roadmaps

David Lo

School of Computing and Information Systems,
Singapore Management University,
Singapore
Email: davidlo@smu.edu.sg



Trustworthy and Synergistic AI4SE:
Vision and the Road Ahead

David Lo

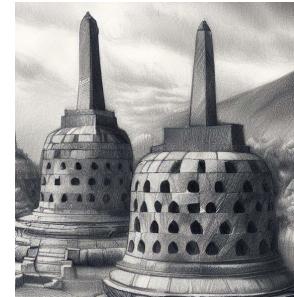
School of
Computing and
Information Systems



ICSE'23 Future of SE Talk

AI for Software Engineering

History



Trust



Roadmap I

Challenges



Synergy



Roadmap II

Vision



Call4Action

Towards Software Engineering 2.0

"If you want to go far, go together" – African Proverb



MONASH



UNIVERSITEIT
STELLENBOSCH
UNIVERSITY



Université
de Montréal



Carnegie
Mellon
University



NC STATE
UNIVERSITY



université
de BORDEAUX

TU/e
EINDHOVEN
UNIVERSITY OF
TECHNOLOGY



Microsoft
Research



Adobe

Alibaba

Tencent

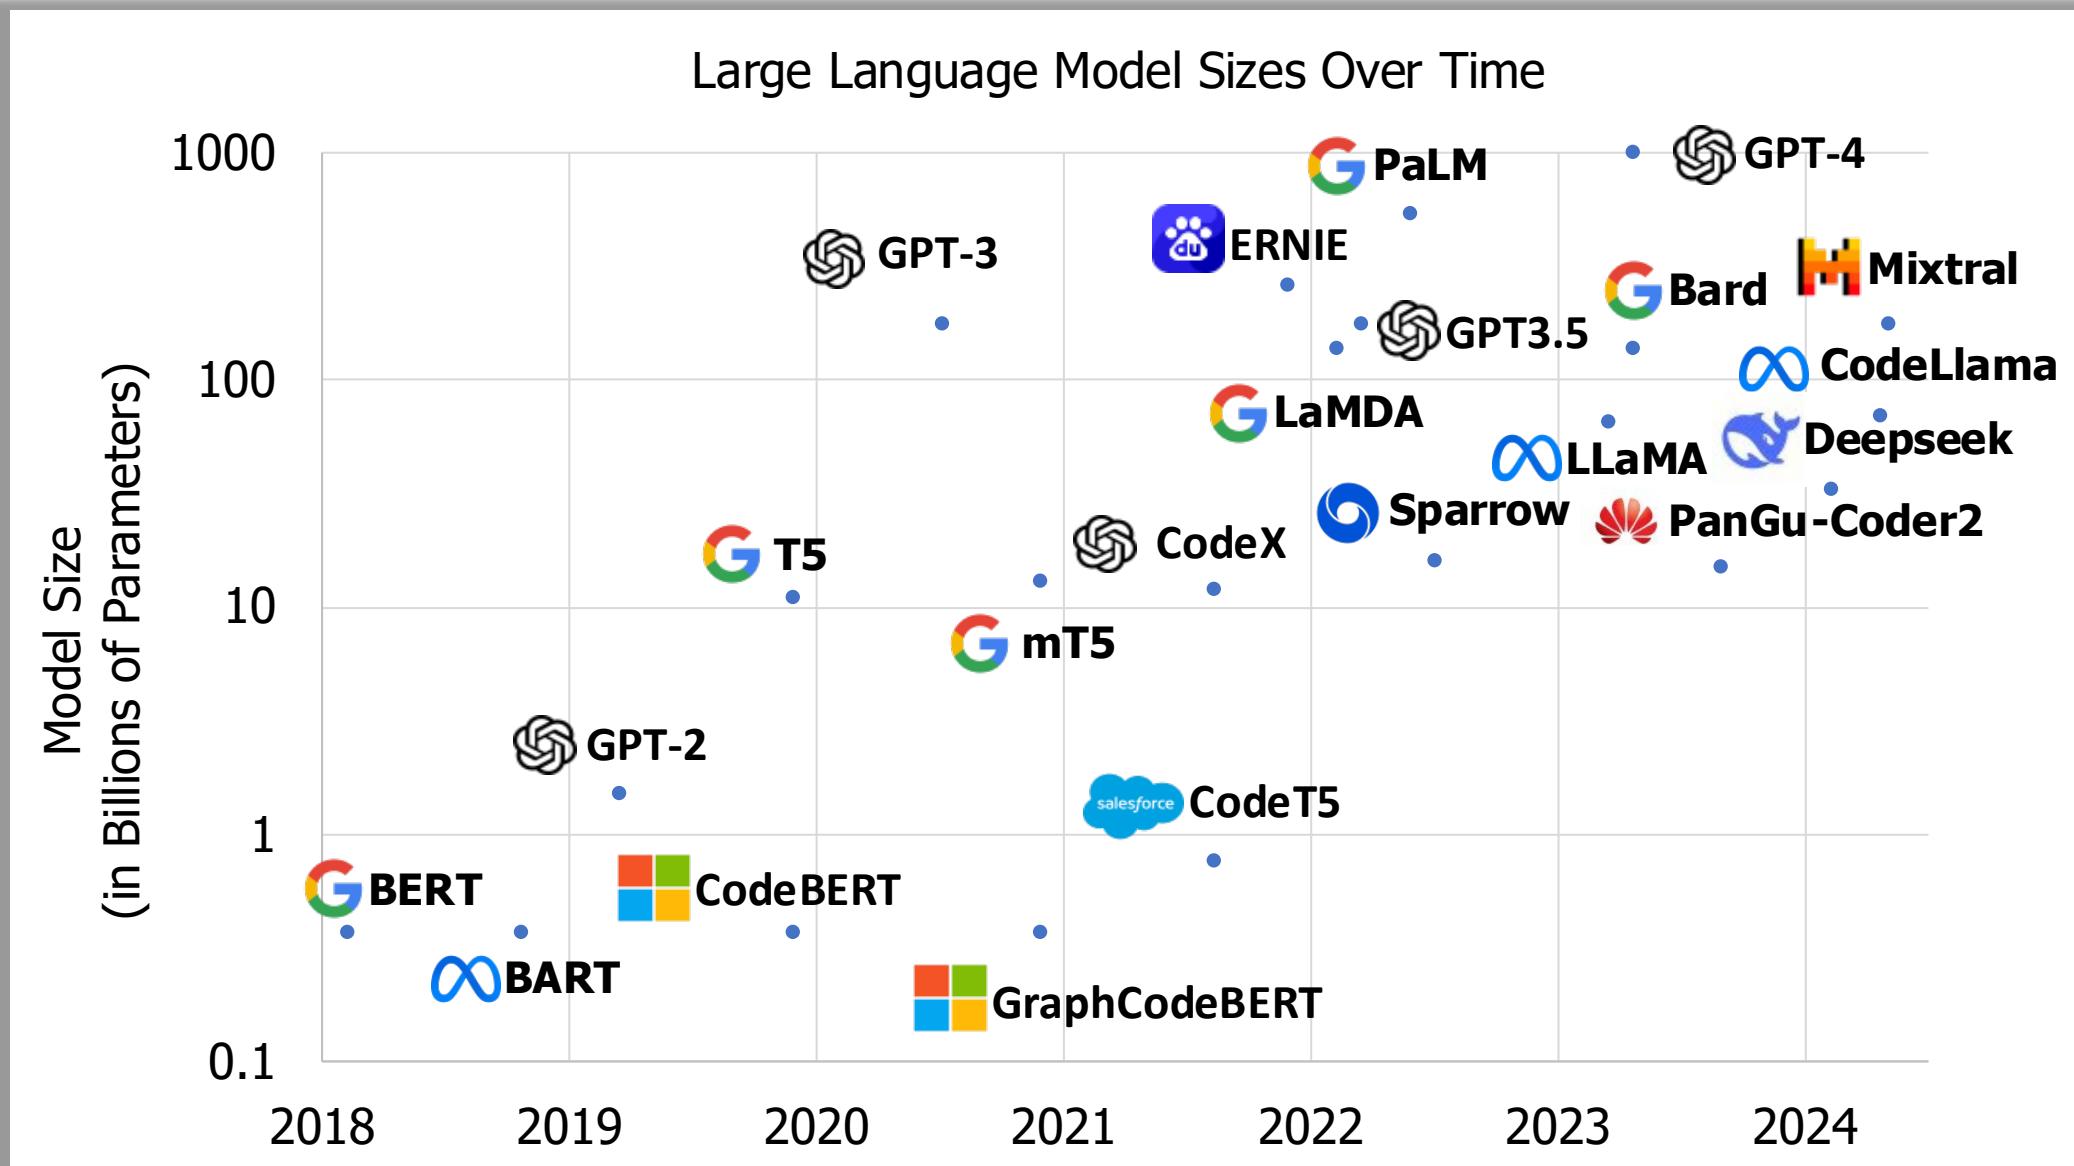


VERACODE
You change the world, we'll secure it.



Code, Critique, Cure: Advancing LLM Reasoning for AI-Augmented Software Maintenance

Large Language Models (LLMs)



LLM Can Greatly Help SE Tasks

ICSMIE 2020

Sentiment Analysis for Software Engineering: How Far Can Pre-trained Transformer Models Go?

Ting Zhang, Bowen Xu*, Ferdian Thung, Stefanus Agus Haryono, David Lo, Lingxiao Jiang

School of Information Systems, Singapore Management University

Email: {tingzhang.2019, bowenxu.2017}@phdcs.smu.edu.sg, {ferdianthung, stefanusah, davidlo, lxjiang}@smu.edu.sg



ICSMIE 2021

Assessing Generalizability of CodeBERT

Xin Zhou, DongGyun Han, and David Lo

School of Computing and Information Systems, Singapore Management University

xinzhou.2020@phdcs.smu.edu.sg, {dhan, davidlo}@smu.edu.sg



Early work on benchmarking code LLM, among most cited papers of ICSME 2021

LLMs Seem to Win for Many SE Scenarios

Large Language Models for Software Engineering: A Systematic Literature Review

XINYI HOU*, Huazhong University of Science and Technology, China

YANJIE ZHAO*, Monash University, Australia

YUE LIU, Monash University, Australia

ZHOU YANG, Singapore Management University, Singapore

KAILONG WANG, Huazhong University of Science and Technology, China

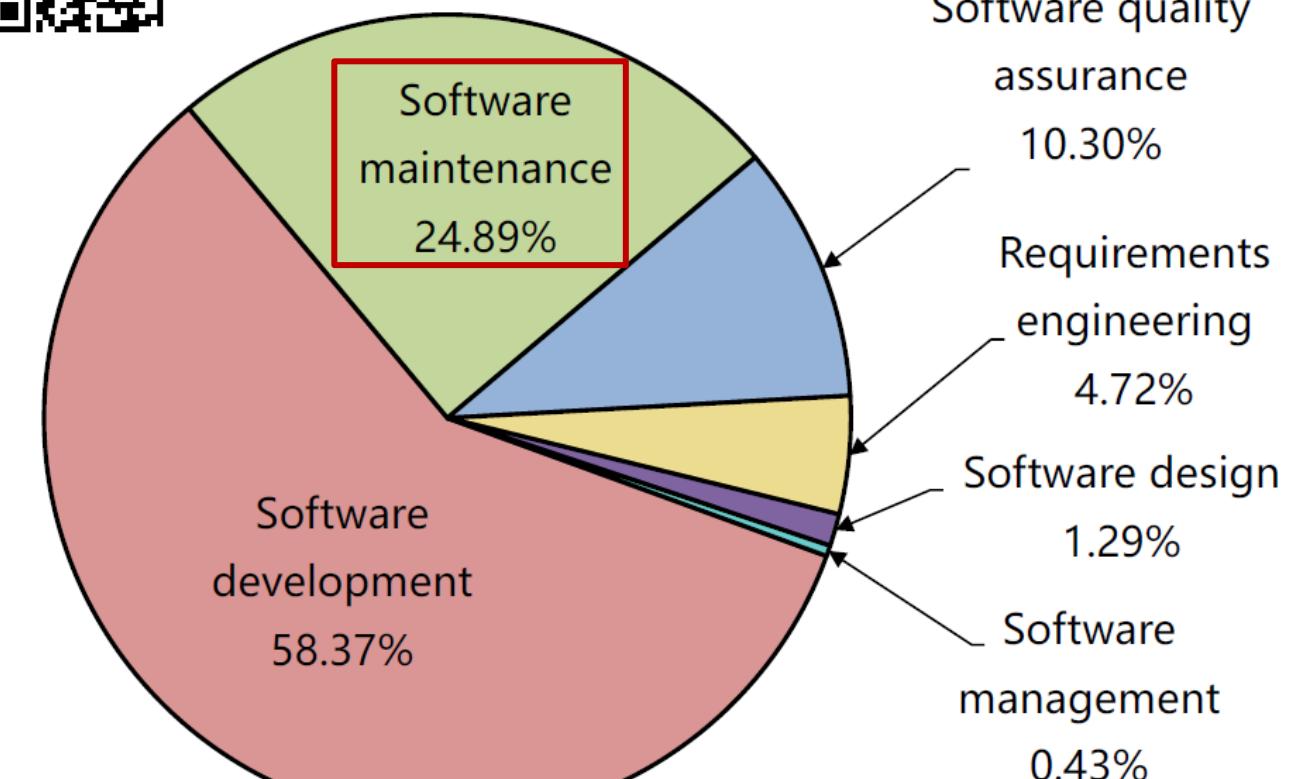
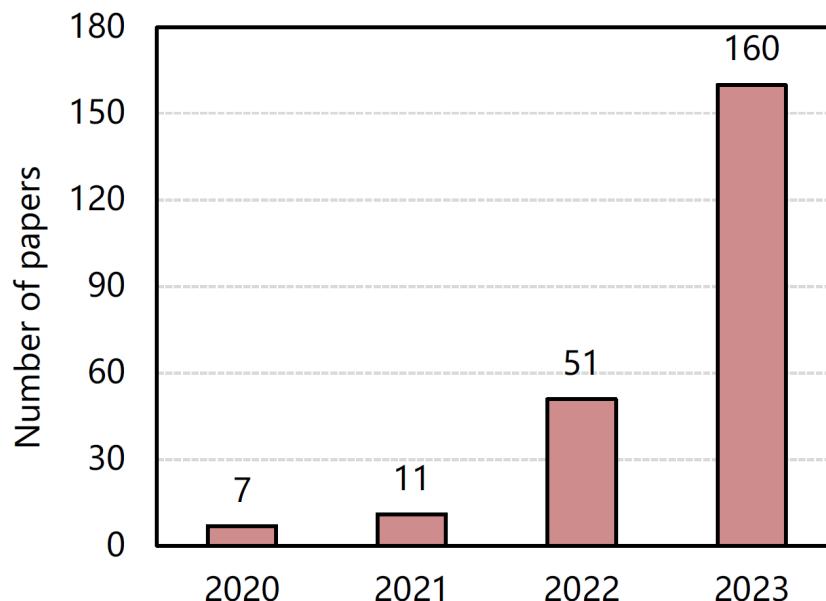
LI LI, Beihang University, China

XIAPU LUO, The Hong Kong Polytechnic University, China

DAVID LO, Singapore Management University, Singapore

JOHN GRUNDY, Monash University, Australia

HAOYU WANG[†], Huazhong University of Science and Technology, China



Despite Successes, Much Work Remains

Measuring the Impact of Early-2025 AI on Experienced Open-Source Developer Productivity

Joel Becker*, Nate Rush*, Beth Barnes, David Rein

Model Evaluation & Threat Research (METR)



InfoWorld

AI coding tools can slow down seasoned developers by 19%

News

Jul 11, 2025 • 6 mins

Despite Successes, Much Work Remains

Factor	Relevant Observations
Over-optimism about AI usefulness (C.1.1)	<ul style="list-style-type: none"> • Developers forecast AI will decrease implementation time by 24% • Developers post hoc estimate AI decreased implementation time by 20%
High developer familiarity with repositories (C.1.2)	<ul style="list-style-type: none"> • Developers slowed down more on issues they are more familiar with • Developers report that their experience makes it difficult for AI to help them • Developers average 5 years experience and 1,500 commits on repositories
Large and complex repositories (C.1.3)	<ul style="list-style-type: none"> • Developers report AI performs worse in large and complex environments • Repositories average 10 years old with >1,100,000 lines of code
Low AI reliability (C.1.4)	<ul style="list-style-type: none"> • Developers accept <44% of AI generations • Majority report making major changes to clean up AI code • 9% of time spent reviewing/cleaning AI outputs
Implicit repository context (C.1.5)	<ul style="list-style-type: none"> • Developers report AI doesn't utilize important tacit knowledge or context

LLM has limited reasoning capabilities for software maintenance tasks

How Can We Improve It?

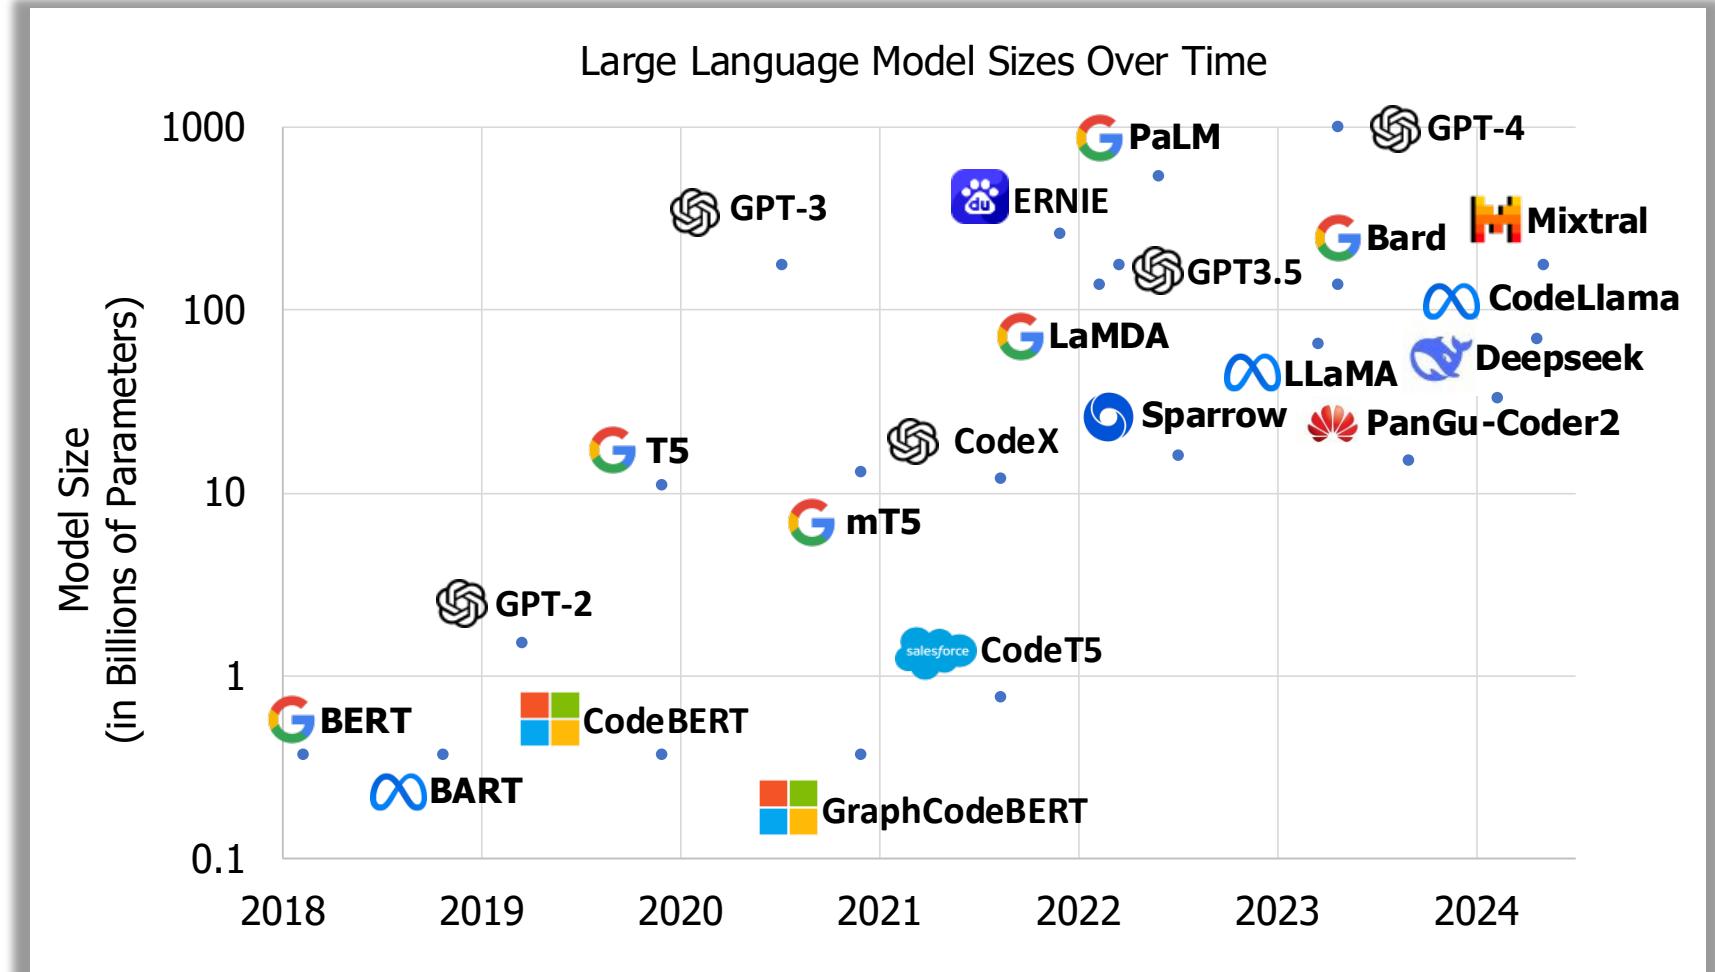
Too Much Focus on Models



“99% of the papers were model-centric with only 1% being data-centric”

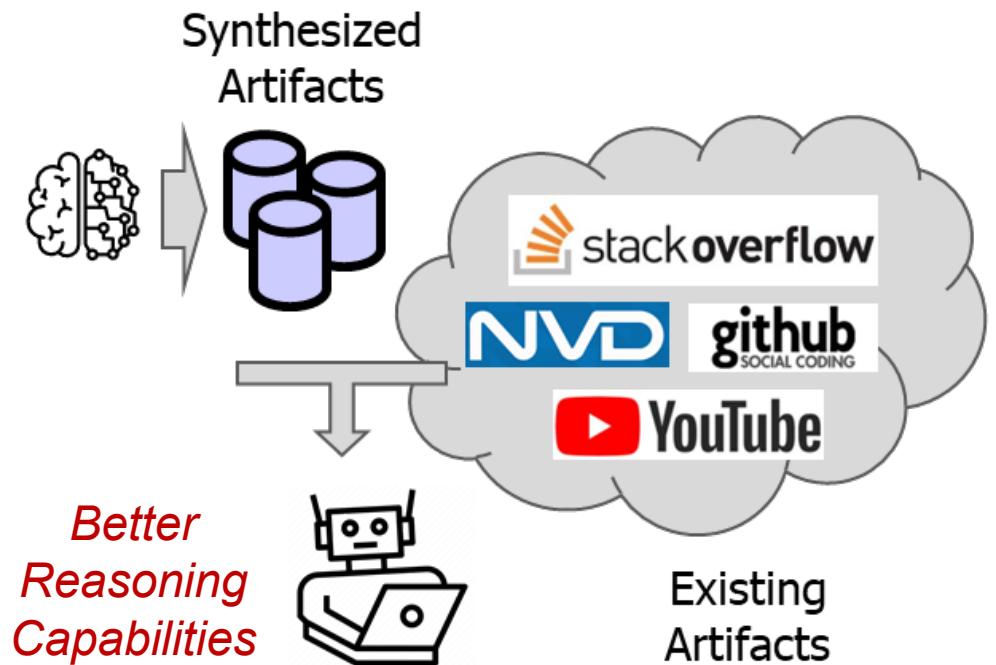
– **Andrew Ng**

Google Brain Co-Founder



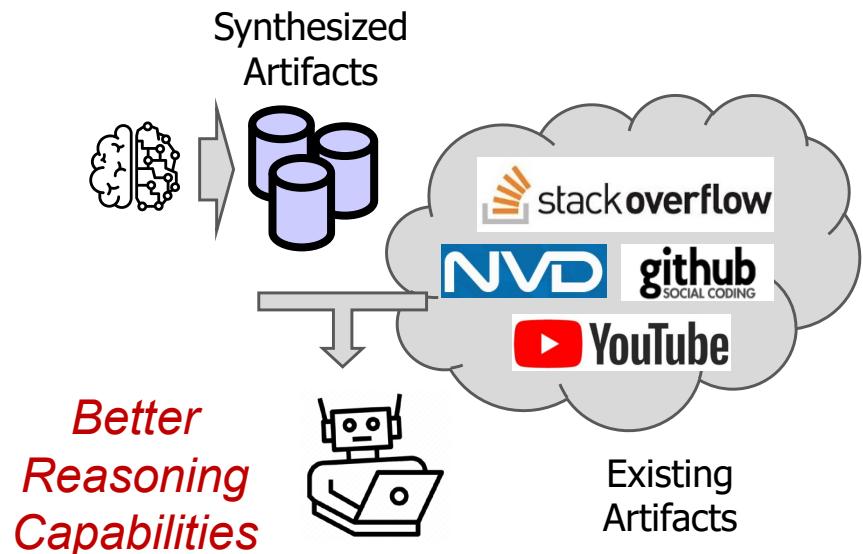
Data-Centric LLM4SE

- Use SE domain knowledge for software **artifact engineering**
- to *guide LLM reasoning*
- for *specific SE tasks*



Data-Centric LLM4SE: Artifact Engineering

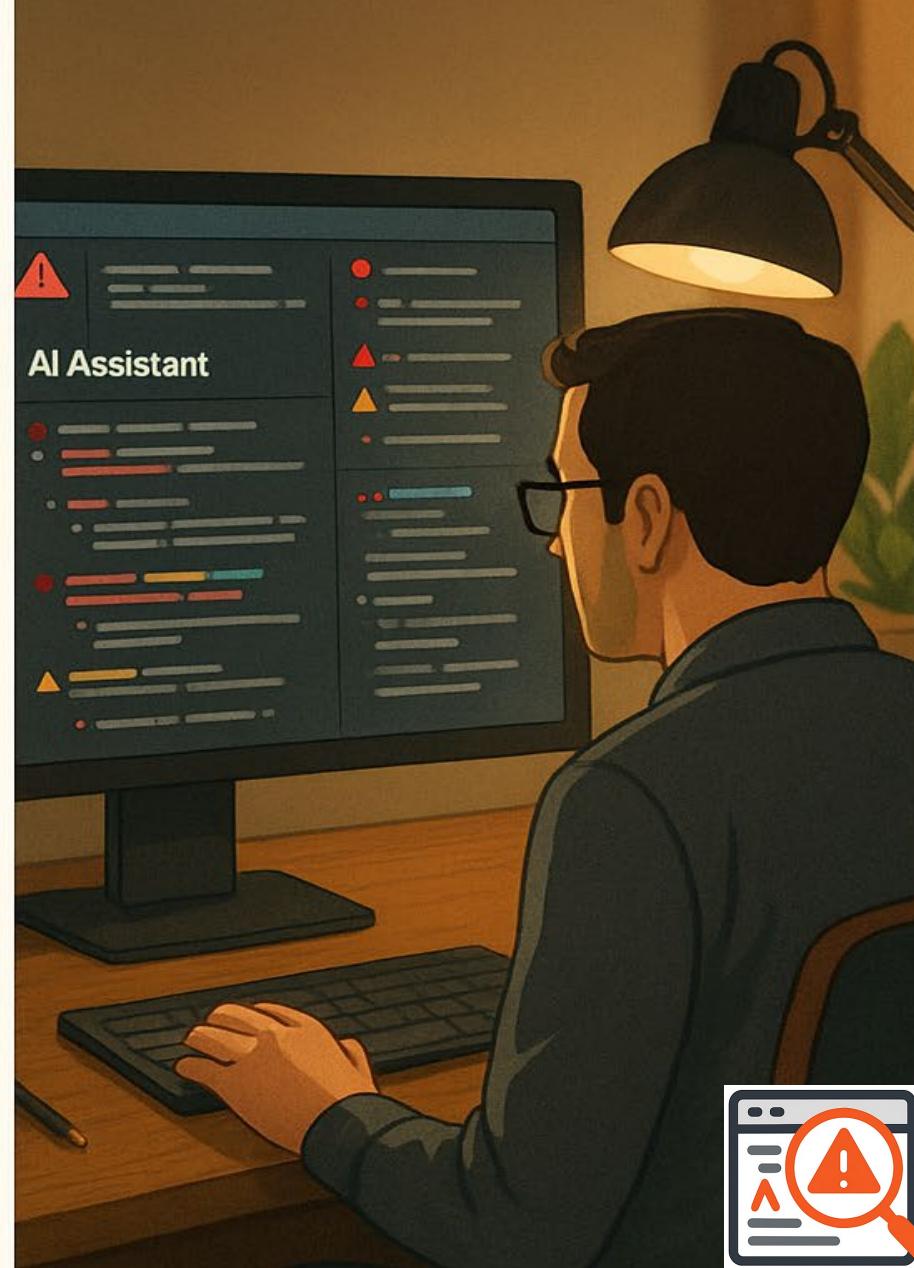
*“How Can We
Systematically Engineer
Software Artifacts to Build
Better LLM4SE Bots?”*



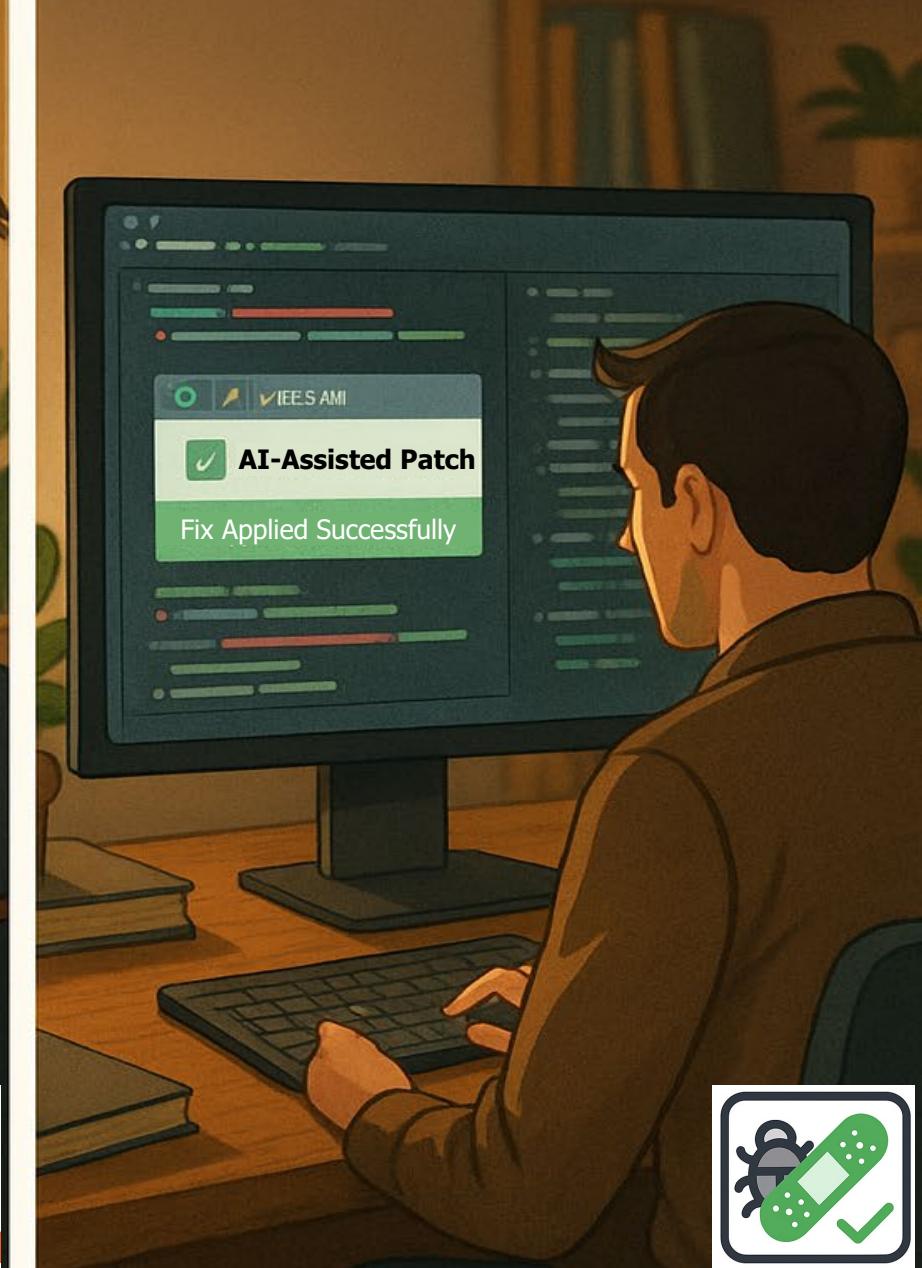
- Construct **LLM-reasoning friendly datasets** by
 - selecting the right artefacts
 - linking related pieces of information,
 - providing examples,
 - transforming them appropriately, etc.
- leveraging **domain knowledge**.
- **Scale up with LLM4(LLM4SE):** LLMs themselves can assist this *SE artifact engineering (crafting)* step



Code



Critique

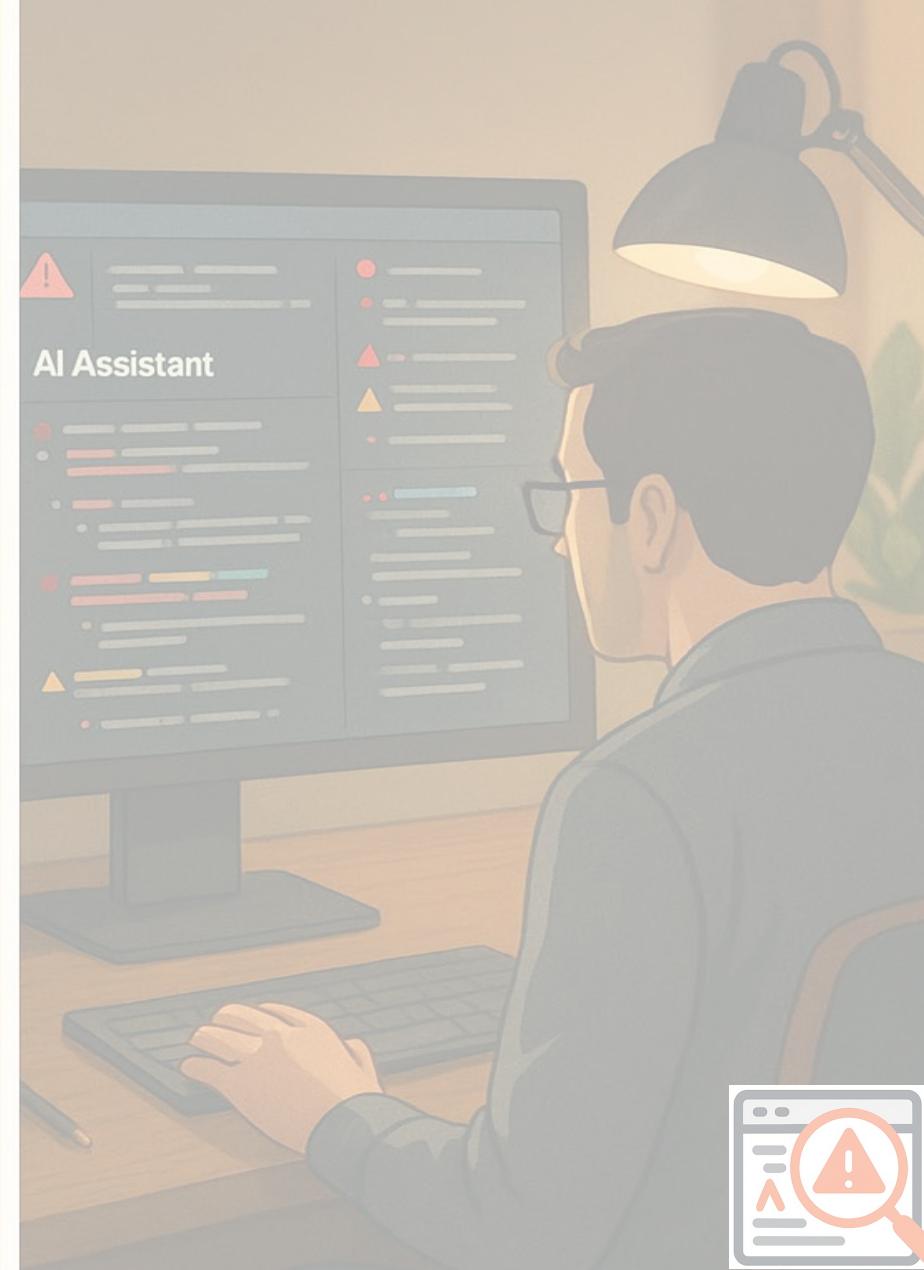


Cure

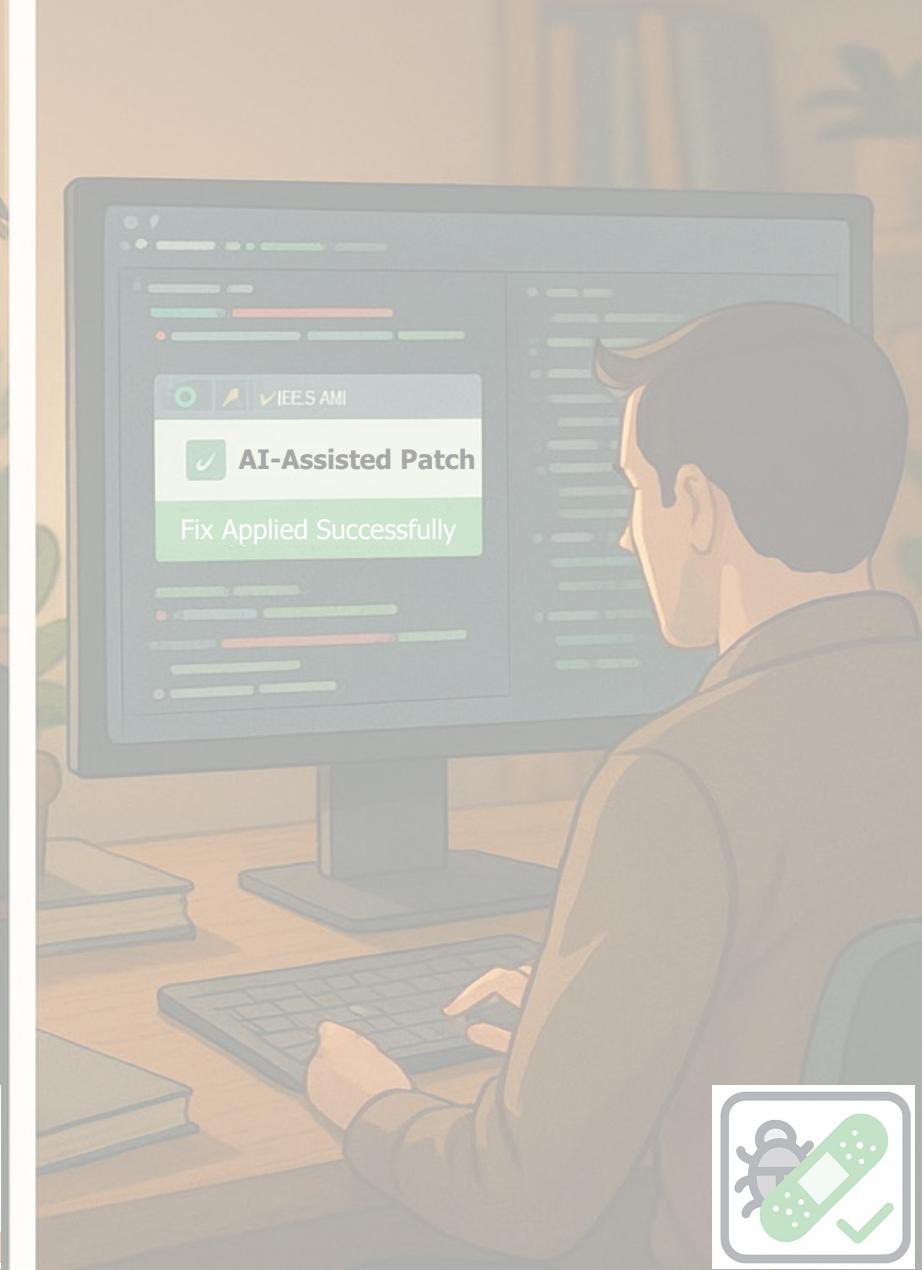




Code



Critique



Cure



Structured Data Can Be Helpful for AI4SE

ICPC 2018

Deep Code Comment Generation*

Xing Hu¹, Ge Li¹, Xin Xia², David Lo³, Zhi Jin¹

¹Key Laboratory of High Confidence Software Technologies (Peking University), MoE, Beijing, China

²Faculty of Information Technology, Monash University, Australia

³School of Information Systems, Singapore Management University, Singapore

¹{huxing0101,lige,zhijin}@pku.edu.cn, ²xin.xia@monash.edu, ³ davidlo@smu.edu.sg



First work that leverages structural information in code AST for better code comment generation; most cited paper of ICPC

Distilling Structured Reasoning Traces from Com. Knowledge

ICSE 2026

Think Like Human Developers: Harnessing Community Knowledge for Structured Code Reasoning

Chengran Yang
Singapore Management University
Singapore
cryang@smu.edu.sg

Zhensu Sun
Singapore Management University
Singapore
zssun@smu.edu.sg

Hong Jin Kang
University of Sydney
Australia
hongjin.kang@sydney.edu.au

Jieke Shi
Singapore Management University
Singapore
jiekeshi@smu.edu.sg

David Lo
Singapore Management University
Singapore
davidlo@smu.edu.sg



*Engineers high-quality **structured reasoning traces** from community knowledge
for better LLM-powered code generation*

Community Knowledge as Code Reasoning Training Resources

558: Time 99.4%, Solution with step by step explanation



Marlen

571 Mar 15, 2023

Divide and Conquer

Tree

Python

Python3

Intuition

Approach

1. Check if either quadTree1 or quadTree2 is a leaf node. If quadTree1 is a leaf node, return quadTree1 if its value is True, else return quadTree2. If quadTree2 is a leaf node, return quadTree2 if its value is True, else return quadTree1.
2. Recursively check the intersection of each child node of quadTree1 and quadTree2 by calling the intersect function on each of them. Store the return values in variables tl (top-left), tr (top-right), bl (bottom-left), and br (bottom-right).

Community discussions offer peer-reviewed insights into developers' problem-solving strategies

“How Can We Leverage Community Knowledge to Help LLM Reason Better?”

RT-Distiller: From Com. Posts to Structured Reasoning Traces

- Distilling high-quality reasoning traces (RTs) for code generation from evolving community knowledge.

✖ Mixed Quality Posts

LeetCode Problem:

Given two binary grids as quad trees, return their logical OR result as a quad tree while maintaining minimal node representation.

User Discussion:

Intuition

Approach

1. Check if either quadTree1 or quadTree2 is a leaf node...
2. Recursively check the intersection of each child node of quadTree1 and quadTree2 by calling the intersect function...
3. Check if all four child nodes (tl, tr, bl, br) are leaves and have the same value...
4. If any of the child nodes have a different value or are not leaf nodes, create a new internal node...

Complexity

Time complexity: ...

Code

```
class Solution:  
    def intersect ...
```

RT-Distiller: From Com. Posts to Structured Reasoning Traces

- Distilling high-quality reasoning traces (RTs) for code generation from evolving community knowledge.

✗ Mixed Quality Posts

✗ Artificial
Site-Specific
Code Structure

LeetCode Problem:

Given two binary grids as quad trees, return their logical OR result as a quad tree while maintaining minimal node representation.

User Discussion:

Intuition

Approach

1. Check if either quadTree1 or quadTree2 is a leaf node...
2. Recursively check the intersection of each child node of quadTree1 and quadTree2 by calling the intersect function...
3. Check if all four child nodes (tl, tr, bl, br) are leaves and have the same value...
4. If any of the child nodes have a different value or are not leaf nodes, create a new internal node...

Complexity

Time complexity: ...

Code

```
class Solution:
```

```
    def intersect ...
```

RT-Distiller: From Com. Posts to Structured Reasoning Traces

- Distilling high-quality reasoning traces (RTs) for code generation from evolving community knowledge.

✖ Mixed Quality Posts

✖ Artificial
Site-Specific
Code Structure

✖ Loosely Structured
Reasoning

LeetCode Problem:

Given two binary grids as quad trees, return their logical OR result as a quad tree while maintaining minimal node representation.

User Discussion:

Intuition

Approach

1. Check if either quadTree1 or quadTree2 is a leaf node...
2. Recursively check the intersection of each child node of quadTree1 and quadTree2 by calling the intersect function...
3. Check if all four child nodes (tl, tr, bl, br) are leaves and have the same value...
4. If any of the child nodes have a different value or are not leaf nodes, create a new internal node...

Complexity

Time complexity: ...

Code

```
class Solution:  
    def intersect ...
```

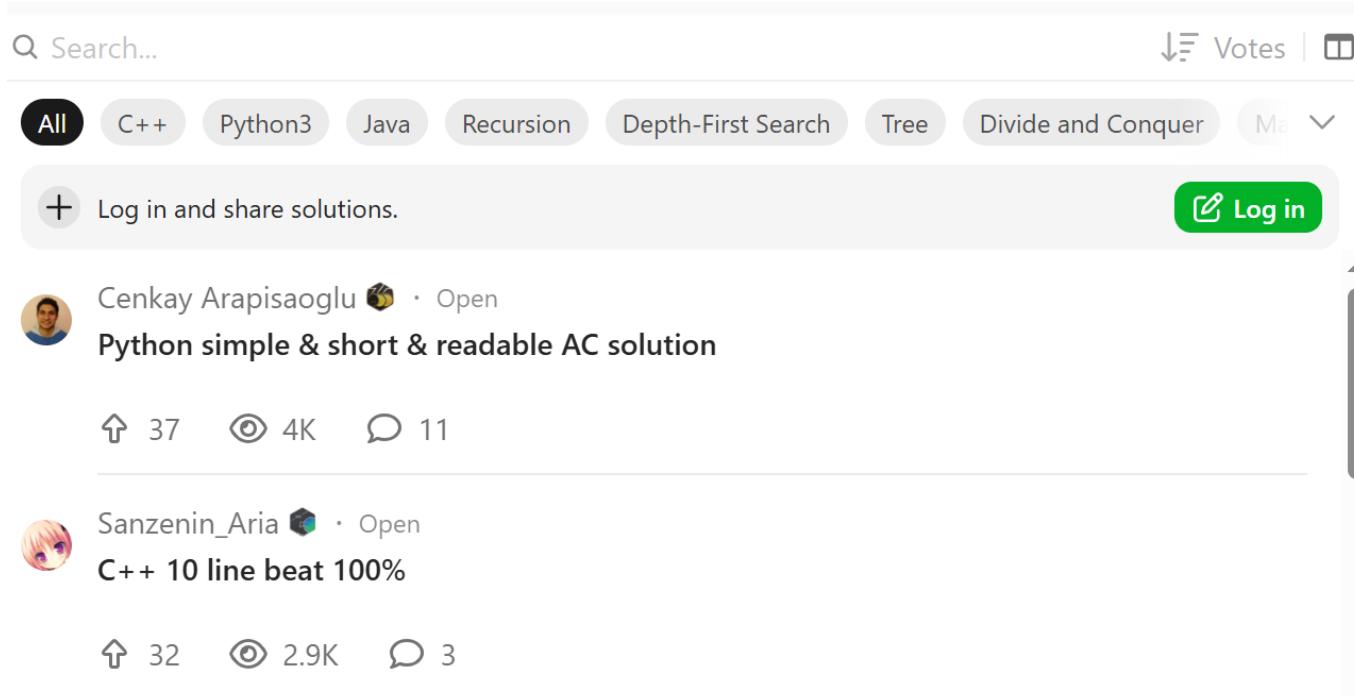
Step I: Addressing Mixed Quality Posts

Identify High Quality Posts

- Identify rating mechanisms
- Select highly-rated posts

Improve Post Quality

- Identify typical information pieces
- Detect missing pieces
- Synthesize missing pieces using LLM



Q Search...

Votes | □

All C++ Python3 Java Recursion Depth-First Search Tree Divide and Conquer Me ▾

+ Log in and share solutions.

 Cenkay Arapisaoglu · Open
Python simple & short & readable AC solution

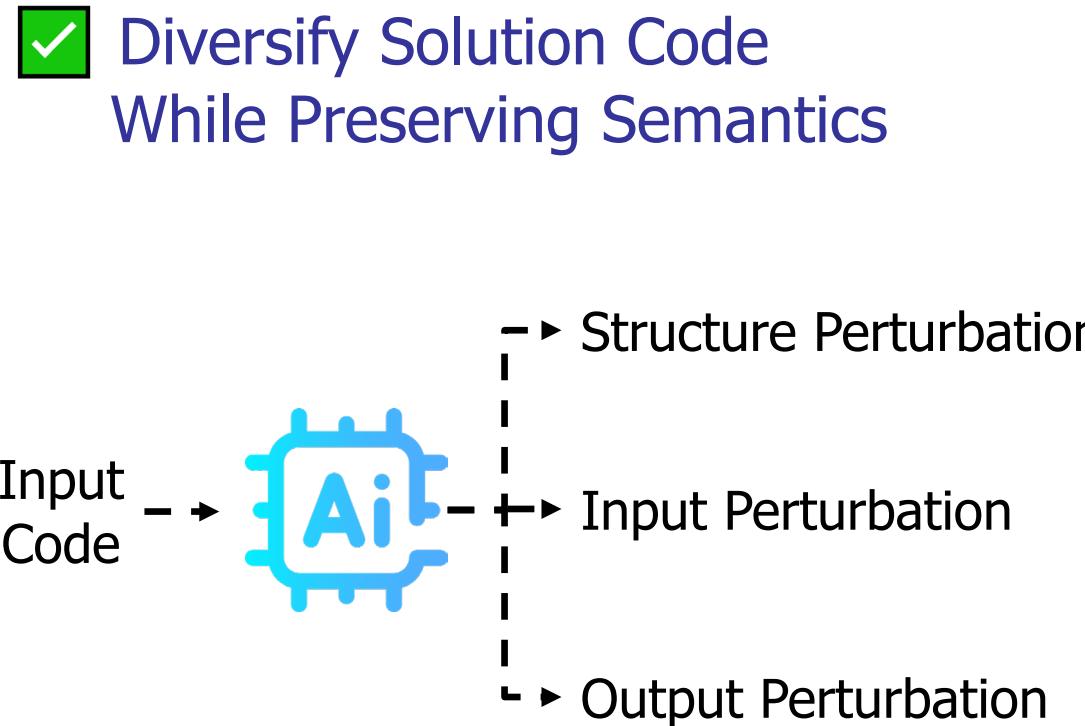
↑ 37 ⏵ 4K ⏷ 11

 Sanzenin_Aria · Open
C++ 10 line beat 100%

↑ 32 ⏵ 2.9K ⏷ 3



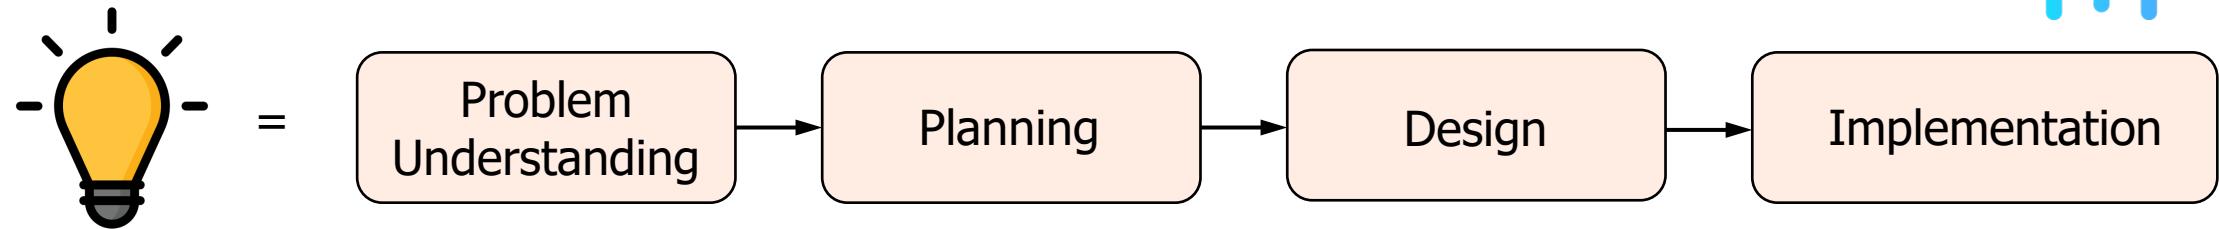
Step II: Dealing with Artificial Site-Specific Code Structure



Category	Variation Type
Input Format	Direct input (e.g., variables, constants)
	Structured input (e.g., list, tuple, dictionary)
	Batch input
	Interactive input()
	Boolean input
	Command-Line Arguments
Output Format	Structured input (e.g., list, tuple, dictionary)
	Boolean output
	Batch output
	Print statement
	Logging Output
Structural Variations	Class-based
	Function-based
	Script-based

Step III: Enriching Loosely Structured Reasoning Traces

- ✓ Rewrite a post into a structured trace inspired by SDLC



Step III: Enriching Loosely Structured Reasoning Traces

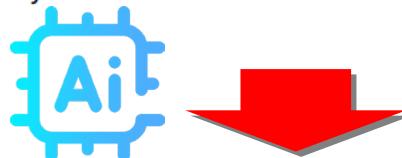
Given two binary grids as quad trees, return their logical OR result as a quad tree while maintaining minimal node representation.

User Discussion:

Intuition

Approach

1. Check if either quadTree1 or quadTree2 is a leaf node...
2. Recursively check the intersection of each child node of



First things first, I need to understand what a quad-tree is. From the description, it's a tree data structure where each internal node has exactly four children ...

Let's think about this step by step. First, I need to understand how the quad-tree represents the matrix. If a node is a leaf and its `val` is true, it means ...

Let me try to outline an approach:
 1. If `quadTree1` is a leaf:
 - If its `val` is true, then the entire grid is 1's, so the OR with anything else is still 1's ...

Let me try to write some solutions:
 function
 intersect(quadTree1, quadTree2):
 if quadTree1.isLeaf:
 if quadTree1.val:
 return quadTree1
 ...

Problem Understanding

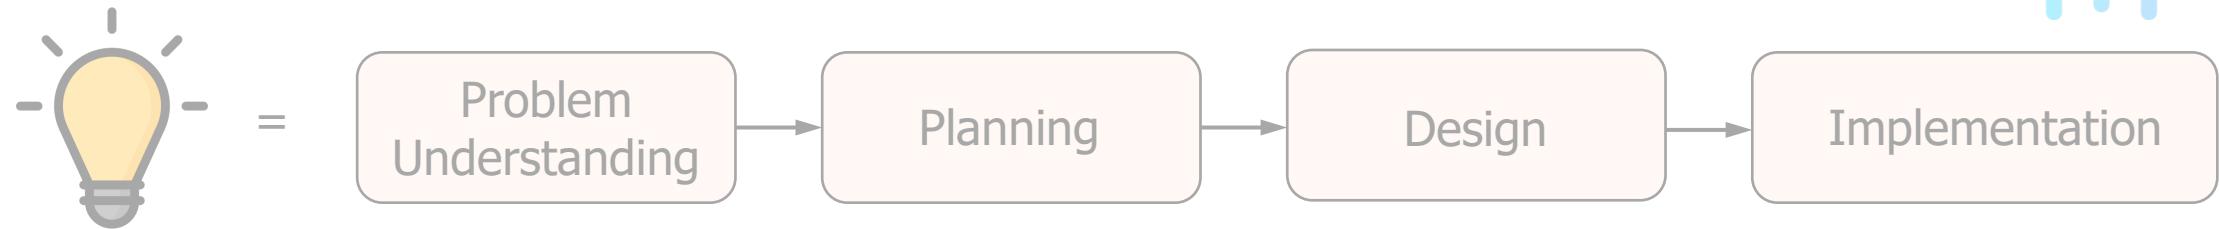
Planning

Design

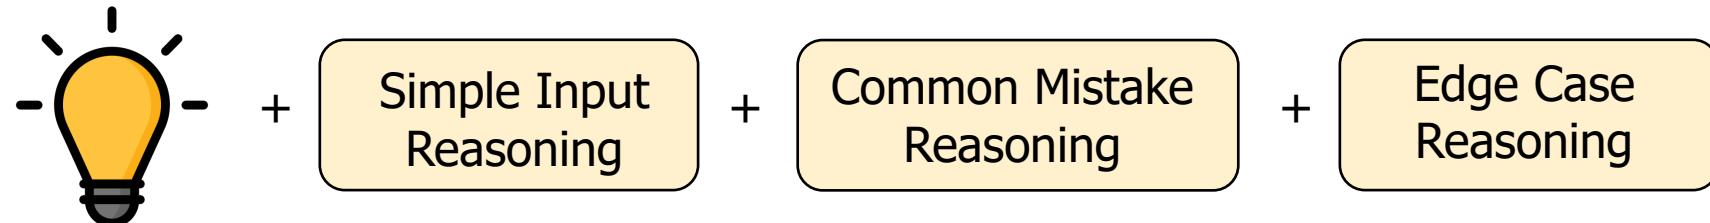
Implementation

Step III: Enriching Loosely Structured Reasoning Traces

-  Rewrite a post into a structured trace inspired by SDLC



-  Add iterative refinement considerations



Step III: Enriching Loosely Structured Reasoning Traces

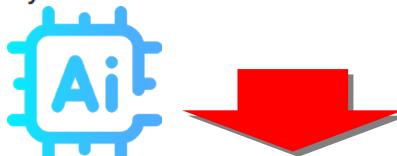
Given two binary grids as quad trees, return their logical OR result as a quad tree while maintaining minimal node representation.

User Discussion:

Intuition

Approach

1. Check if either quadTree1 or quadTree2 is a leaf node...
2. Recursively check the intersection of each child node of



First things first, I need to understand what a quad-tree is. From the description, it's a tree data structure where each internal node has exactly four children ...

Problem Understanding

Let's think about this step by step. First, I need to understand how the quad-tree represents the matrix. If a node is a leaf and its `val` is true, it means ...

Planning

Let me try to outline an approach:
 1. If `quadTree1` is a leaf:
 - If its `val` is true, then the entire grid is 1's, so the OR with anything else is still 1's ...

Design

Let me try to write some solutions:
 function intersect(quadTree1, quadTree2):
 if quadTree1.isLeaf:
 if quadTree1.val:
 return quadTree1
 ...

Implementation

I should think about edge cases:

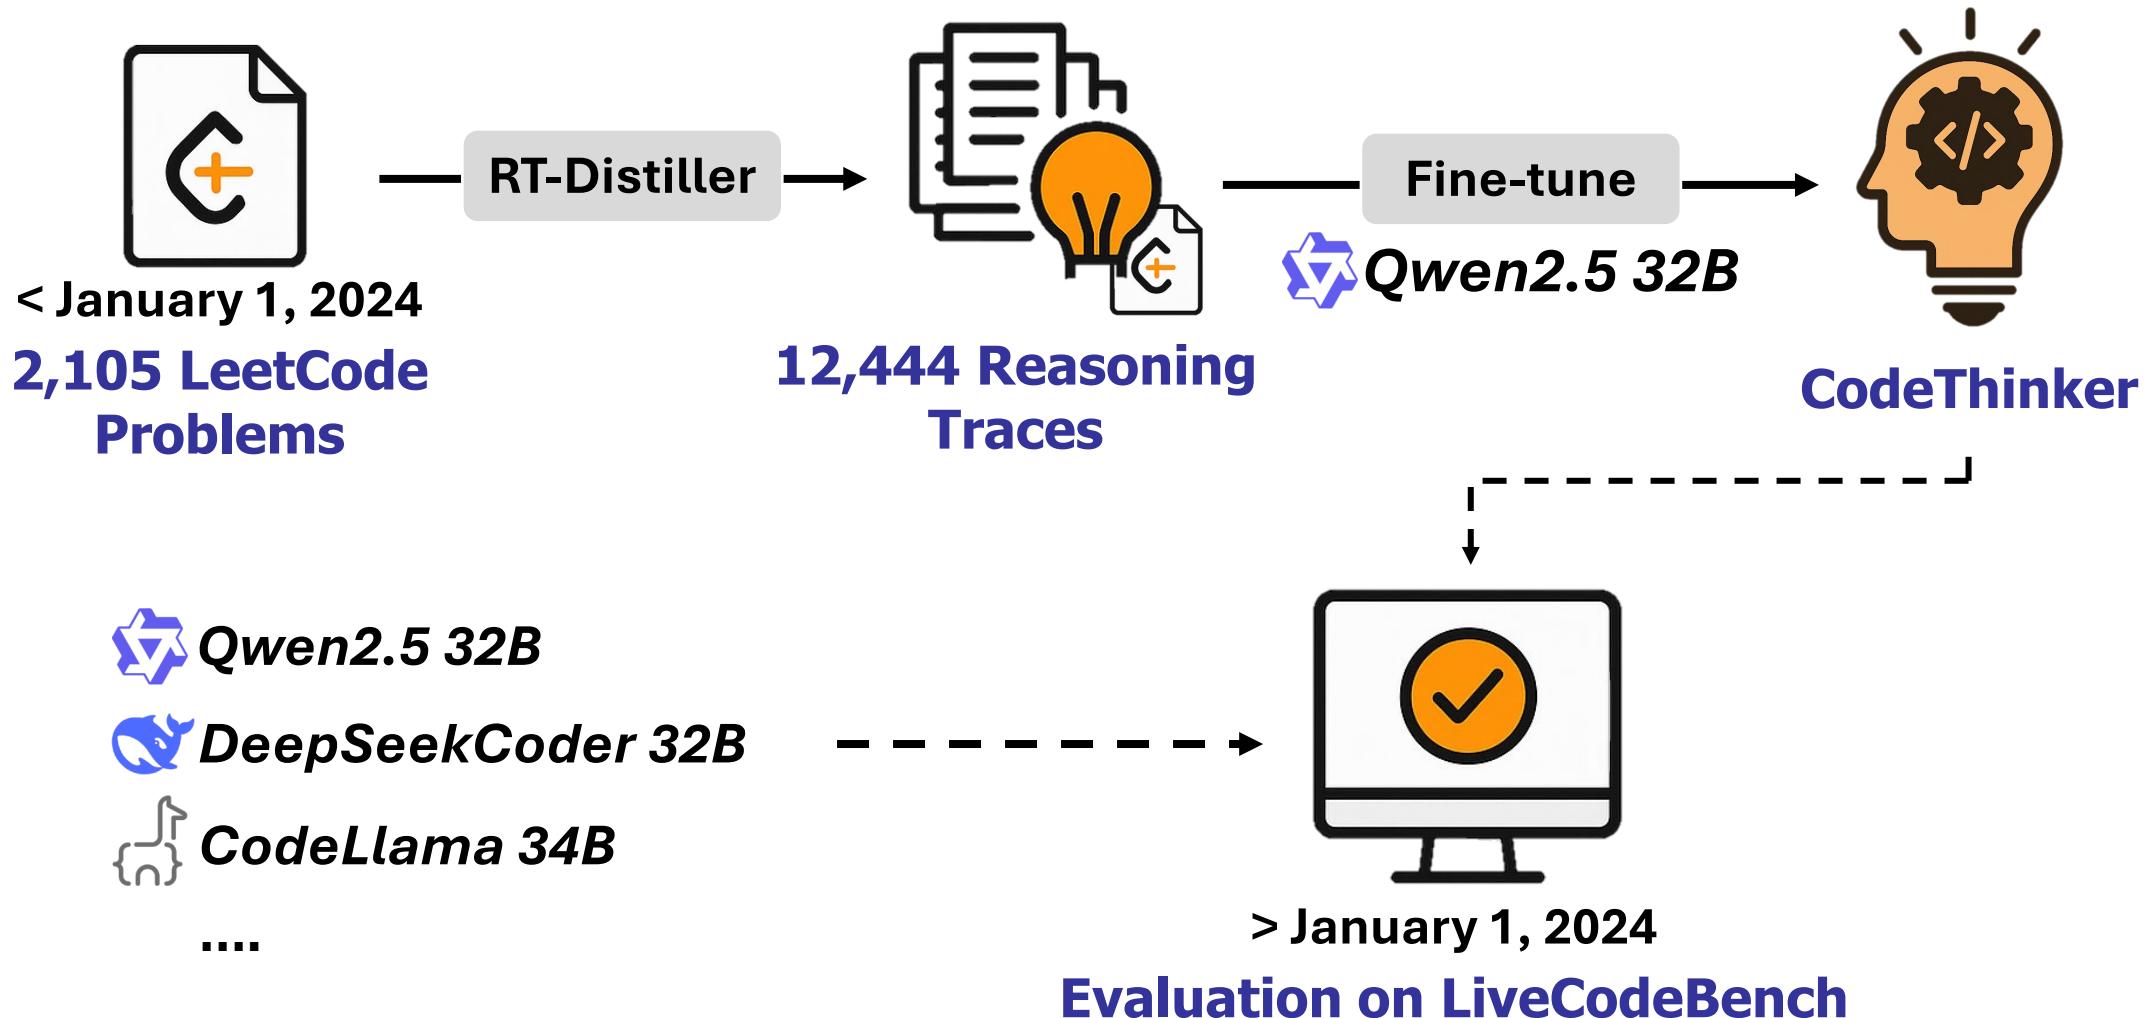
- Both trees are leaf nodes.
- One tree is a leaf, the other is ...

I need to make sure that the function handles all these cases correctly.
 Let me consider a simple 2x2 grid...

Now, I need to think about whether there are any potential issues or edge cases that this doesn't handle...

Iterative Refinement

Experiment Settings

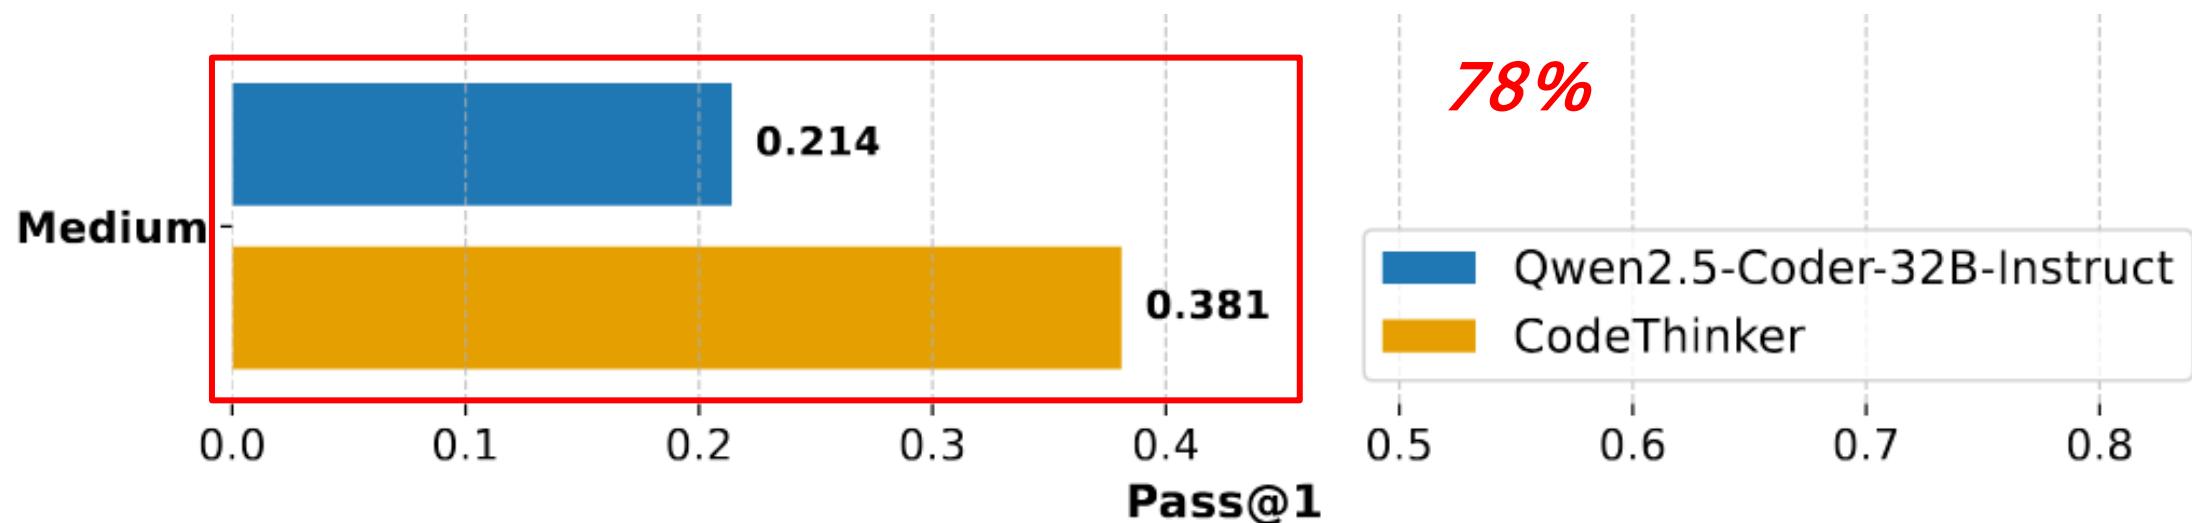


Experiment Results – Comparison with LLMs of Similar Sizes

Model	Prompting	Size	Easy pass@1	Medium pass@1	Hard pass@1	Overall pass@1
Open-source LLMs						
CodeLlama-Instruct	Standard	34B	0.290	0.028	0	0.078
CodeLlama-Instruct	CoT	34B	0.298	0.028	0	0.106
s1-32B	Standard	32B	0.306	0.098	0	0.134
s1-32B	CoT	32B	0.315	0.112	0	0.142
DeepSeekCoder	Standard	32B	0.589	0.126	0	0.235
DeepSeekCoder	CoT	32B	0.565	0.133	0	0.233
Qwen2.5-Instruct	Standard	32B	0.847	0.343	0	0.398
Owen2.5-Instruct	CoT	32B	0.823	0.371	0	0.401
CodeThinker	CodeThinker Style	32B	0.831 (↓1.89%)	0.490(↑42.86%)	0	0.447(↑12.31%)

Experiment Results – Generalizability Beyond LeetCode

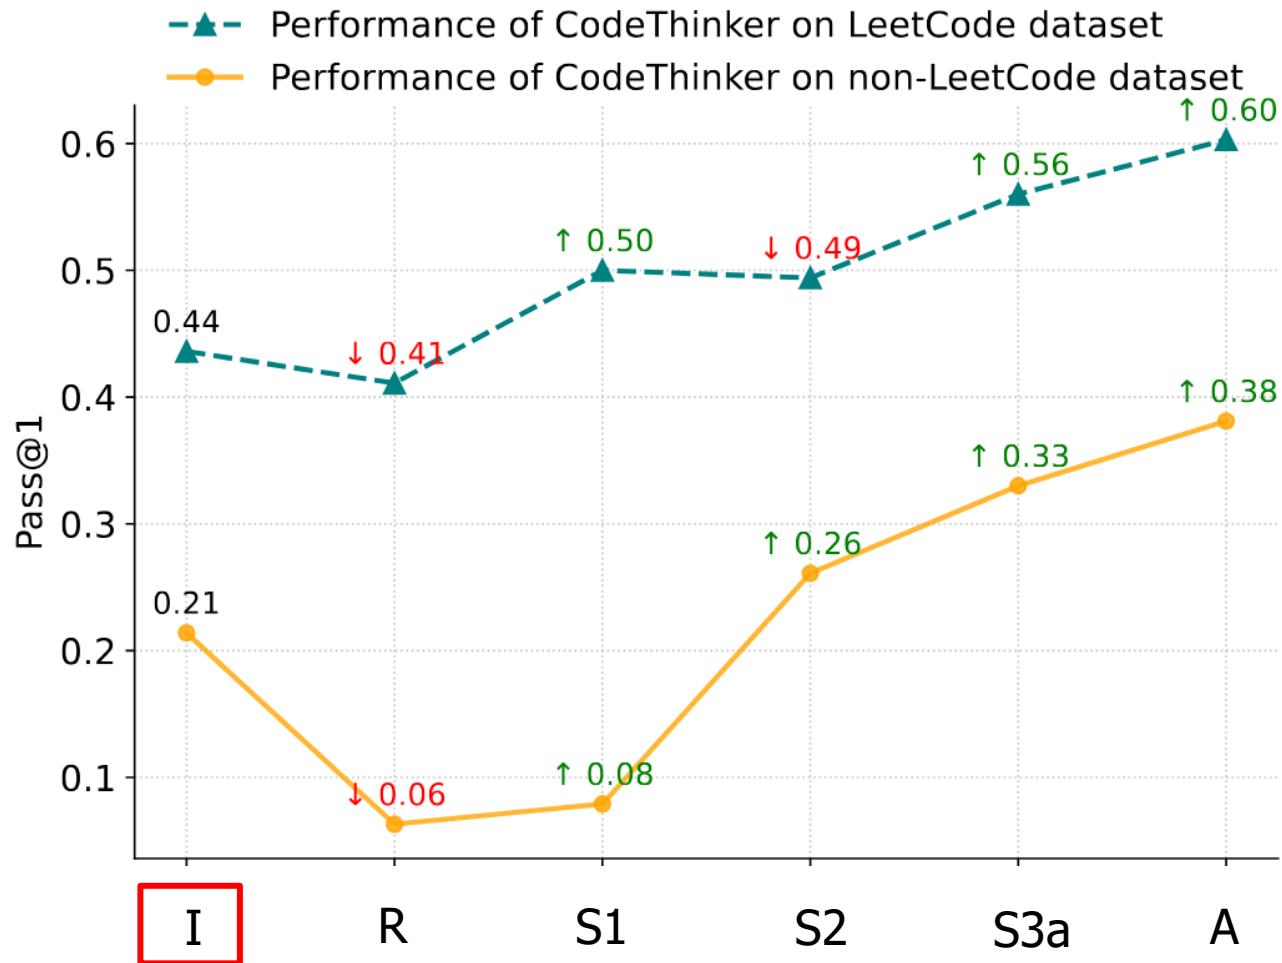
Figure 3: Comparison of CodeThinker and its base model on non-LeetCode coding problems (AtCoder and Codeforces), highlighting improved reasoning performance on medium-difficulty tasks.



CodeThinker shows **strong generalization ability** to problems from **non-LeetCode platforms**, achieving significant improvements

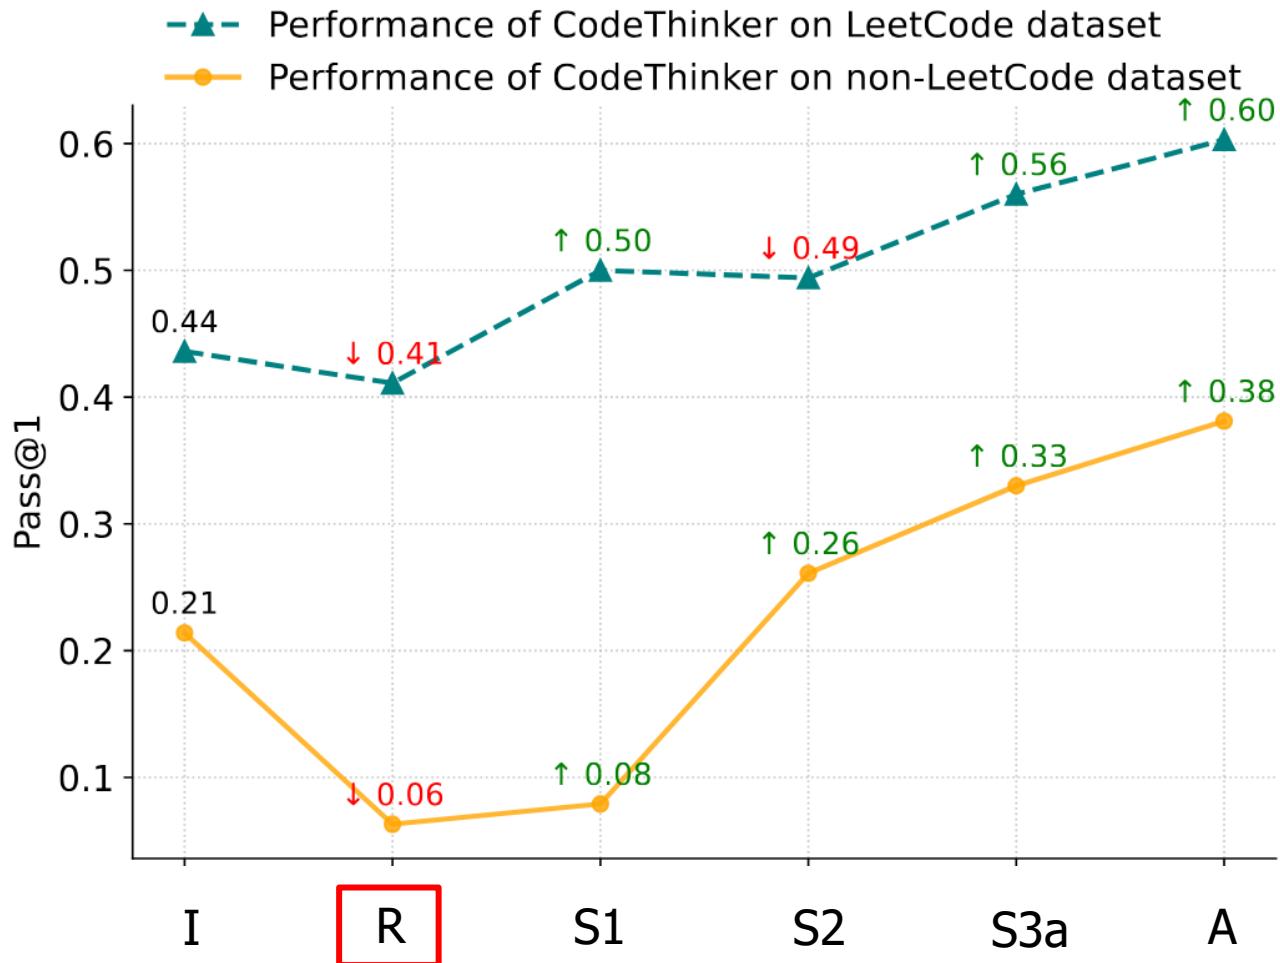
Experiment Results – Ablation Study

*I: Initial checkpoint; R: Fine tune using raw LeetCode discussion posts;
S1: Step 1 only; S2: Step 1+2 only; S3a: Step 1+2+3a only; A: All steps of RT-Distiller



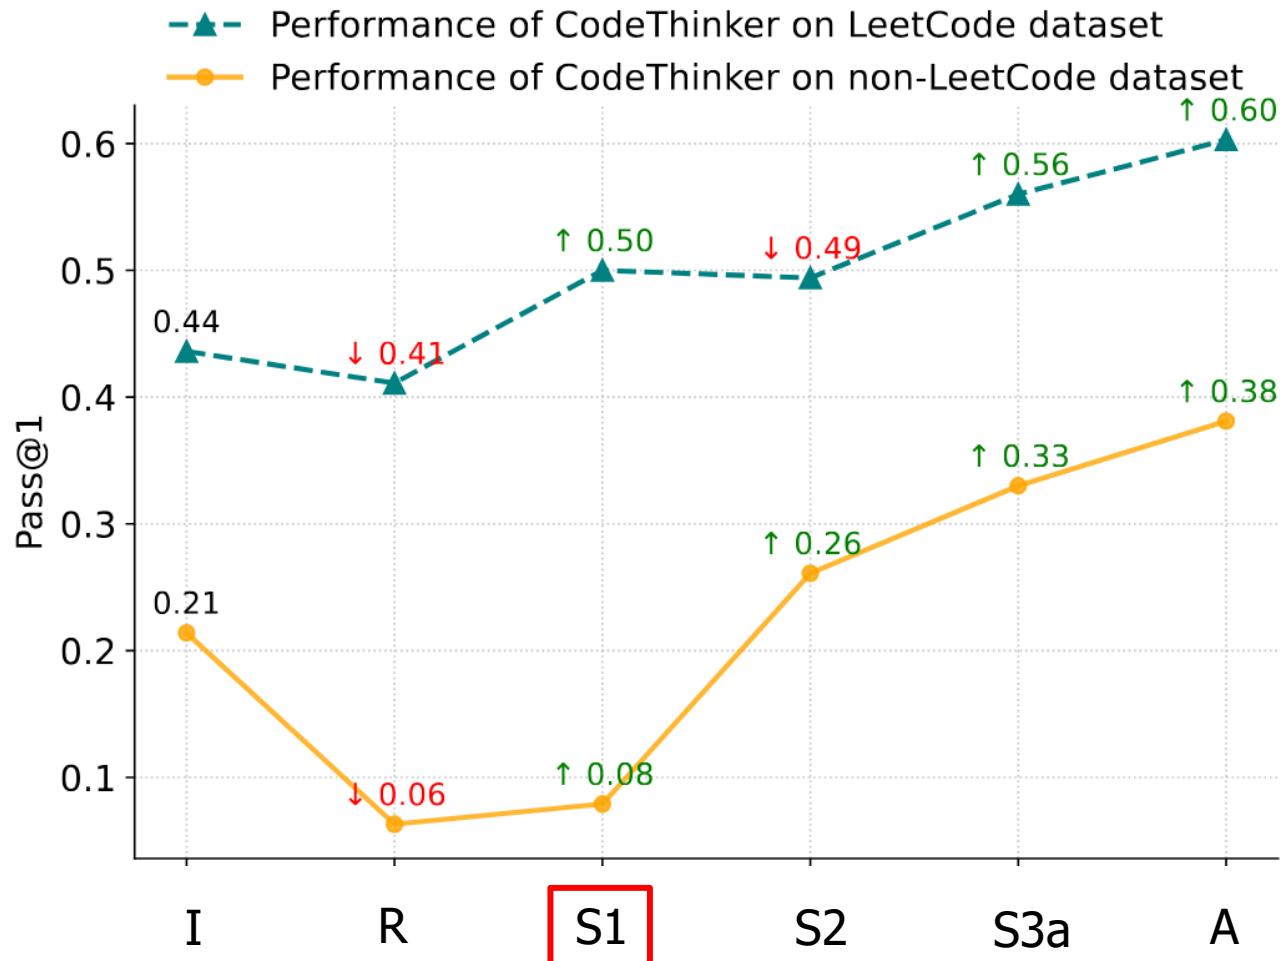
Experiment Results – Ablation Study

*I: Initial checkpoint; R: Fine tune using raw LeetCode discussion posts;
S1: Step 1 only; S2: Step 1+2 only; S3a: Step 1+2+3a only; A: All steps of RT-Distiller



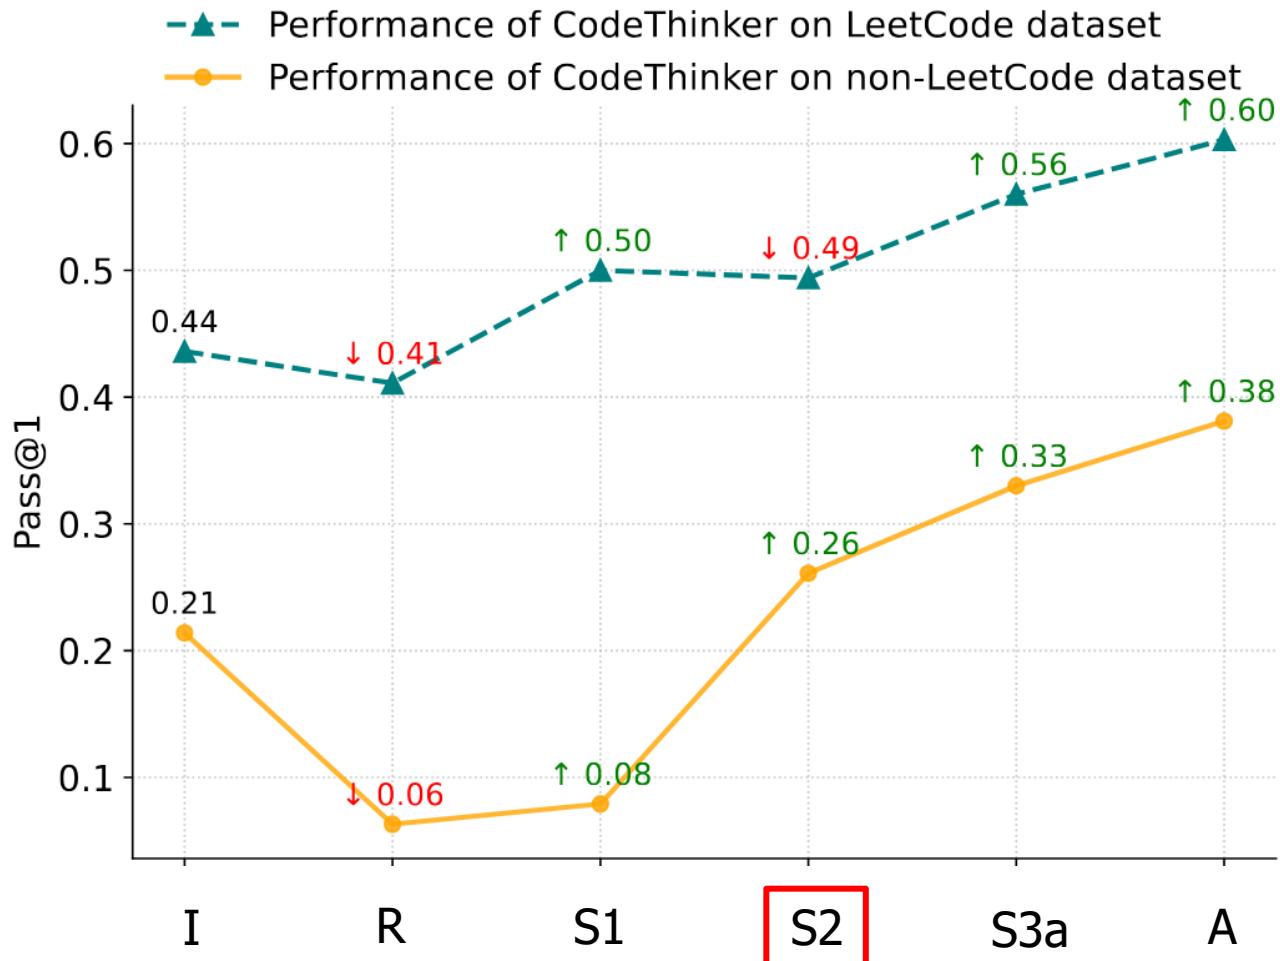
Experiment Results – Ablation Study

*I: Initial checkpoint; R: Fine tune using raw LeetCode discussion posts;
S1: Step 1 only; S2: Step 1+2 only; S3a: Step 1+2+3a only; A: All steps of RT-Distiller



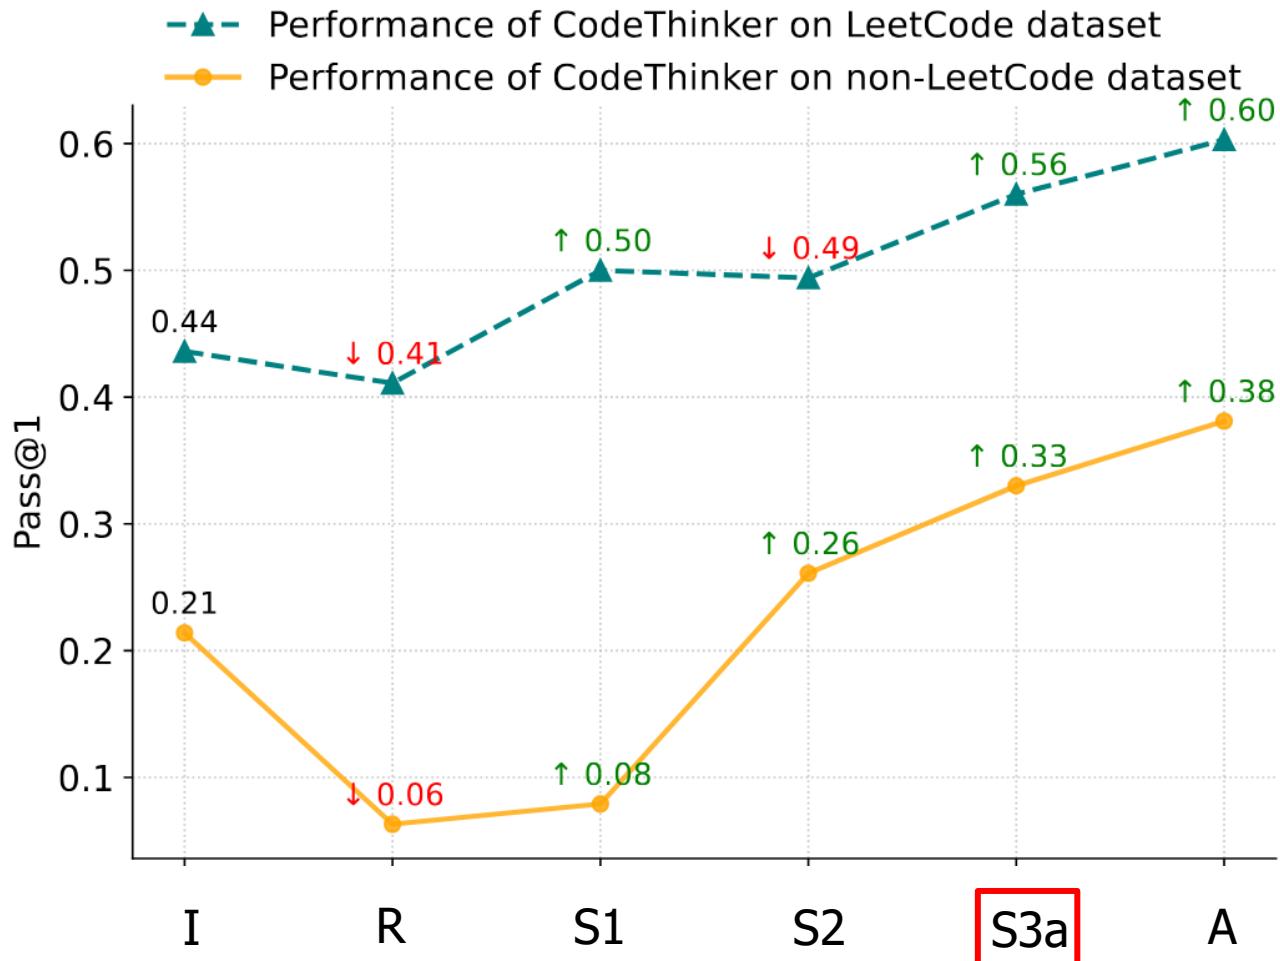
Experiment Results – Ablation Study

*I: Initial checkpoint; R: Fine tune using raw LeetCode discussion posts;
S1: Step 1 only; S2: Step 1+2 only; S3a: Step 1+2+3a only; A: All steps of RT-Distiller



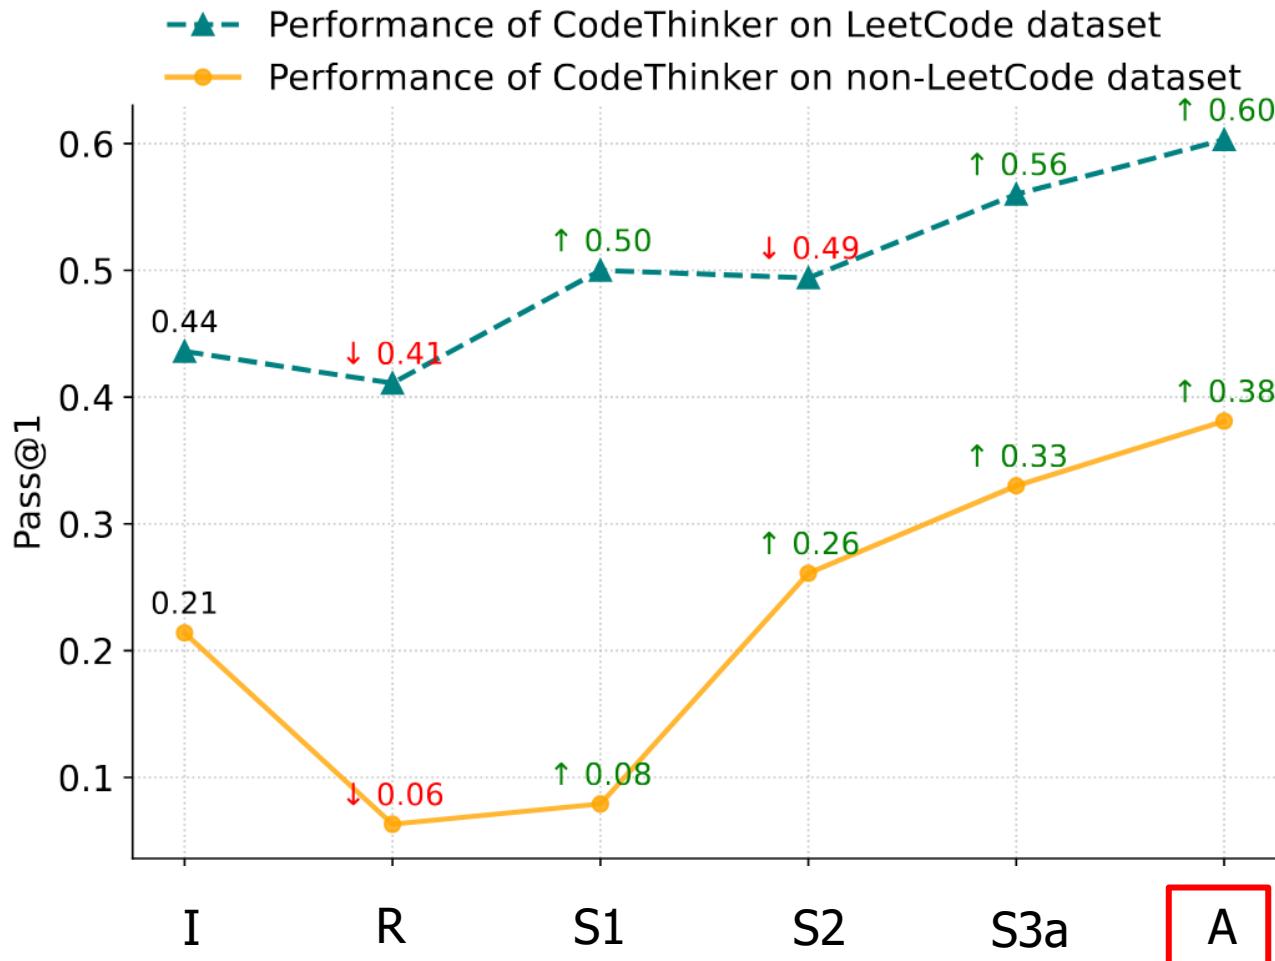
Experiment Results – Ablation Study

*I: Initial checkpoint; R: Fine tune using raw LeetCode discussion posts;
S1: Step 1 only; S2: Step 1+2 only; S3a: Step 1+2+3a only; A: All steps of RT-Distiller



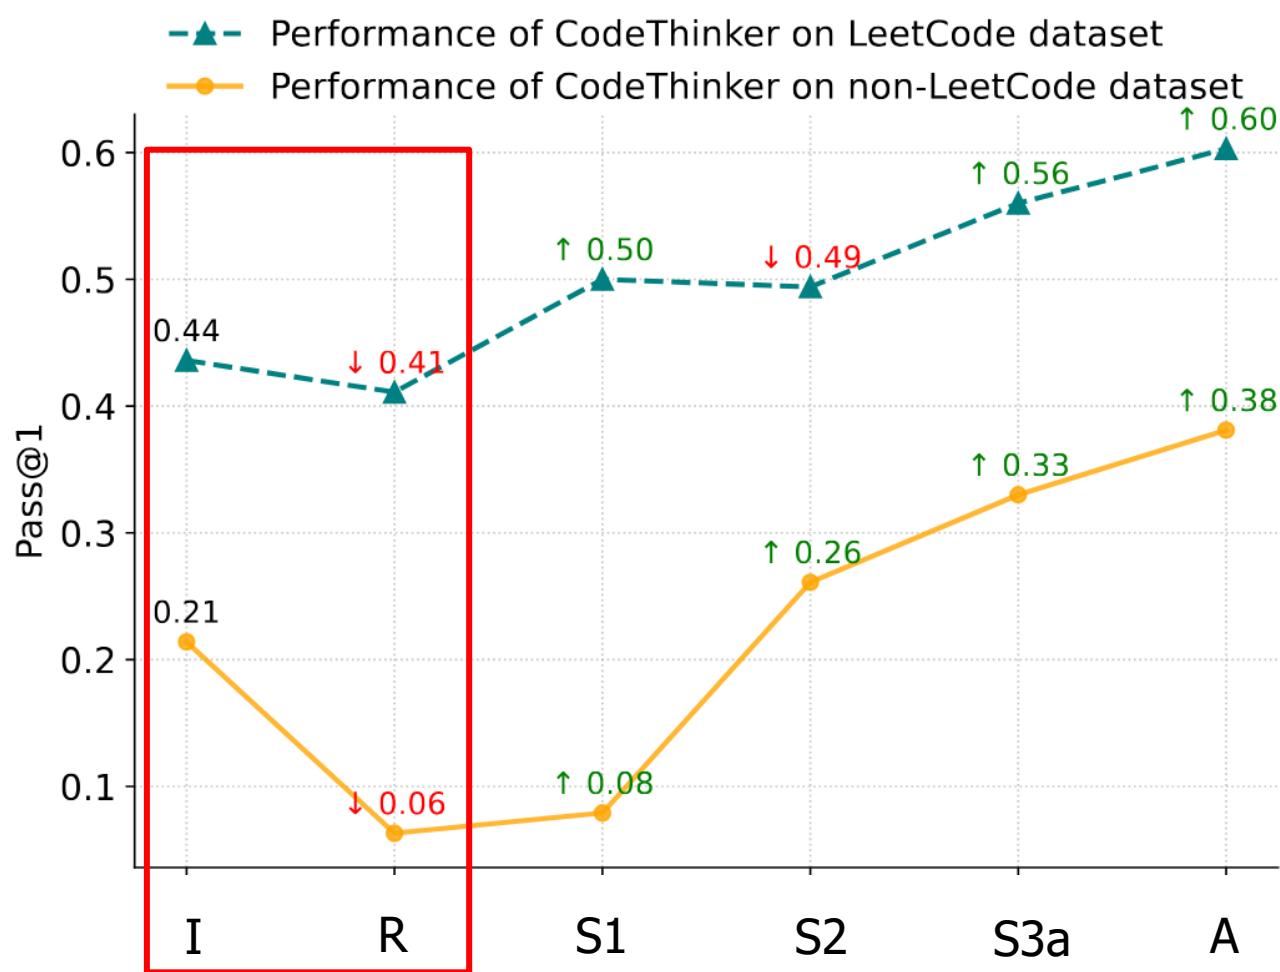
Experiment Results – Ablation Study

*I: Initial checkpoint; R: Fine tune using raw LeetCode discussion posts;
S1: Step 1 only; S2: Step 1+2 only; S3a: Step 1+2+3a only; A: All steps of RT-Distiller



Experiment Results – Ablation Study

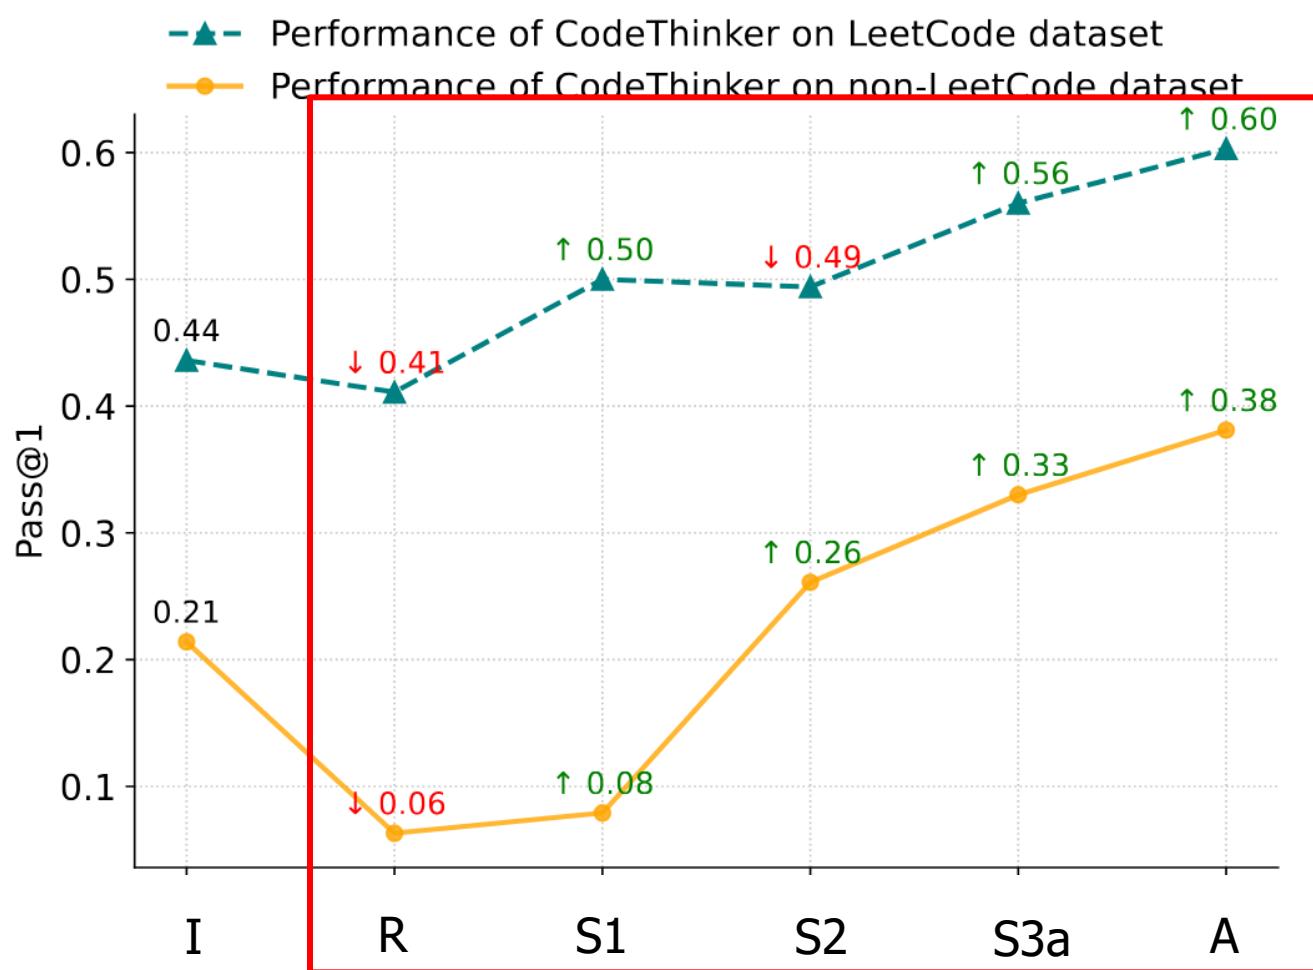
*I: Initial checkpoint; R: Fine tune using raw LeetCode discussion posts;
S1: Step 1 only; S2: Step 1+2 only; S3a: Step 1+2+3a only; A: All steps of RT-Distiller



I vs. R: Training with raw LeetCode discussion posts results in performance drop

Experiment Results – Ablation Study

*I: Initial checkpoint; R: Fine tune using raw LeetCode discussion posts;
 S1: Step 1 only; S2: Step 1+2 only; S3a: Step 1+2+3a only; A: All steps of RT-Distiller



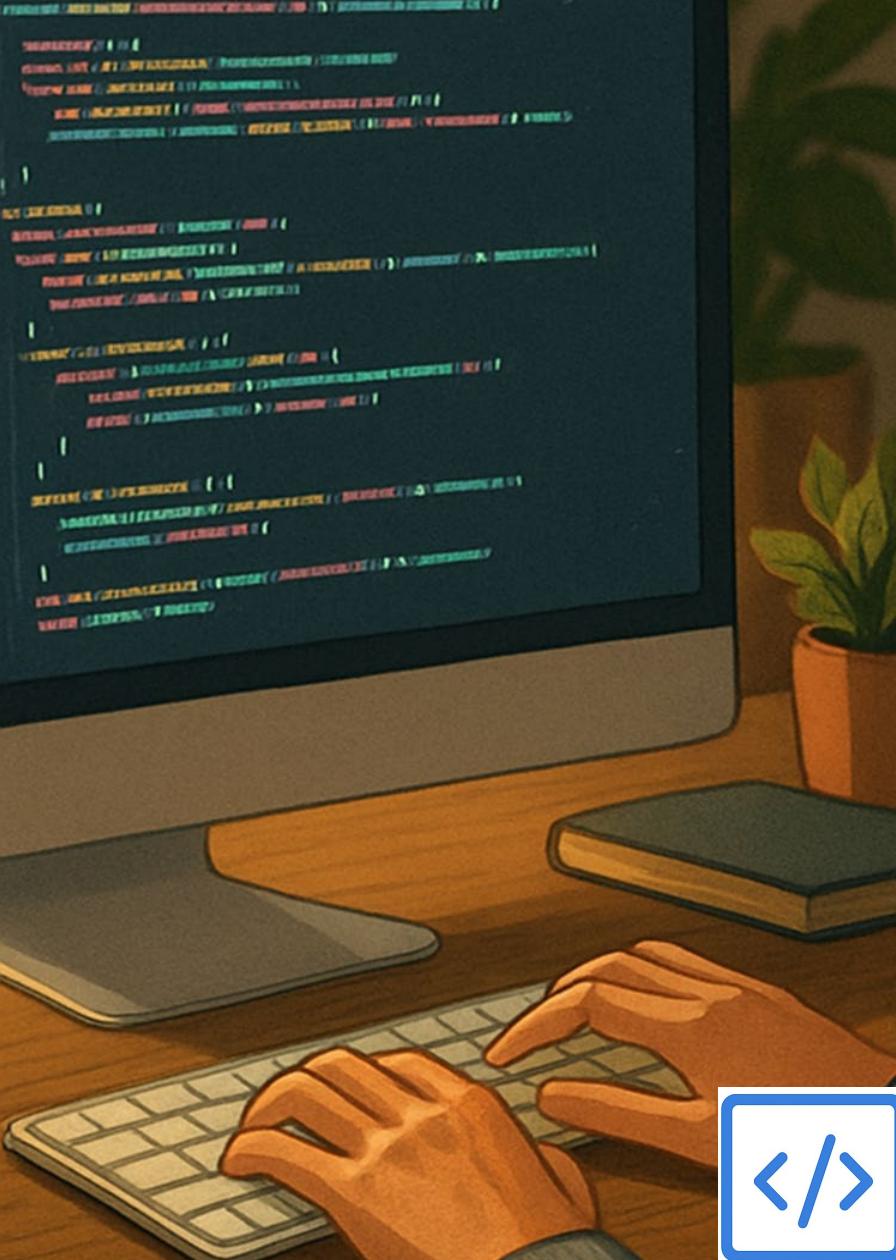
I vs. R: Training with raw LeetCode discussion posts results in performance drop

R vs. S1: Synthesizing missing information in LeetCode discussion posts helps

S1 vs. S2: Applying code perturbation improves generalization to non-LeetCode

S2 vs. S3a: Introducing SDLC-inspired structure to reasoning traces helps

S3a vs. A: Adding iterative refinement considerations helps



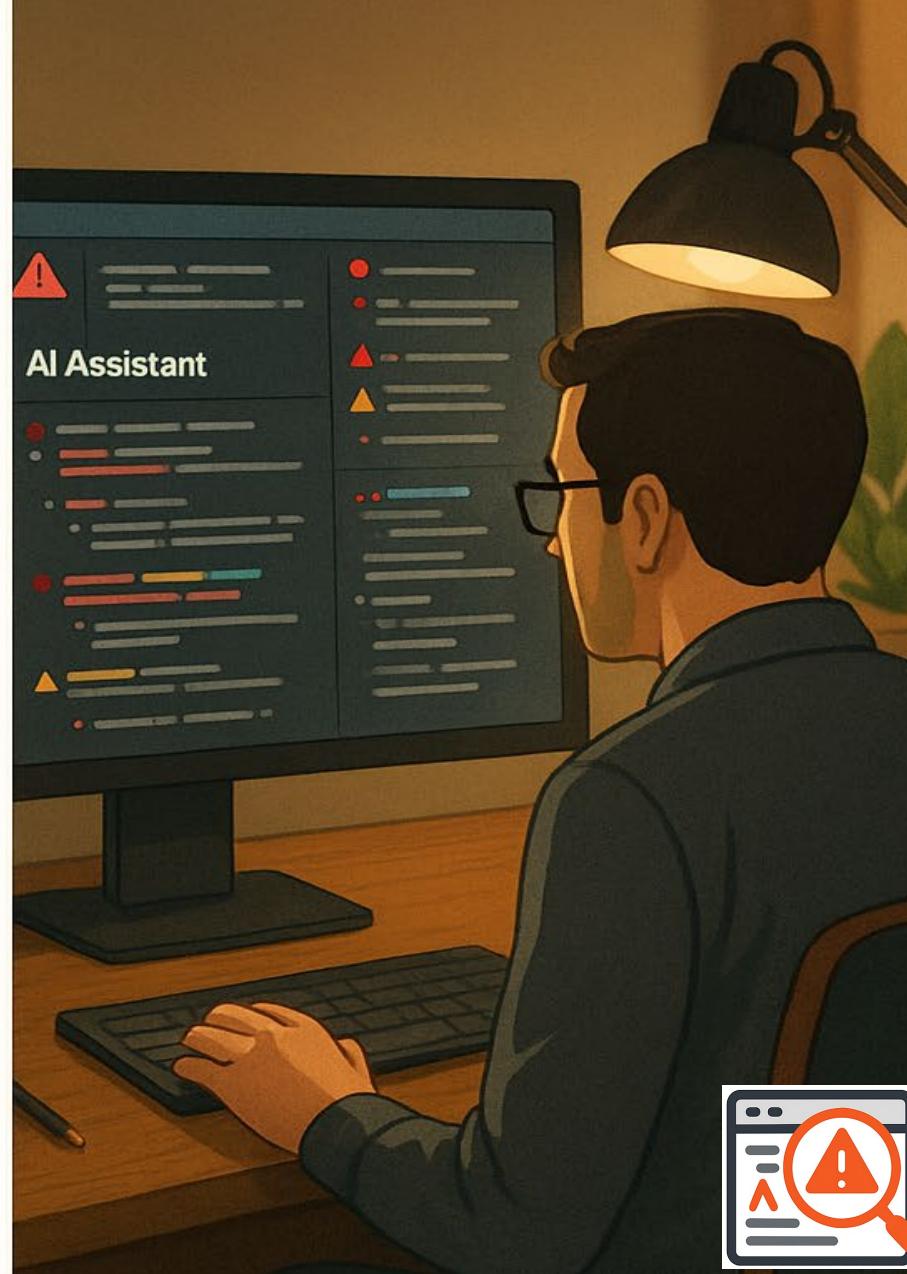
Code



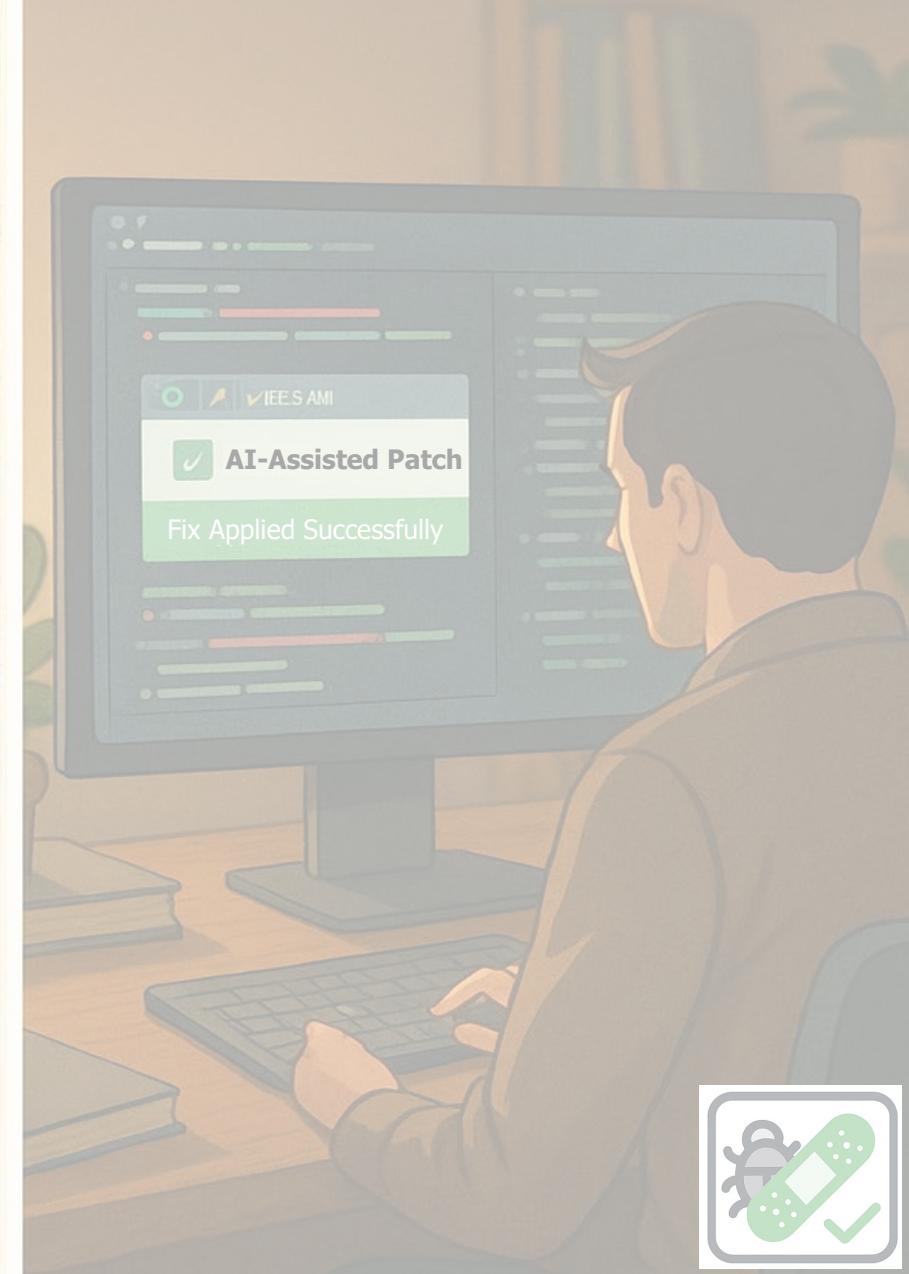
Engineering **structured reasoning traces** help LLM reasoning for code generation



Code



Critique



Cure



Contrastive Patterns Can Help AI Reasoning

KDD 2009

Classification of Software Behaviors for Failure Detection: A Discriminative Pattern Mining Approach

David Lo
Singapore Management University
davidlo@smu.edu.sg

Hong Cheng^{*}
Chinese University of Hong Kong
hcheng@se.cuhk.edu.hk

Jiawei Han
University of Illinois at Urbana-Champaign
hanj@cs.uiuc.edu

Siau-Cheng Khoo and Chengnian Sun
National University of Singapore
{khoosc,suncn}@comp.nus.edu.sg



*Distills rich discriminative (contrastive) patterns for AI-powered failure detection
(aka. test oracle generation); early SE paper in a top AI venue*

Distilling Contrastive Reasoning Trace Pairs

R2VUL: Learning to Reason about Software Vulnerabilities with Reinforcement Learning and Structured Reasoning Distillation

Martin Weyssow^{1*}, Chengran Yang¹, Junkai Chen¹, Ratnadira Widyasari¹, Ting Zhang¹, Huihui Huang¹, Huu Hung Nguyen¹, Yan Naing Tun¹, Tan Bui¹, Yikun Li¹, Ang Han Wei², Frank Liauw², Eng Lieh Ouh¹, Lwin Khin Shar¹, David Lo¹

¹Singapore Management University

²GovTech Singapore



*Engineers contrastive reasoning trace pairs
to help LLM reason about vulnerabilities*

Vulnerability Detection with LLMs

- Vulnerability detection demands a *binary judgement*:
 - Is the input code **safe** or **vulnerable**?

TOSEM 2025

Large Language Model for Vulnerability Detection and Repair: Literature Review and the Road Ahead

XIN ZHOU, School of Computing and Information Systems, Singapore Management University, Singapore, Singapore

SICONG CAO and **XIAOBING SUN**, Yangzhou University, Yangzhou, China

DAVID LO, School of Computing and Information Systems, Singapore Management University, Singapore, Singapore



Vulnerability Detection with LLMs

- Limited efficacy

ICSE 2024

Large Language Model for Vulnerability Detection: Emerging Results and Future Directions

Xin Zhou, Ting Zhang, and David Lo

School of Computing and Information Systems, Singapore Management University, Singapore
[xinzhou.2020,tingzhang.2019}@phdcs.smu.edu.sg](mailto:{xinzhou.2020,tingzhang.2019}@phdcs.smu.edu.sg), davidlo@smu.edu.sg



- From a developer viewpoint, a binary judgement is also not sufficient:

Why is this input code vulnerable or safe?

What is the mechanism of a vulnerability and its **impact**?

R2Vul: Beyond Binary Judgement

Fine-tune an LLM to **better detect vulnerabilities** and generate **structured reasoning**.

- The LLM outputs a structured reasoning covering key aspects of code safety and vulnerability.

Vulnerable Code

- Discuss specific code constructs responsible for the vulnerability
- Explain the mechanism of the vulnerability
- Discuss its potential impact
- Relate the vulnerability to a relevant CWE

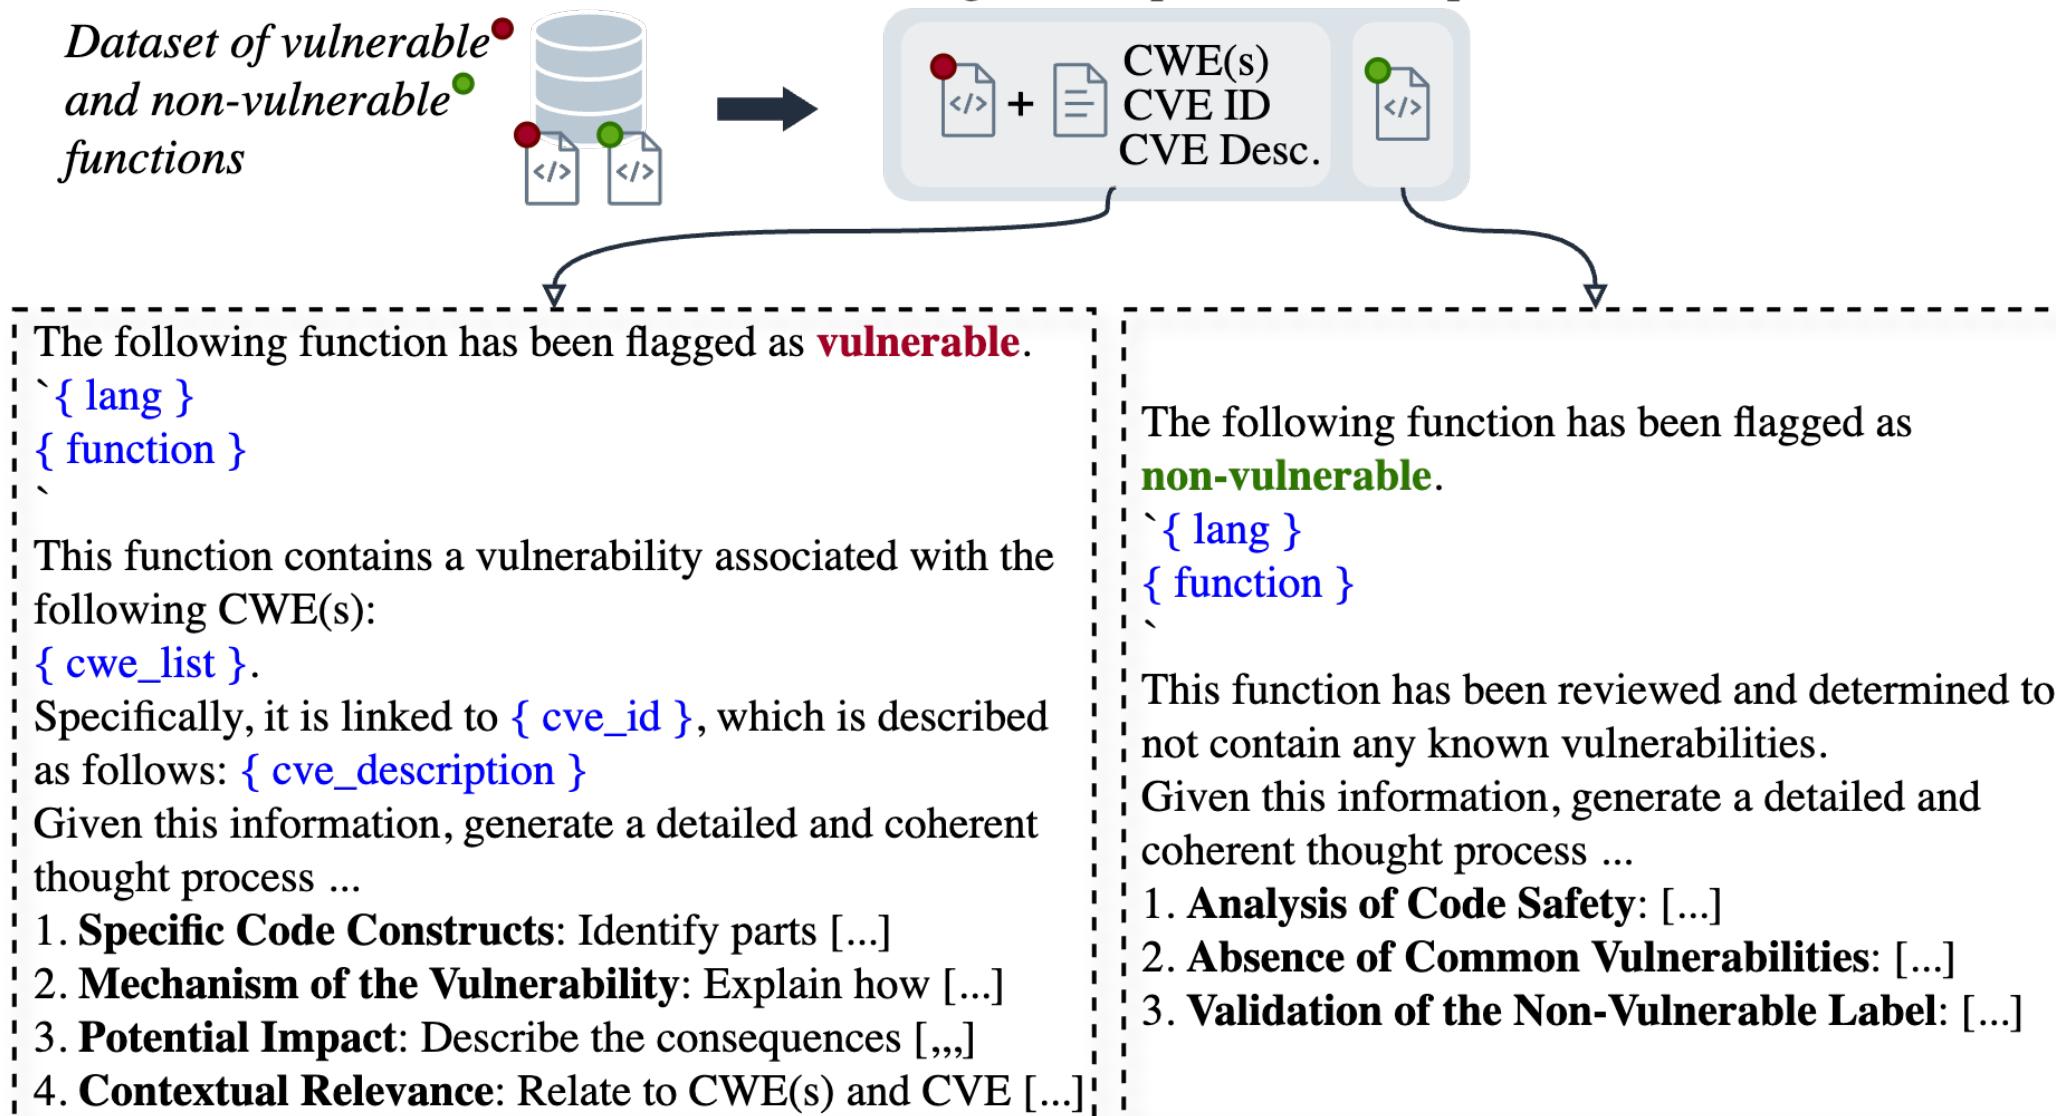
Safe Code

- Discuss key aspects contributing to code safety
- Discuss the absence of key vulnerabilities
- Provide evidence-based justification why a code is safe

Step 1: Generate Contrastive Reasoning Trace Pairs

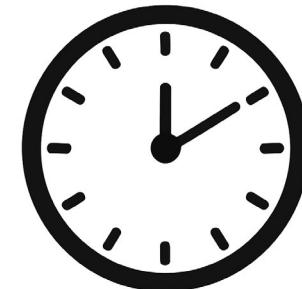
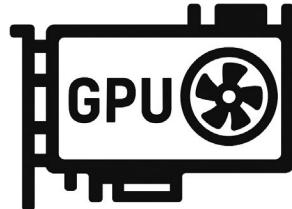
- Given a labeled dataset $\mathcal{D} = \{(x_i, y_i)\}_{i=1}^N$
 - (x_i : function, $y_i \in \{Vuln, Safe\}$)
- Prompt a base LLM to generate:
 - A **valid** reasoning r_i^+ conditioned on the *true* sample label y_i
 - A **flawed** reasoning r_i^- conditioned on the *flipped* sample label $-y_i$ $\rightarrow \mathcal{D} = \{(x_i, y_i, r_i^+, r_i^-)\}_{i=1}^N$ = preference dataset

Step 1: Generate Contrastive Reasoning Trace Pairs



Step 2: Use Contrastive Reasoning Trace Pairs

- Combines supervised fine-tuning (SFT) and odds-ratio (OR) losses
 - SFT: trains the model to generate high-quality structured reasoning
 - **OR: distinguish valid vs. flawed reasoning**
- Additional feature: Distill **smaller** models
 - constraints from deployment environment

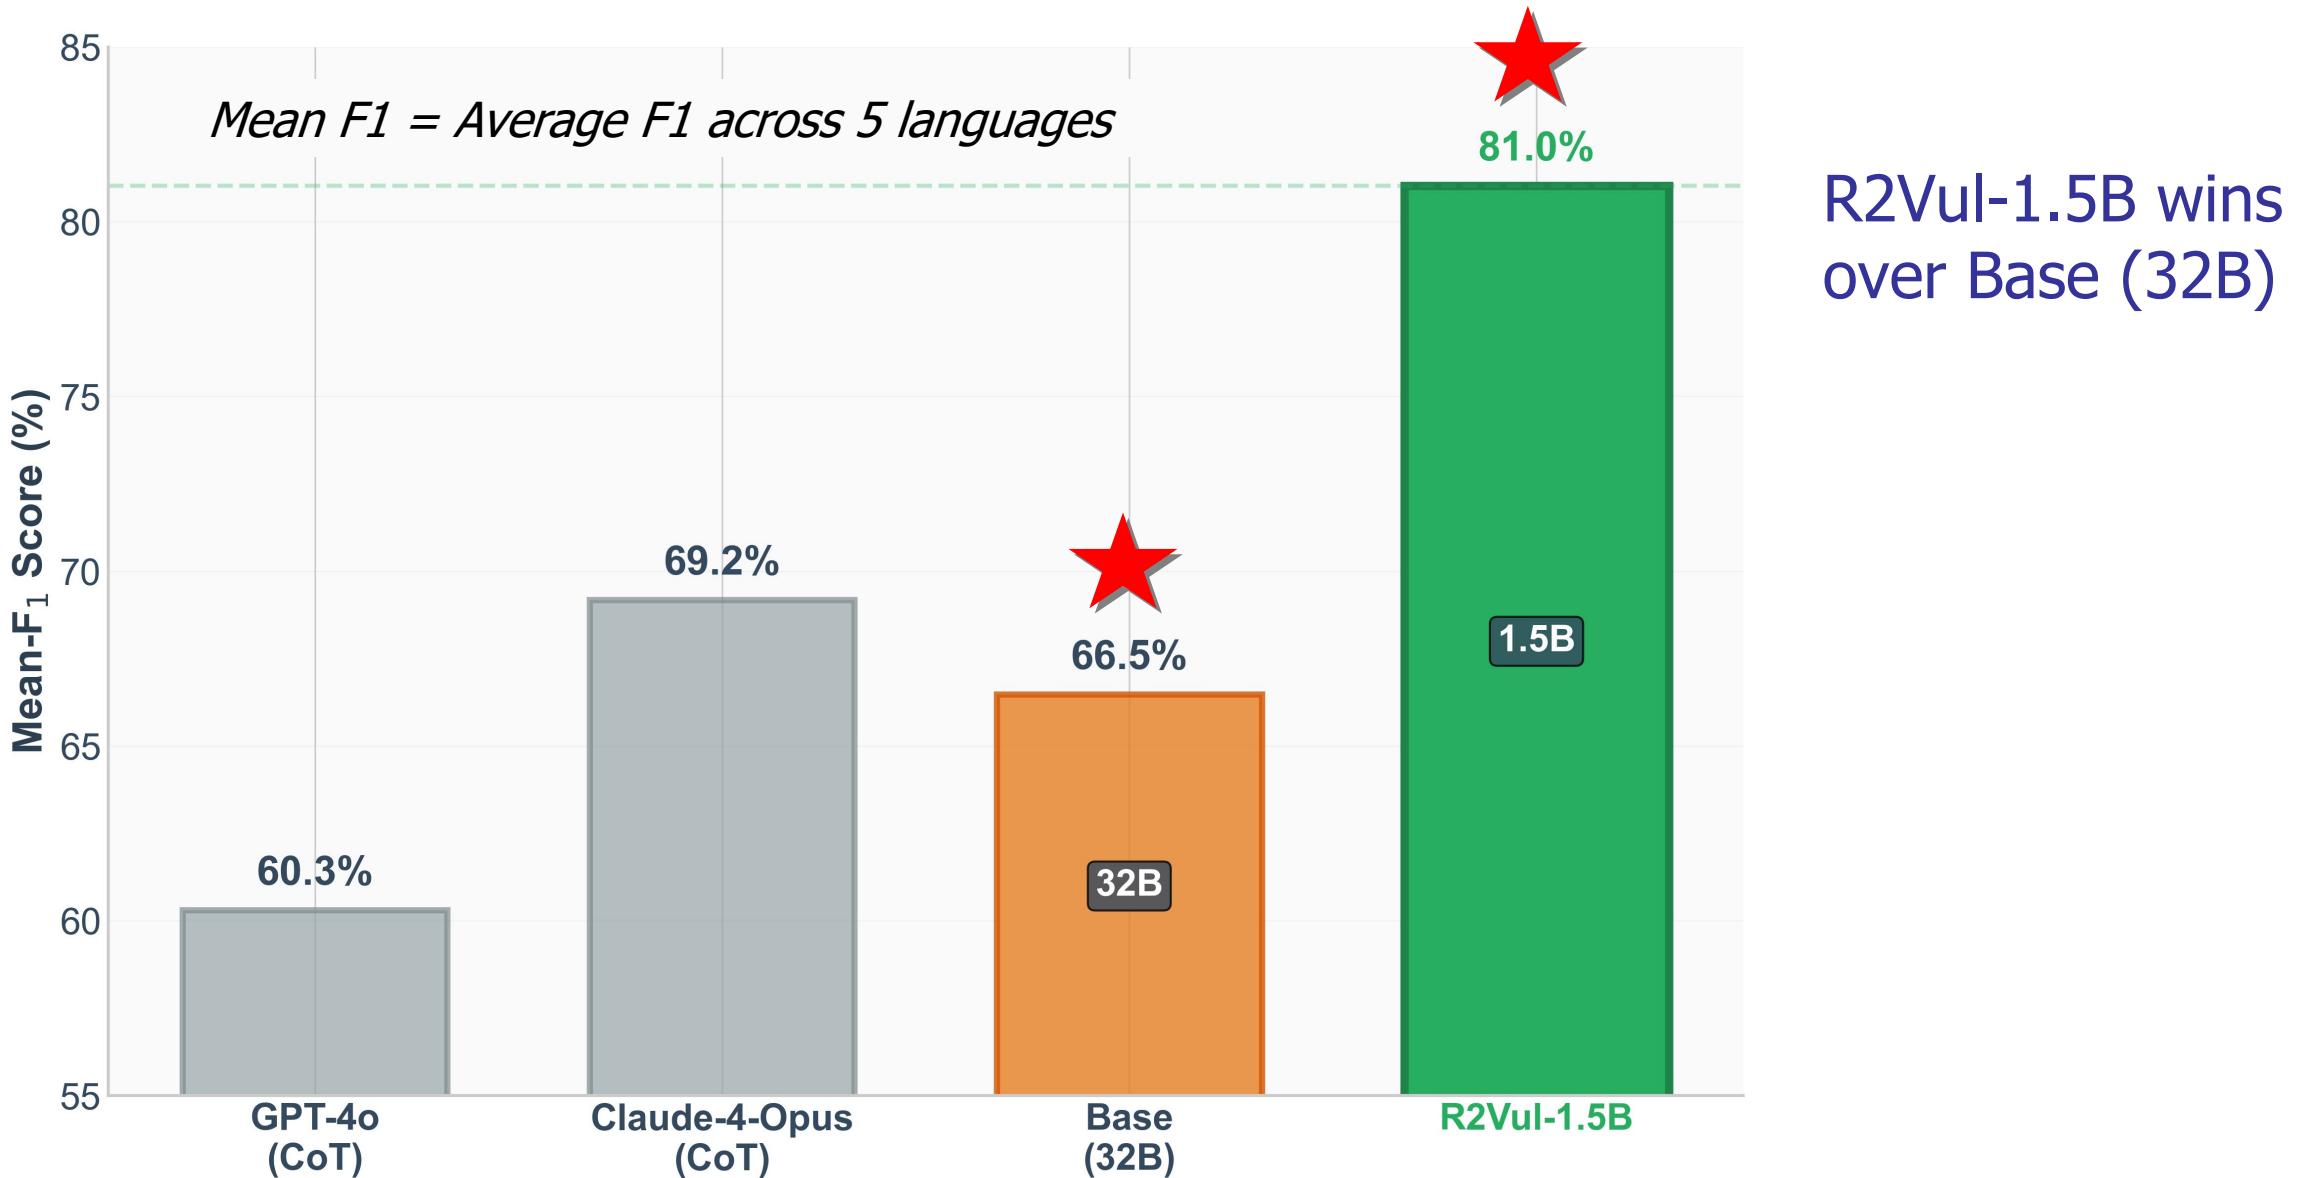


Experiment Details

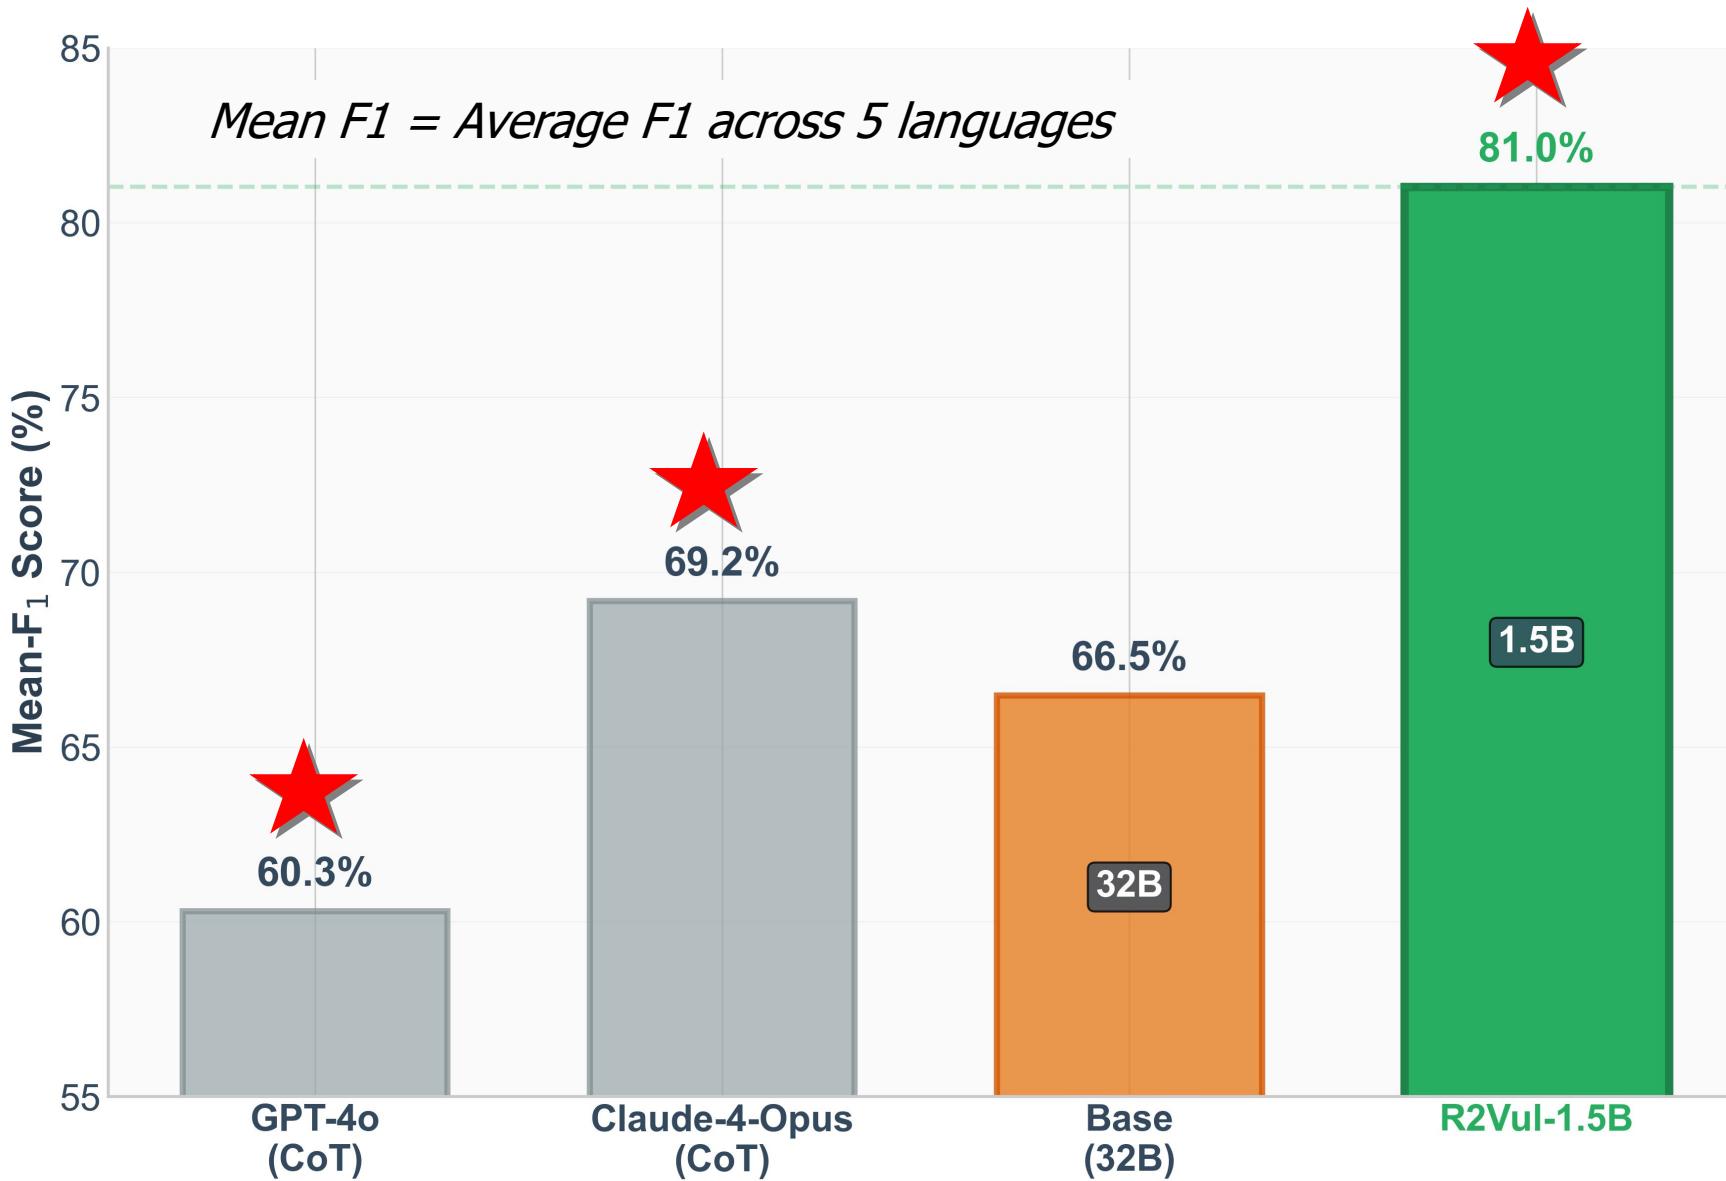
- Step 1:
 - Base LLM: Qwen2.5-Coder-Instruct with 32B parameters
 - About **18,000 pairs** of contrastive traces
 - **Five** programming languages: C#, JavaScript, Java, Python, and C

- Step 2:
 - Target LLMs: Qwen2.5-Coder-Instruct with 0.5B, 1.5B, and 7B parameters

Main Results: Detection Performance

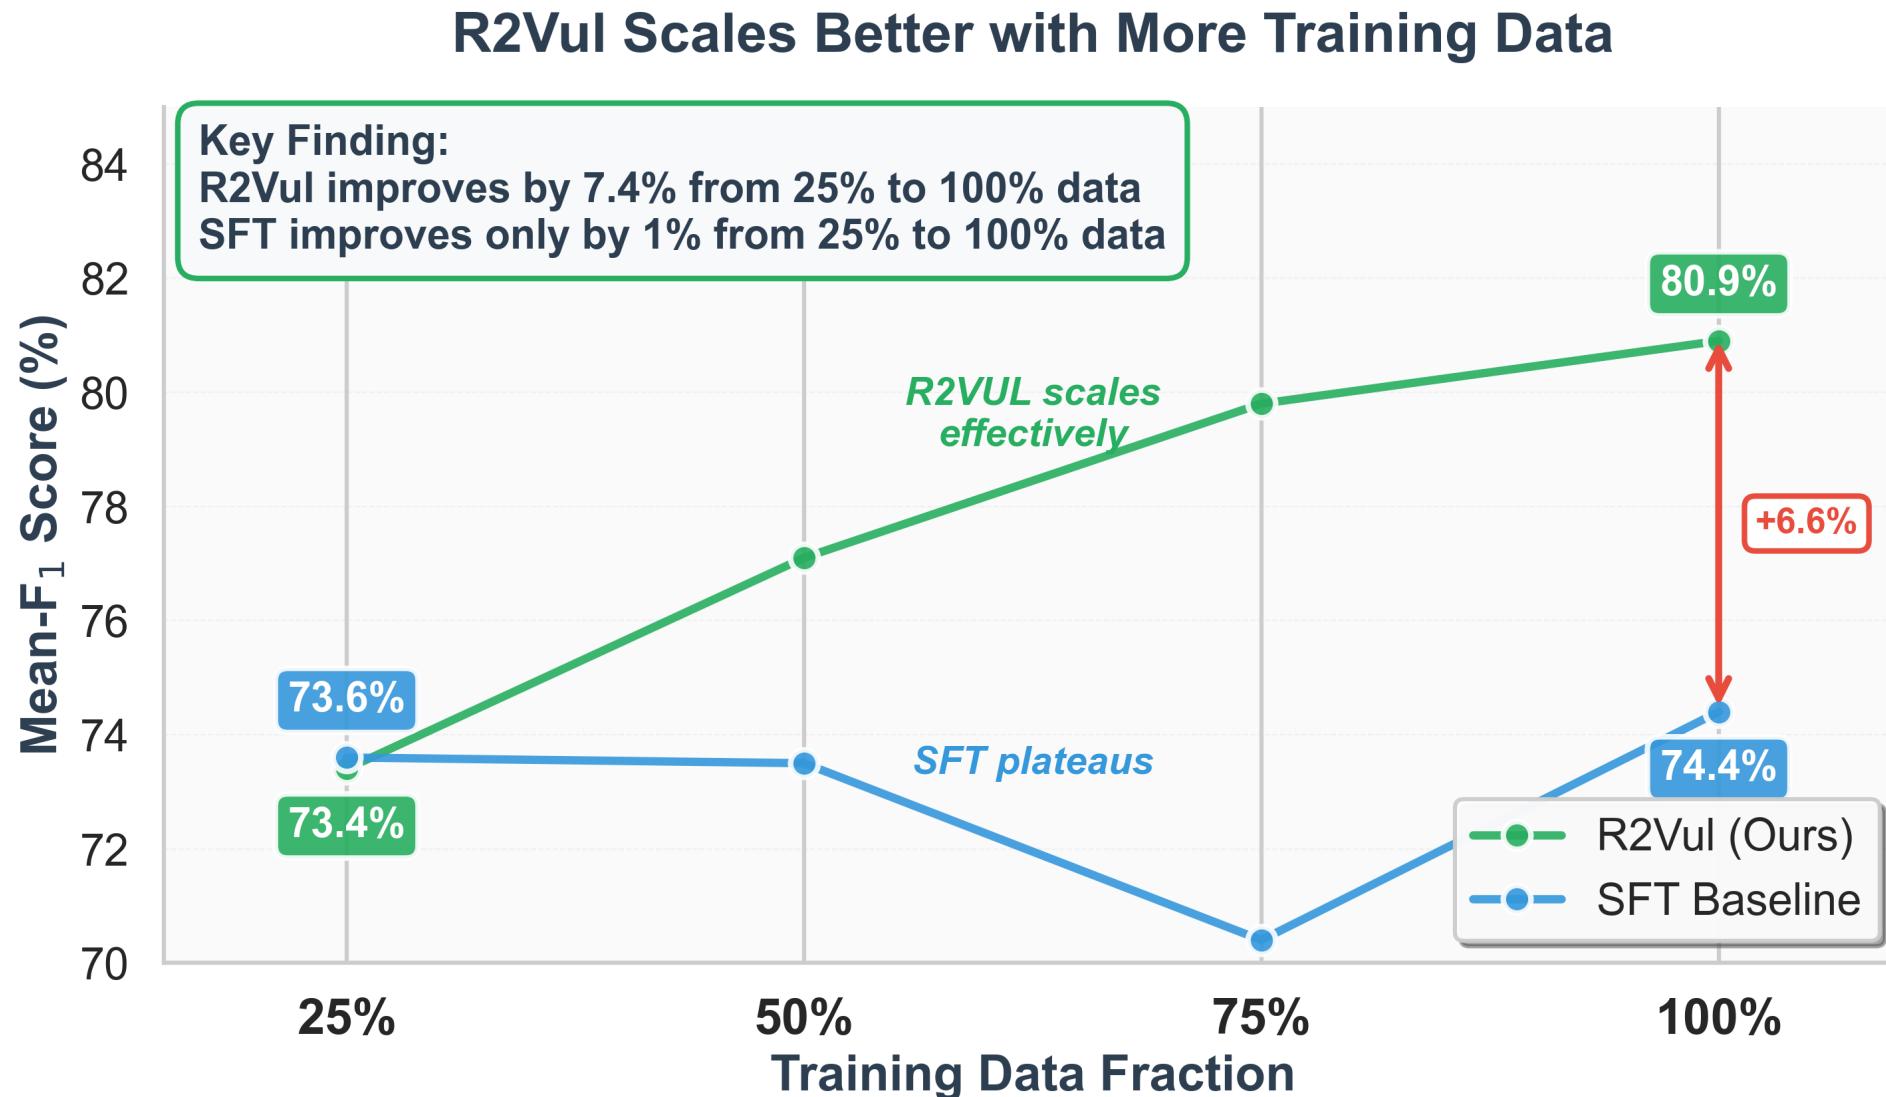


Main Results: Detection Performance



R2Vul-1.5B wins over commercial LLMs

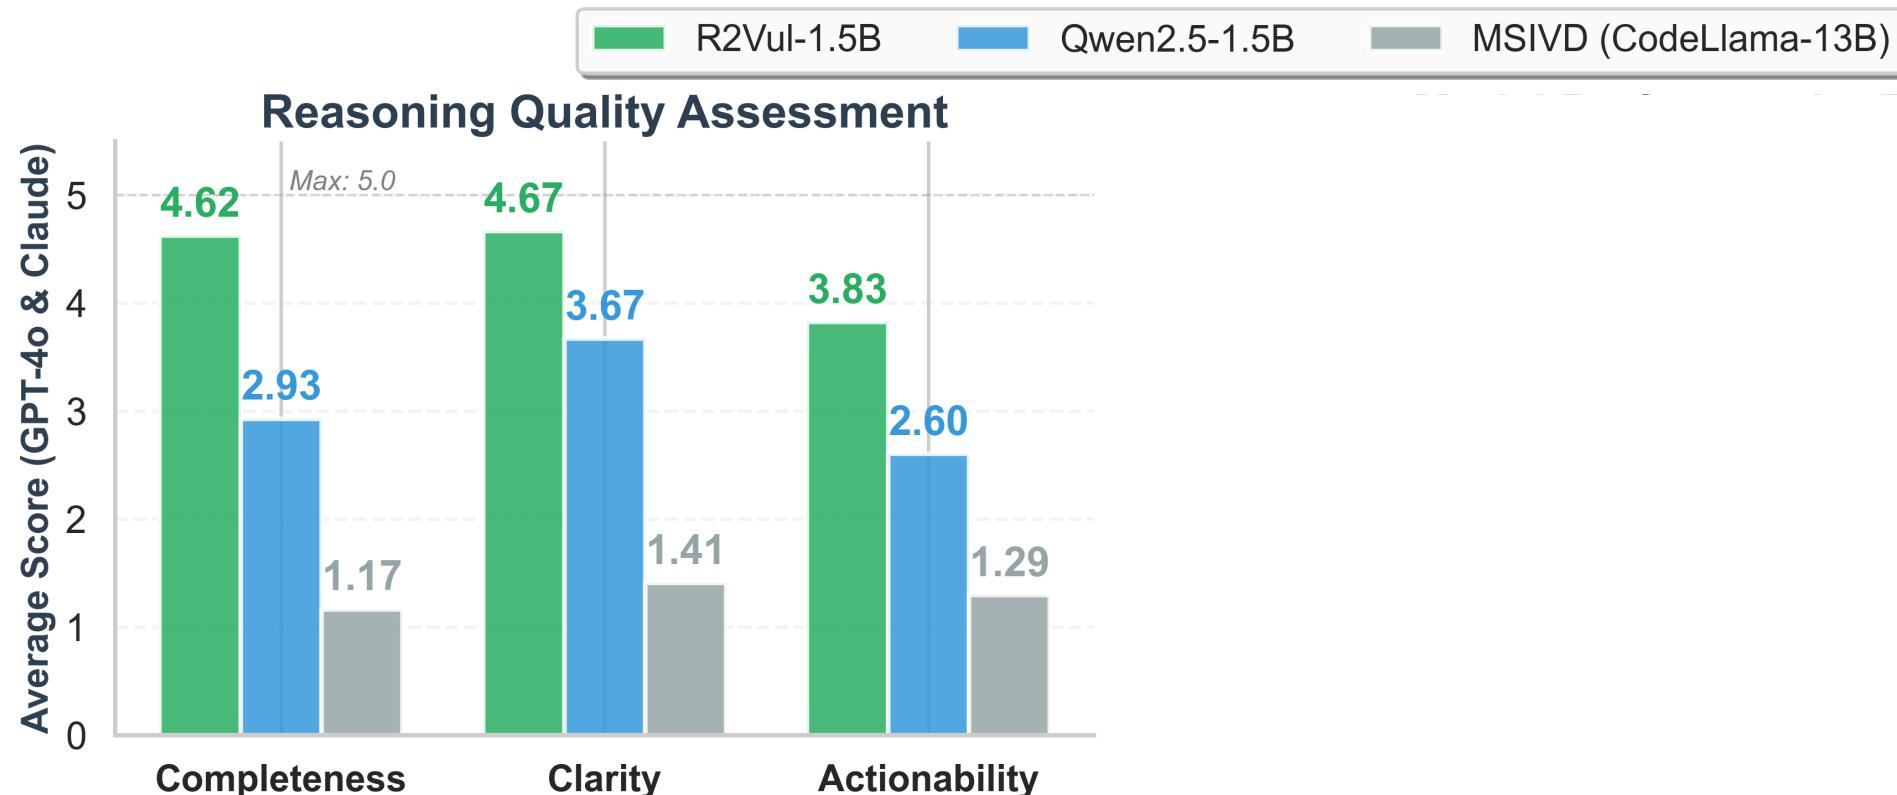
R2Vul vs. SFT: Impact of Contrastive Reasoning Trace Pairs

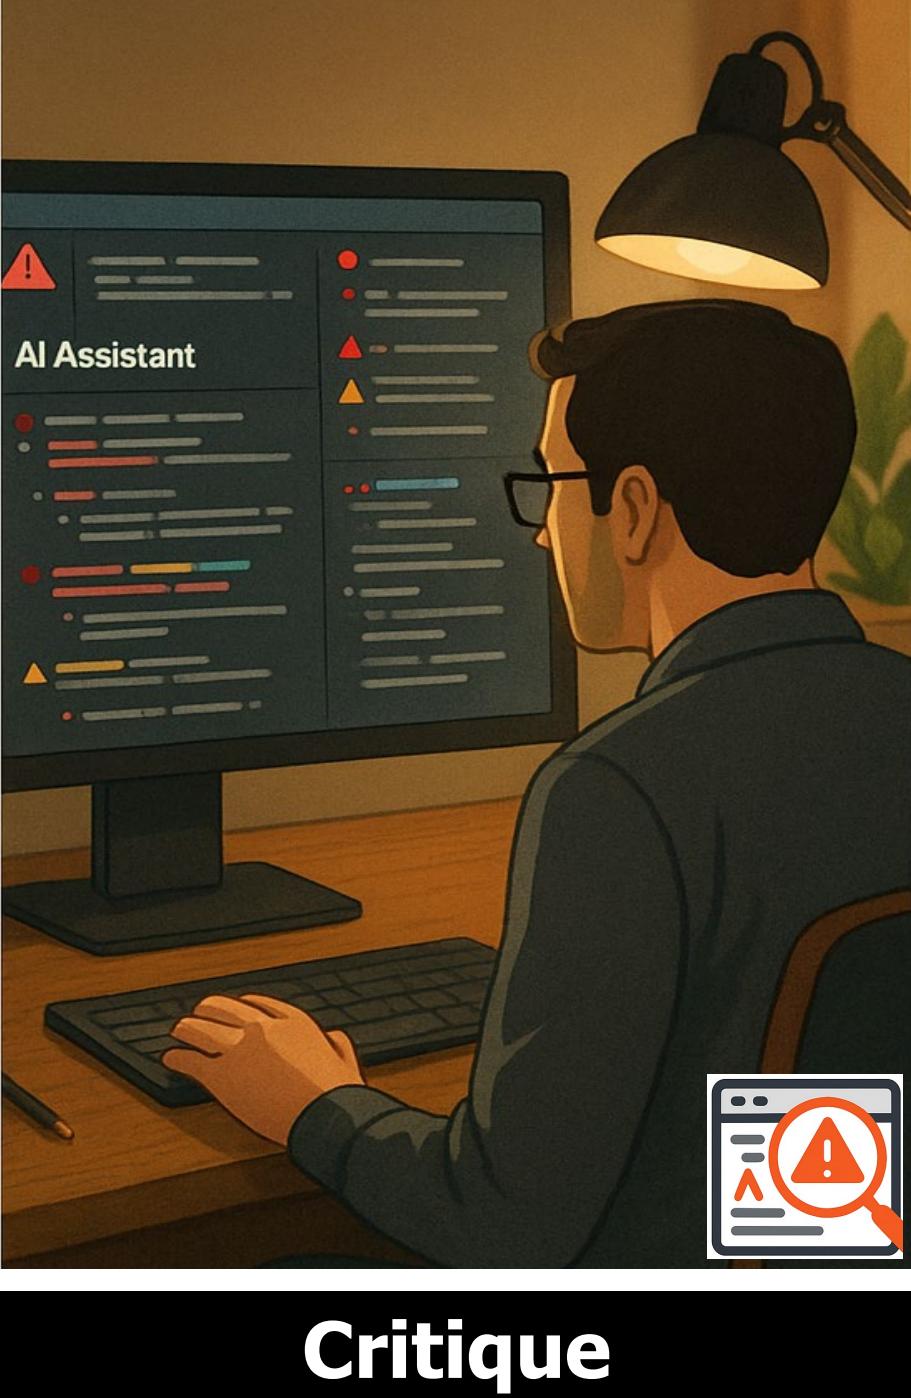


Reasoning Quality Assessment

Baselines: Qwen2.5-1.5B-Coder-Instruct and MSIVD (CodeLlama-13B): reasoning-based LLM [1]

Evaluators: GPT-4o, Claude-3.5-Sonnet, and two human experts



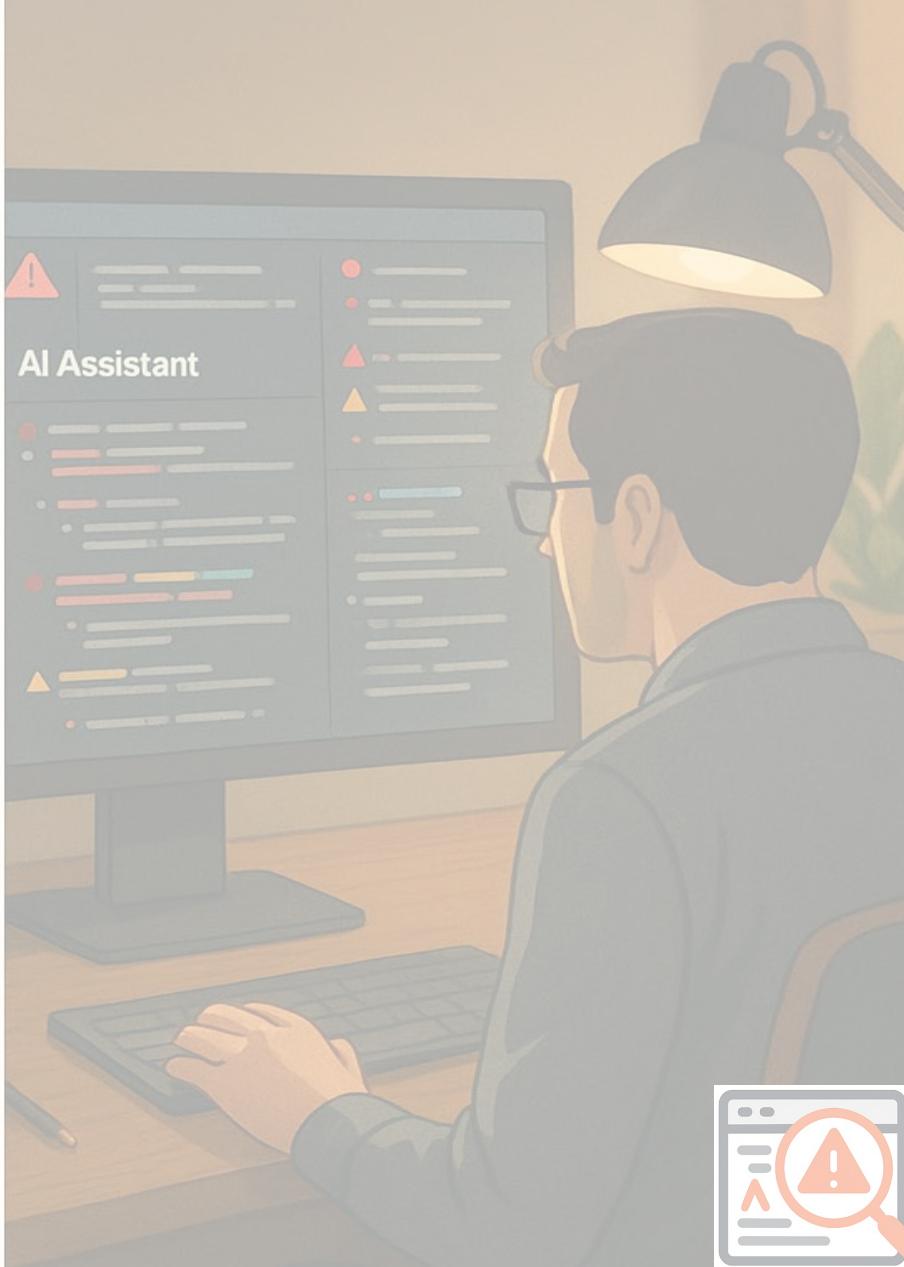
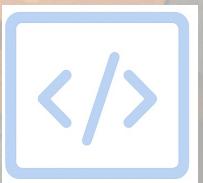


Critique

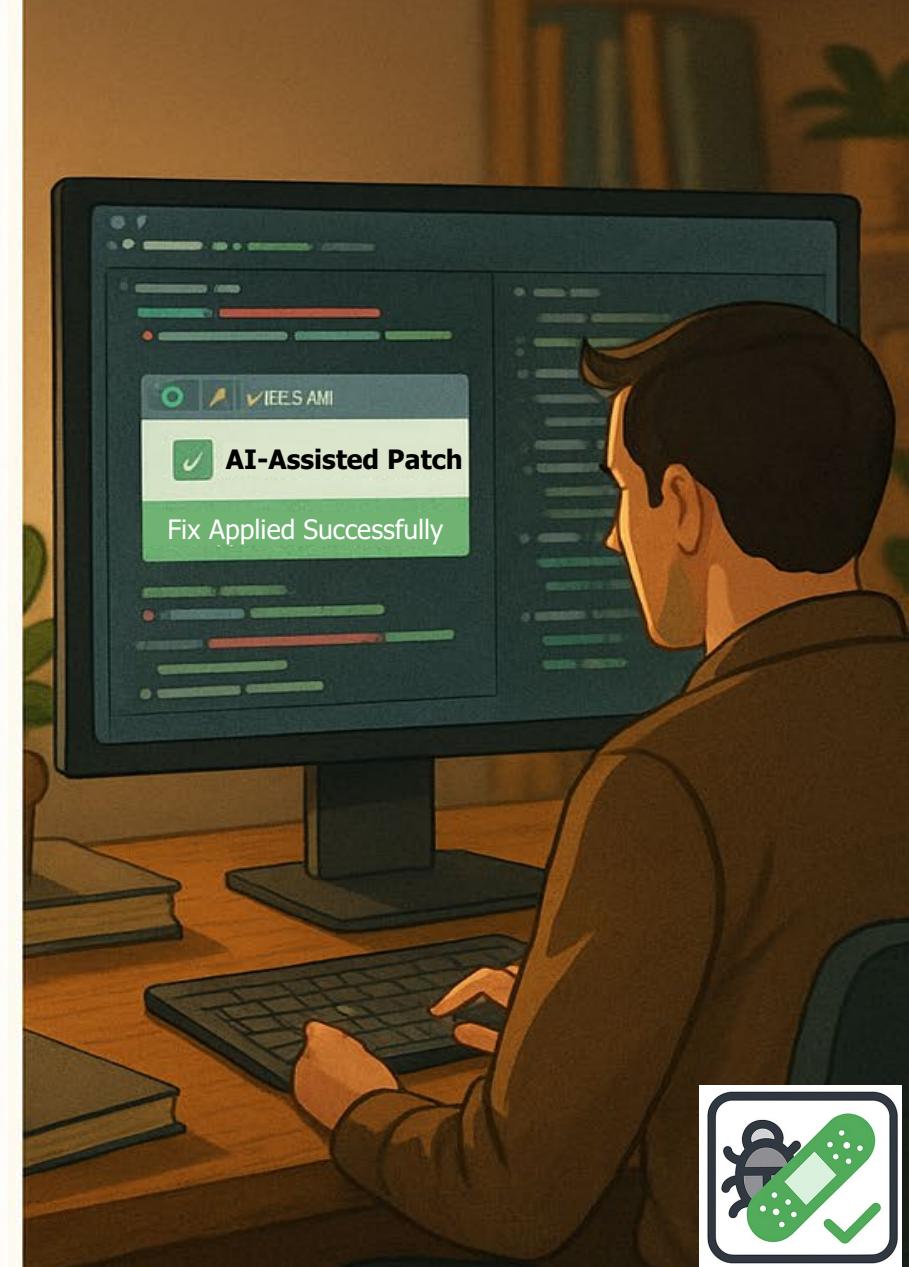
Engineering **contrastive reasoning trace pairs** help LLM reasoning for vulnerability detection



Code



Critique



Cure



Integrating Diverse Inputs from Diverse Sources Can Help AI

SANER 2016

History Driven Program Repair

Xuan-Bach D. Le, David Lo
School of Information Systems
Singapore Management University
{dxb.le.2013,davidlo}@smu.edu.sg

Claire Le Goues
School of Computer Science
Carnegie Mellon University
clegoues@cs.cmu.edu



First MSR-powered APR work: mines hundreds of GitHub repositories for repair patterns serving as augmented inputs to test cases for search-based automated program repair (APR); most-cited SANER 2016 paper that inspired Facebook

Integrating Diverse Inputs for LLM-Powered Vulnerability Repair

ICSE 2024

Out of Sight, Out of Mind: Better Automatic Vulnerability Repair by Broadening Input Ranges and Sources

Xin Zhou
Singapore Management University
Singapore
xinzhou.2020@phdcs.smu.edu.sg

Kisub Kim*
Singapore Management University
Singapore
kisubkim@smu.edu.sg

Bowen Xu
North Carolina State University
USA
bxu22@ncsu.edu

DongGyun Han
Royal Holloway, University of London
United Kingdom
donggyun.han@rhul.ac.uk

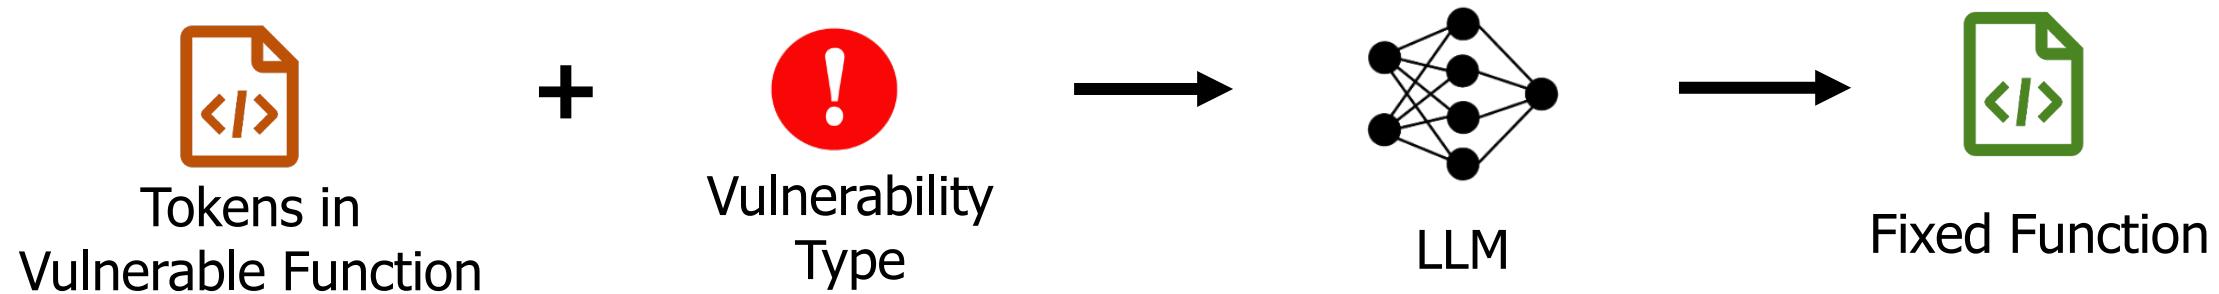
David Lo
Singapore Management University
Singapore
davidlo@smu.edu.sg



First LLM-powered automated program repair work that (i) links vulnerable code to LLM-enriched CWE knowledge and (ii) leverages AST contents to double the efficacy of prior work

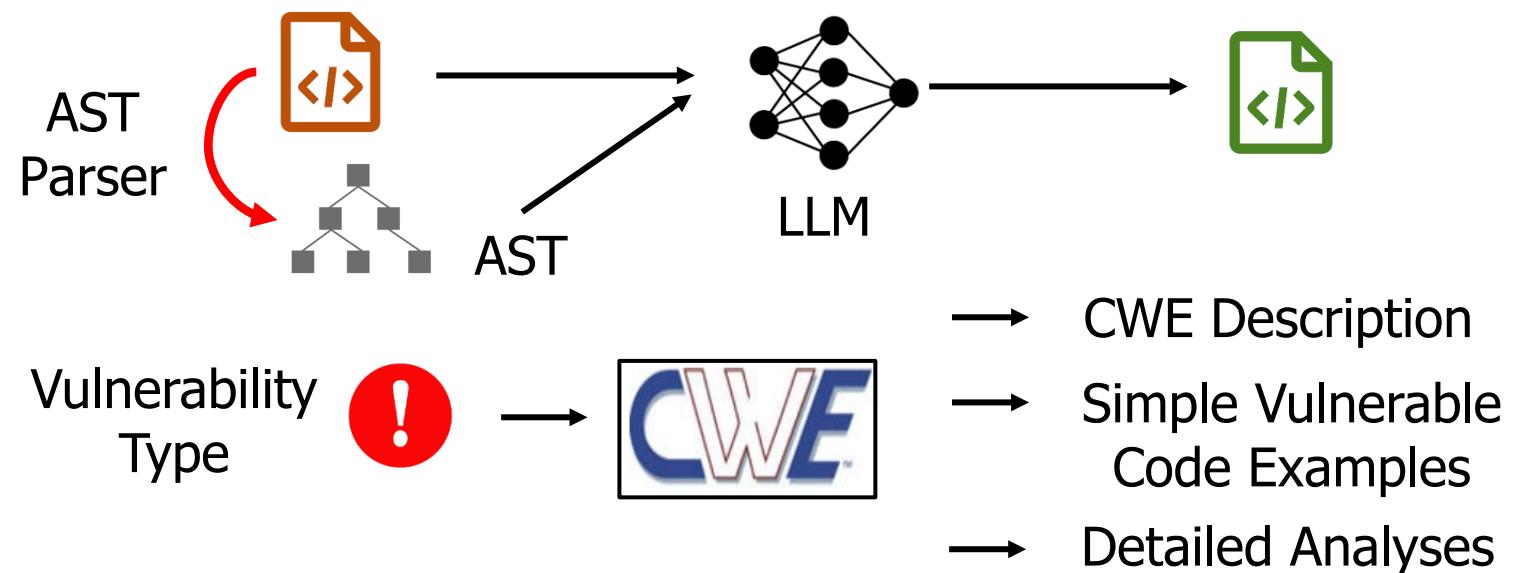
Integrating Diverse Inputs for LLM-Powered Vulnerability Repair

Previous solutions:



Many **other inputs** have not been leveraged:

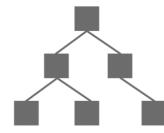
- *Abstract Syntax Tree*
- *CWE Knowledge*



VulMaster's Design

Data-Centric Innovations

Incorporate AST



Incorporate CWE knowledge



Address lengthy inputs

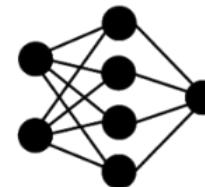


+

Multi-LLM Collaboration



GPT-3.5



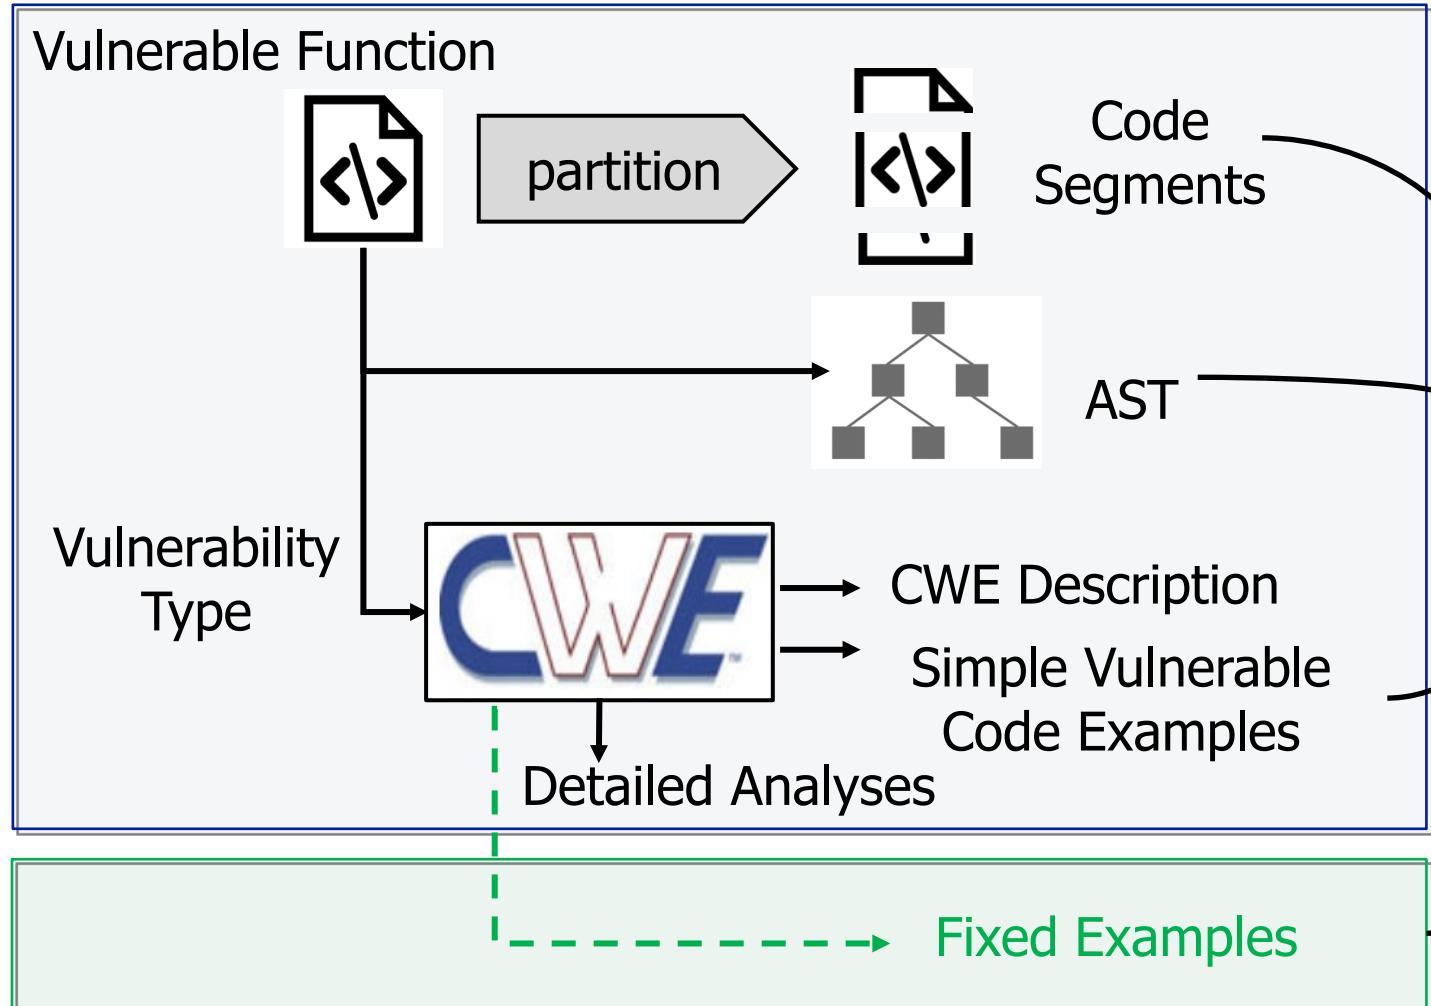
CodeT5

=

2x Fixed Vulnerabilities

Overall Framework

Step 1: Collect Diverse Inputs



Step 2: Fill in Missing Data with GPT-3.5

Step 3: Fuse Diverse Inputs

LLM (CodeT5) + Data Fusion



Fixed Function

Additional Model Innovations

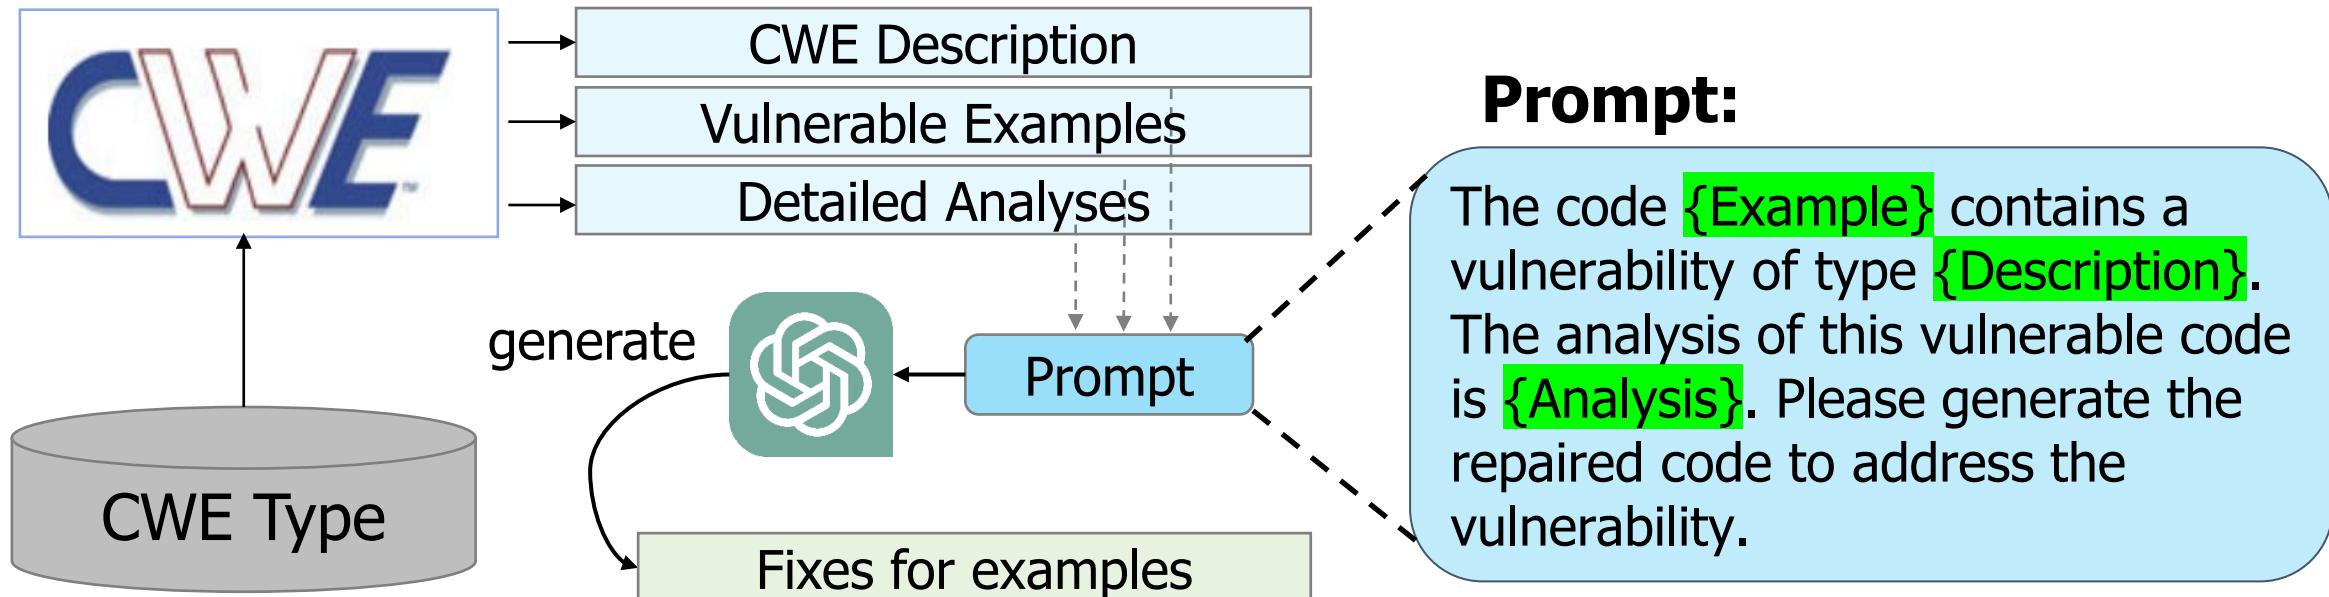
Model Adaptation for AST

Multi-Task Learning

CWE Tree Exploration

Step 2: Fill In Missing Data with GPT-3.5

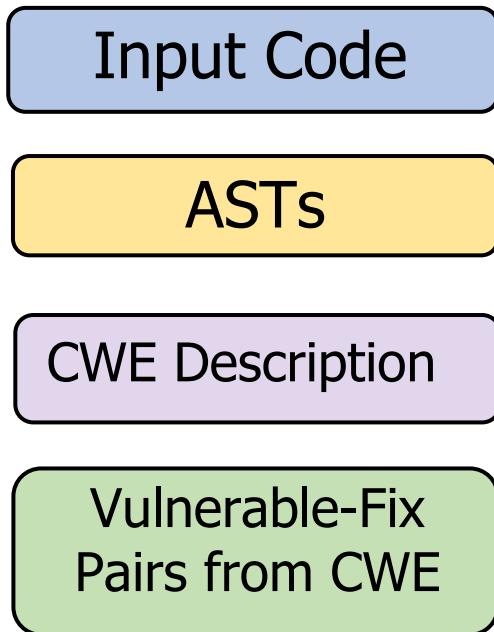
- General-purpose LLMs are effective for
 - Fixing simple (toy) vulnerable examples when *detailed analysis* is given.
- Thus, we use **GPT-3.5** to generate the fixed CWE examples:



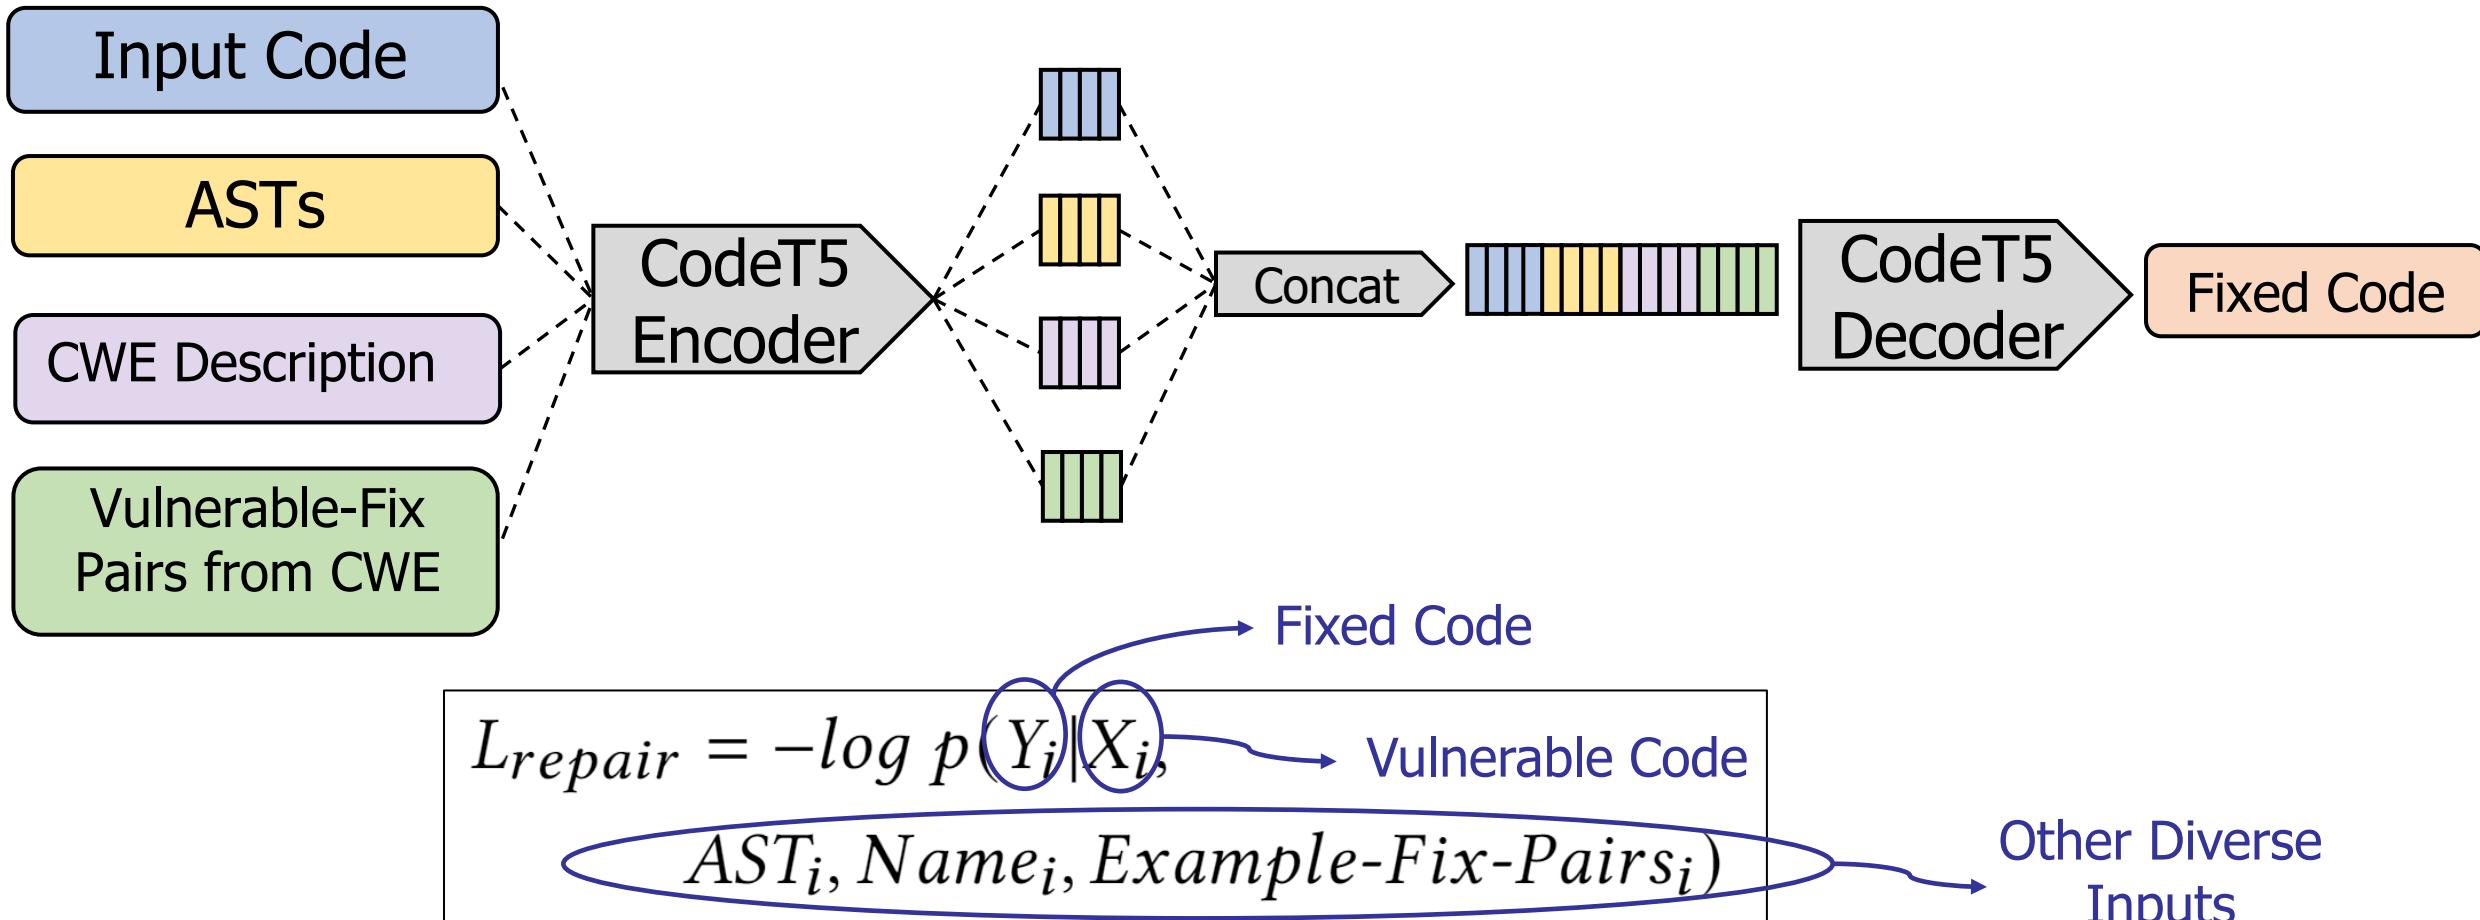
- **Note:** Real vulnerability fixes are not simple or come with detailed analysis

Step 3: Fusing Diverse Input Data Components

Diverse Inputs



Contextual Embedding and Concatenation



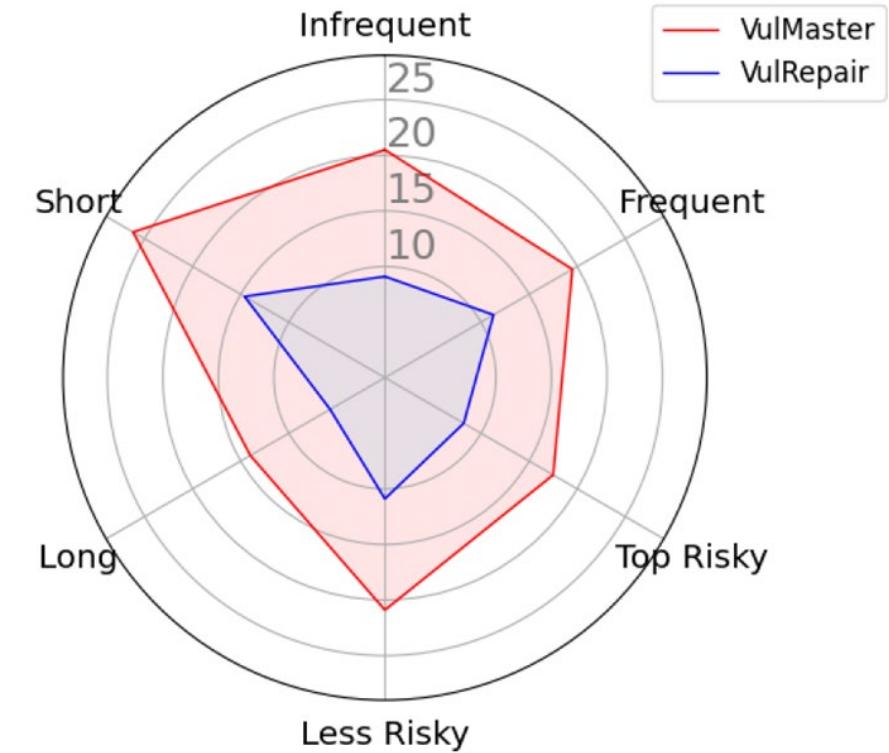
Repair Generation

Results: Comparisons with SOTA

Main Results

Type	Approach	EM	BLEU
LLM	GPT-3.5 [55]	3.6	8.8
	GPT-4 [56]	5.3	9.7
task-specific	VRepair [9]	8.9	11.3
	VulRepair [19] (SOTA)	10.2	21.3
Ours	VulMaster	20.0	29.3

- VulMaster **doubles the Exact Match (EM) score**
- VulMaster consistently outperforms for vulnerabilities of different characteristics

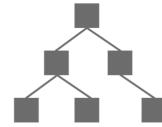


- *long/short*: the length of the code
- *frequent/infrequent*: the vulnerability type frequencies
- *top/less risky*: top 10 most dangerous CWEs or not

Latest Extension

More Diverse Data +

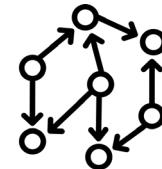
Incorporate AST



Incorporate
CWE knowledge



Incorporate
Data Flow



Reasoning
steps for repair



GPT-4o

**Reinforcement
Learning**

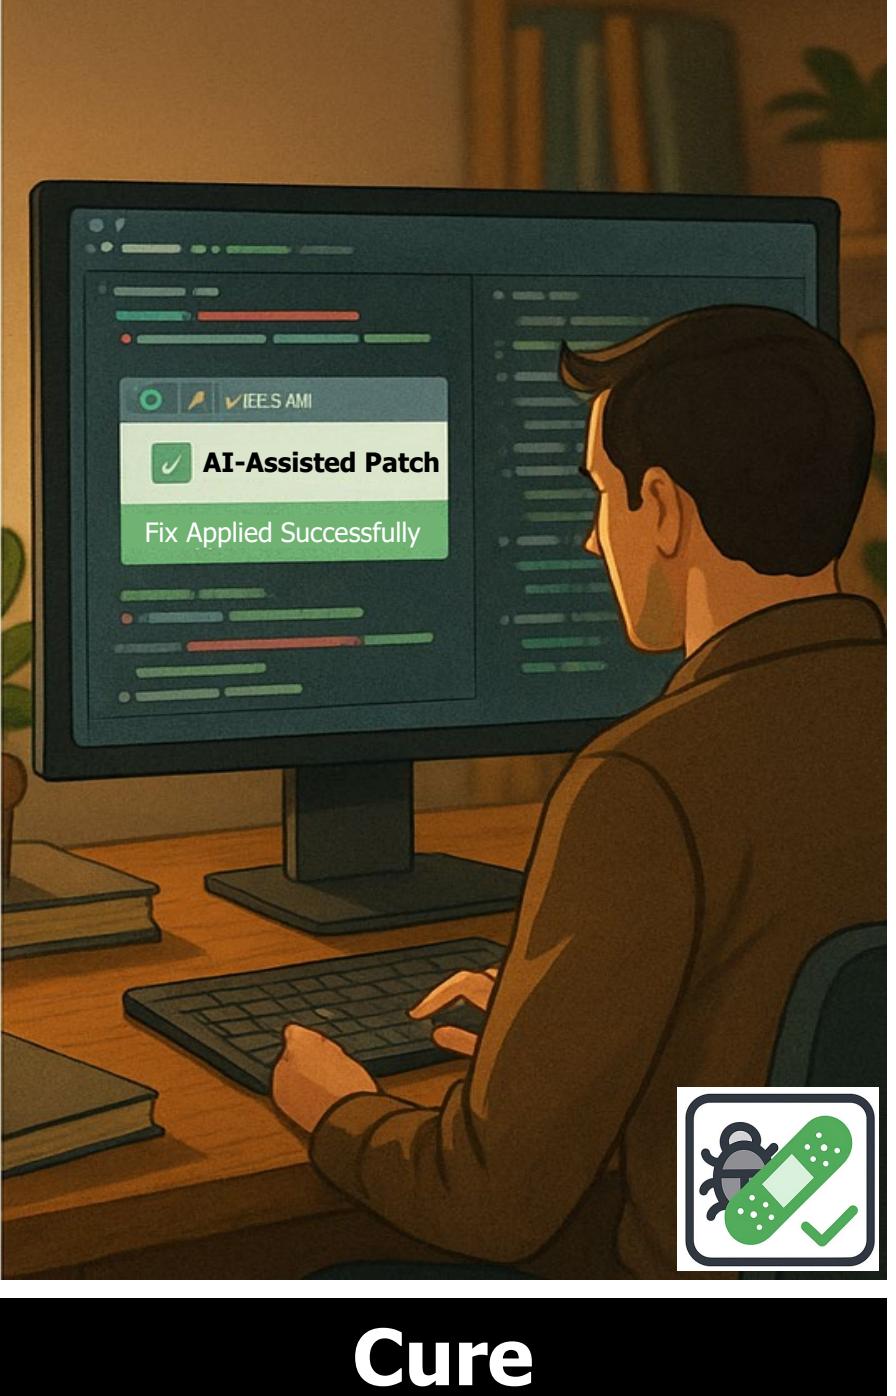


Qwen-2.5 7B Instruct

=

**+34.96%
Performance Gain**

Please stay tuned: Paper will be released on arXiv soon

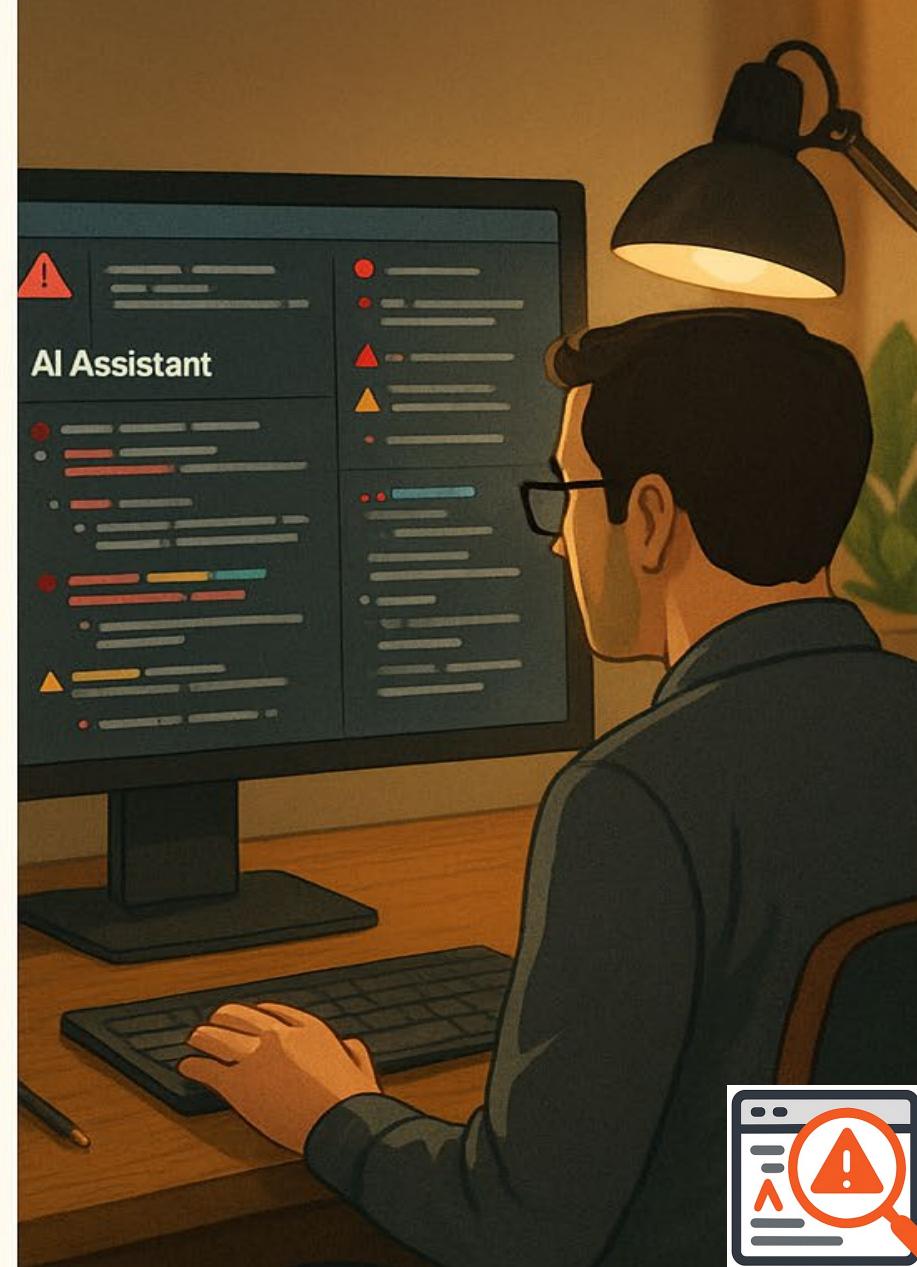


Cure

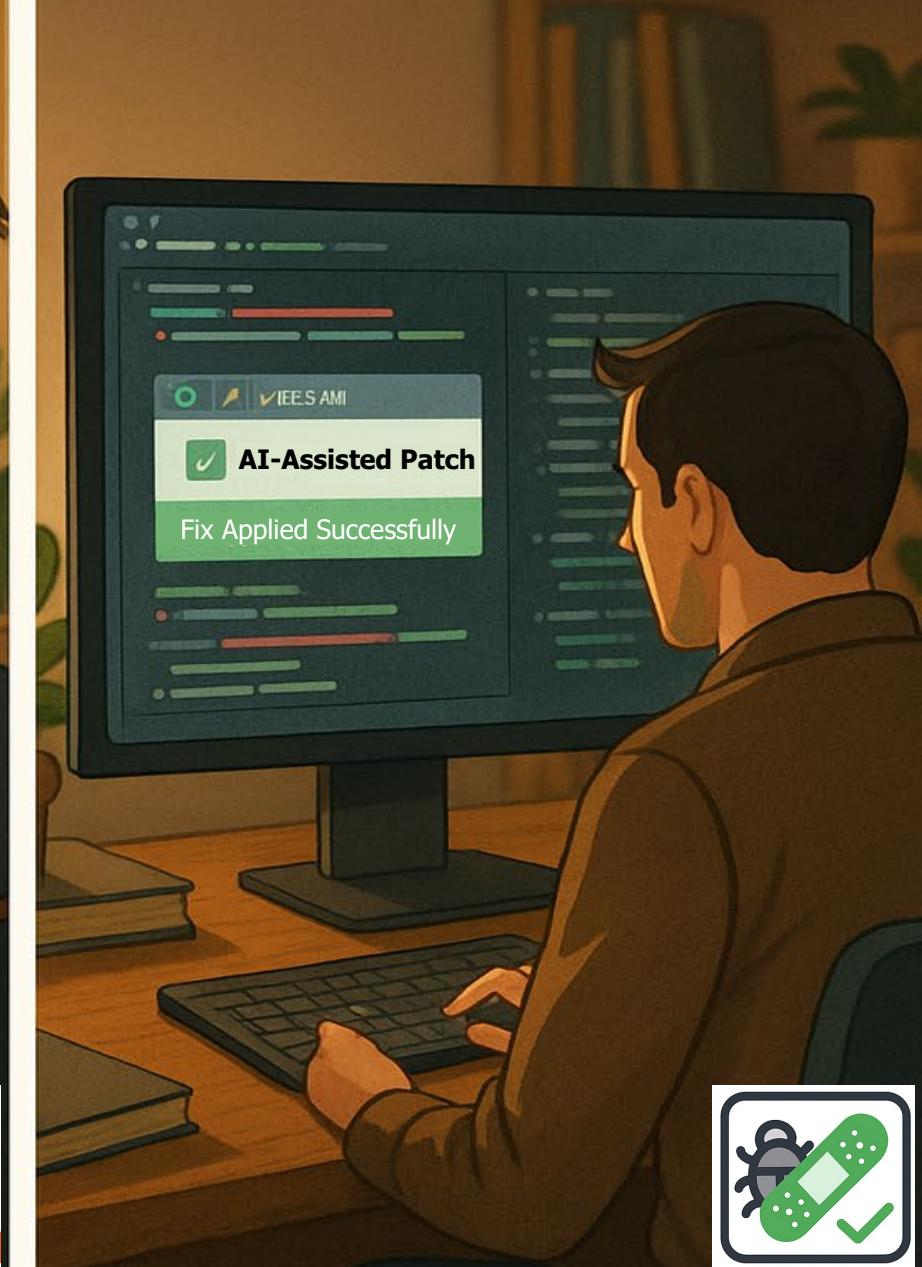
Engineering **holistic** representation
from **diverse inputs and sources**
help LLM reasoning for repair



Code



Critique

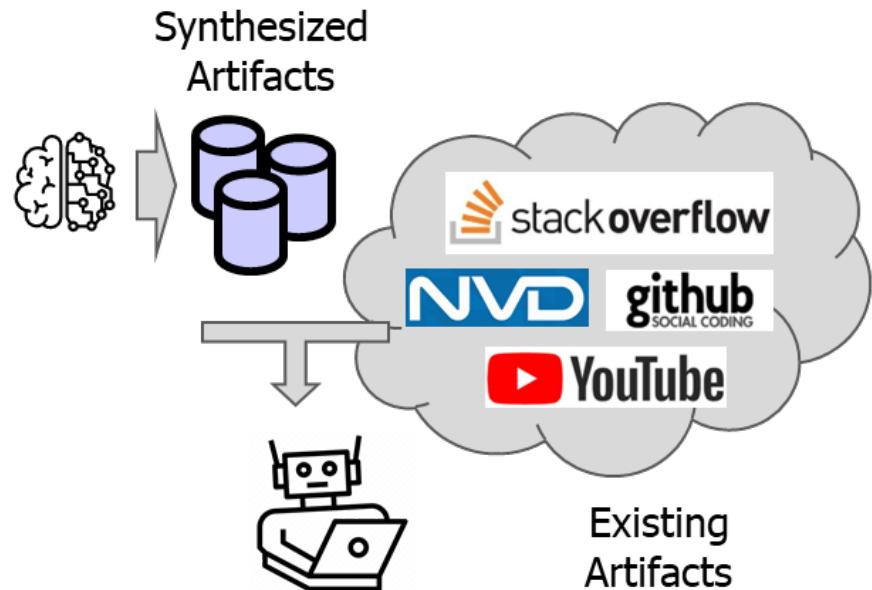


Cure

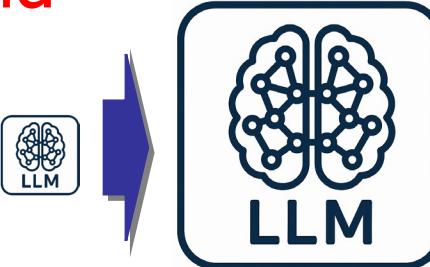


Data-Centric LLM4SE: Summary

*“How Can We
Systematically Engineer
Software Artifacts to Build
Better LLM4SE Bots?”*



Beyond



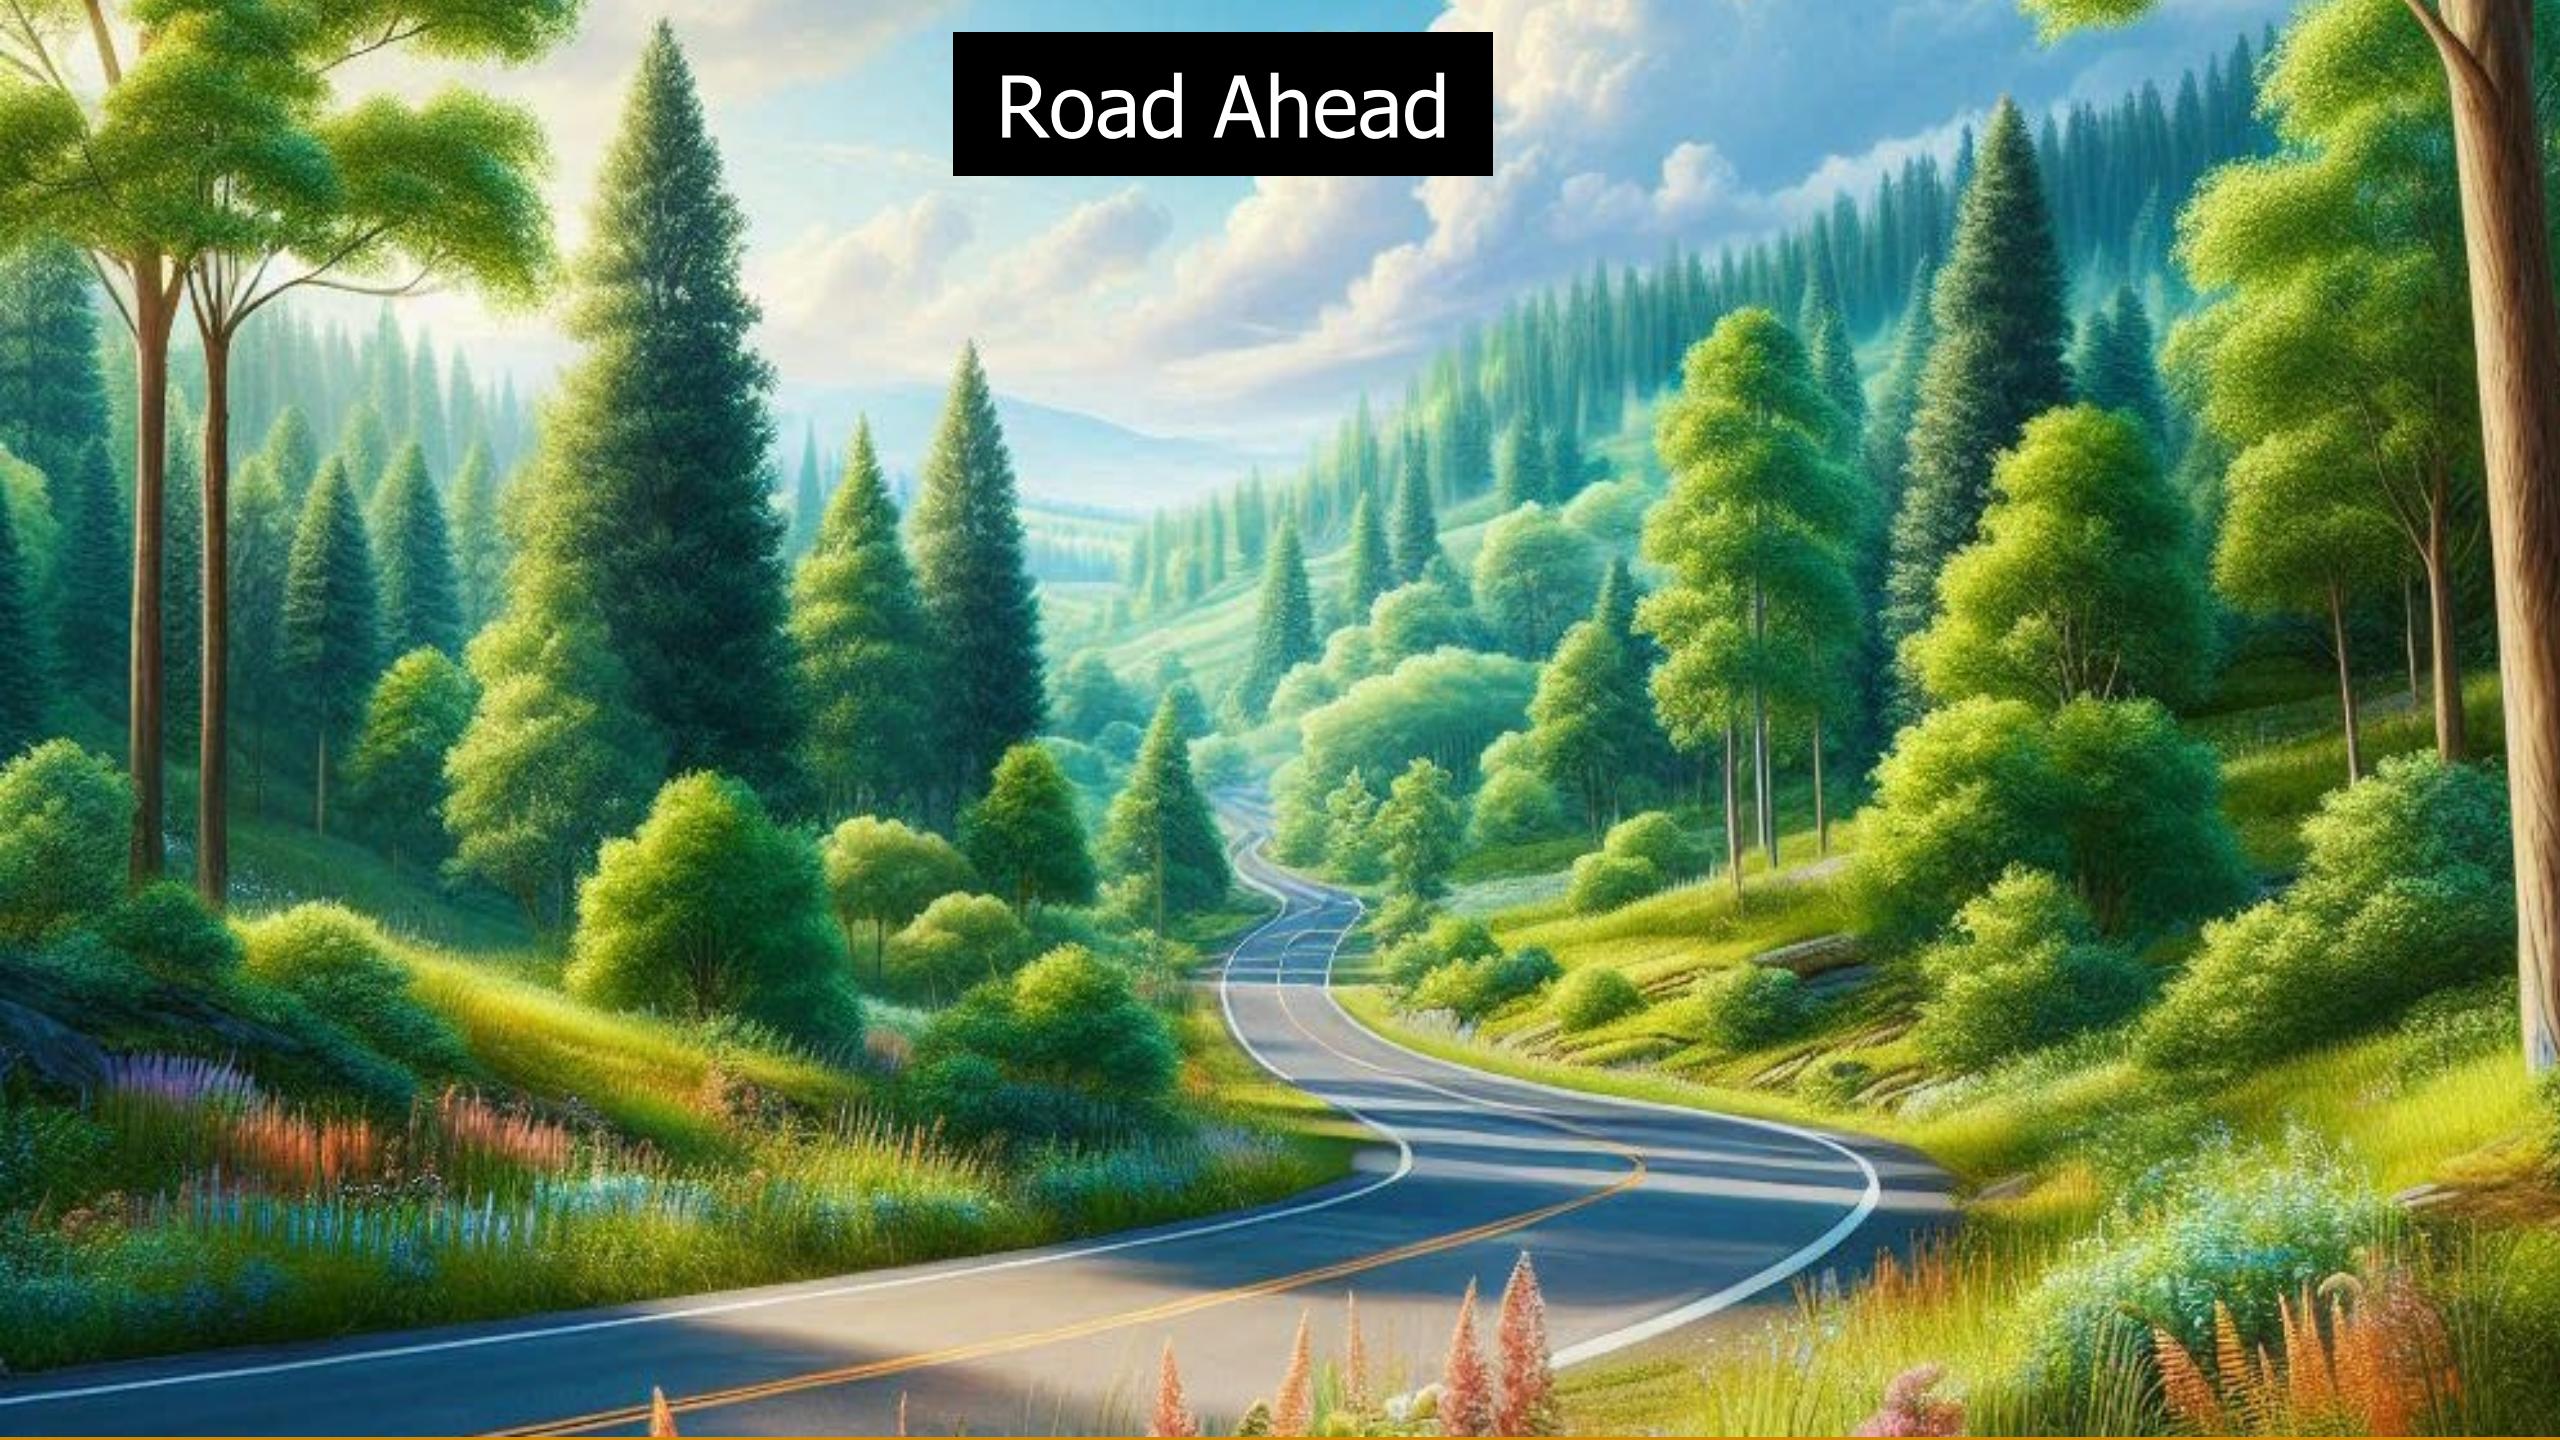
*Prompt
Engineering*

- We want to **elevate the data**:
 - enriching it, restructuring it, contrasting it, and linking it in ways that *promote deeper reasoning*
- And **utilizing them effectively** to train LLMs to *reason and do better for specific SE tasks*

Data-Centric LLM4SE: Summary



“Because in software engineering,
how we **craft the data**
is how we **shape the intelligence**”

A wide-angle photograph of a winding asphalt road through a dense forest. The road curves from the bottom left towards the center of the frame. The surrounding landscape is filled with tall, green coniferous trees and smaller deciduous trees. The terrain is hilly and covered in lush greenery. The sky above is bright and blue with scattered white clouds.

Road Ahead

Road Ahead



Multi-Agent LLM4SE

Real-world problems demand "synergistic collaboration". Multi-agent LLM4SE transforms isolated agents into a coordinated group of experts.

LLM-Based Multi-Agent Systems for Software Engineering: Literature Review, Vision and the Road Ahead

JUNDA HE, Singapore Management University, Singapore

CHRISTOPH TREUDE, Singapore Management University, Singapore

DAVID LO, Singapore Management University, Singapore

TOSEM 2025



Multi-Agent LLM4SE: Vulnerability Detection

Let the Trial Begin: A Mock-Court Approach to Vulnerability Detection using LLM-Based Agents

Ratnadira Widyasari

ratnadiraw@smu.edu.sg

Singapore Management University
Singapore

Han Wei Ang

ang_han_wei@tech.gov.sg

GovTech
Singapore

Lwin Khin Shar

lkshar@smu.edu.sg

Singapore Management University
Singapore

Martin Weyssow

mweyssow@smu.edu.sg

Singapore Management University
Singapore

Frank Liauw

frank_liauw@tech.gov.sg

GovTech
Singapore

Hong Jin Kang

hongjin.kang@sydney.edu.au

University of Sydney
Australia

Ivana Clairine Irsan

ivanairsan@smu.edu.sg

Singapore Management University
Singapore

Eng Lieh Ouh

elouh@smu.edu.sg

Singapore Management University
Singapore

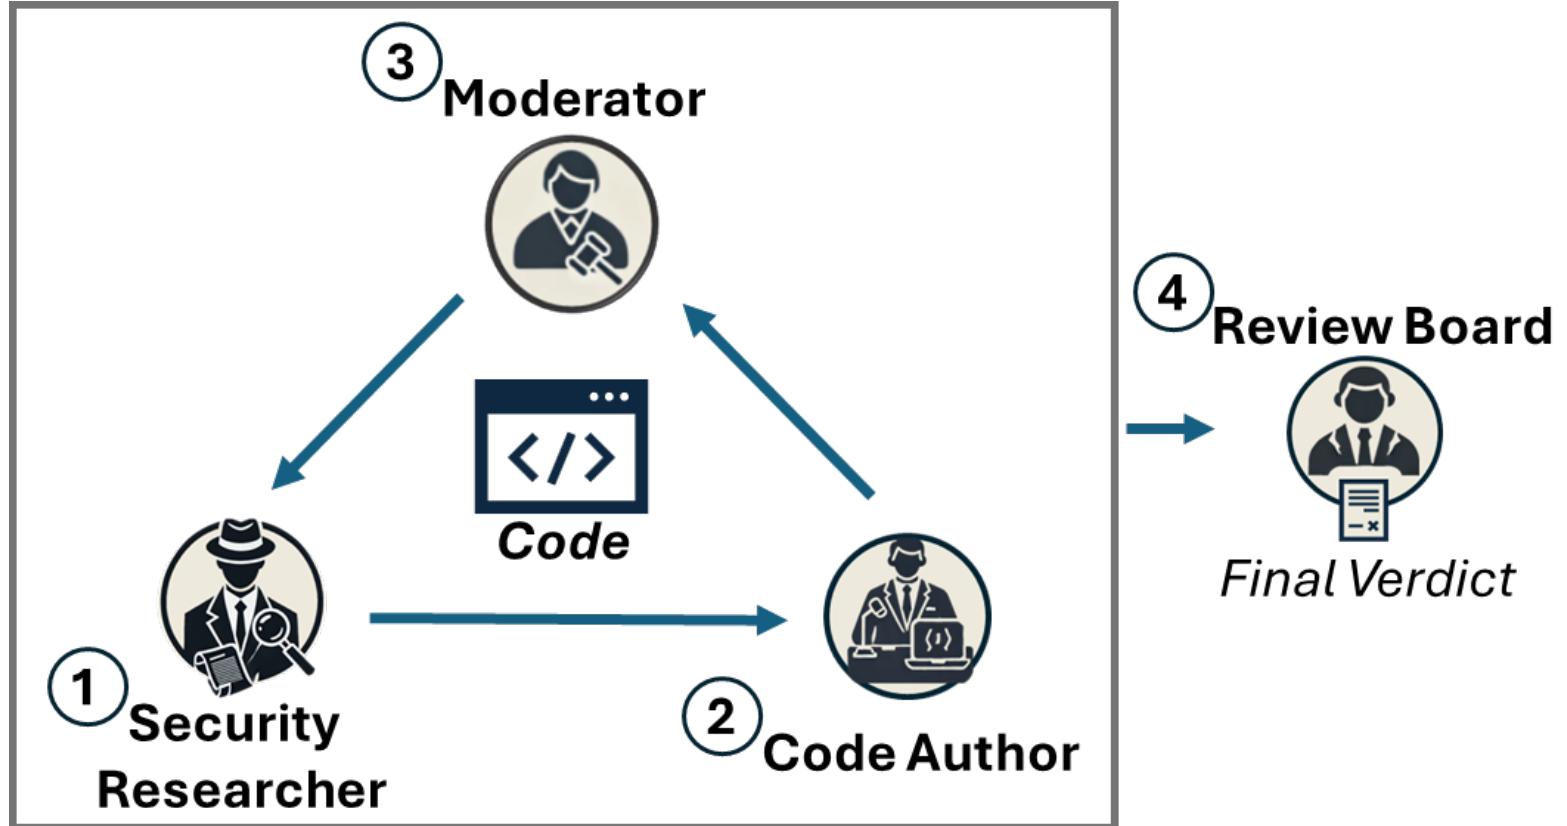
David Lo

davidlo@smu.edu.sg

Singapore Management University
Singapore



Multi-Agent LLM4SE: VulTrial



VulTrial's Design

```
TfLiteStatus EvalGatherNd(TfLiteContext* context,
    const TfLiteTensor* params,
    const TfLiteTensor* comment indices,
    TfLiteTensor* output) {
    // Vulnerability: There's no check to ensure that 'indices' are
    // within valid bounds for 'params'. Negative or out-of-range
    // indices lead to out-of-bounds reads.
    switch (params->type) {
        case kTfLiteFloat32:
            return GatherNd<float, IndicesT>(params, indices, output);
    // ... other types ...
        default:
            context->ReportError(
                context, "Params type '%s' not supported by gather_nd.", ,
                TfLiteTypeGetName(params->type));
            return kTfLiteError;
    }
}
```

Example of Vulnerable Code *CWE-125 (Out-of-bounds Read)*

	Security Researcher
vulnerability	: Out-of-Bounds Read
reason	: The function 'EvalGatherNd' does not perform any bounds checking on the 'indices' tensor. [...].
impact	: Exploiting this vulnerability could lead to information disclosure, [...] could be leveraged for further attacks such as Denial of Service (DoS).
vulnerability	: Type Handling
reason	: [...] if the new type is not properly validated [...], it could lead to unexpected behavior or crashes.
impact	: [...] could be exploited by an attacker to cause a DoS

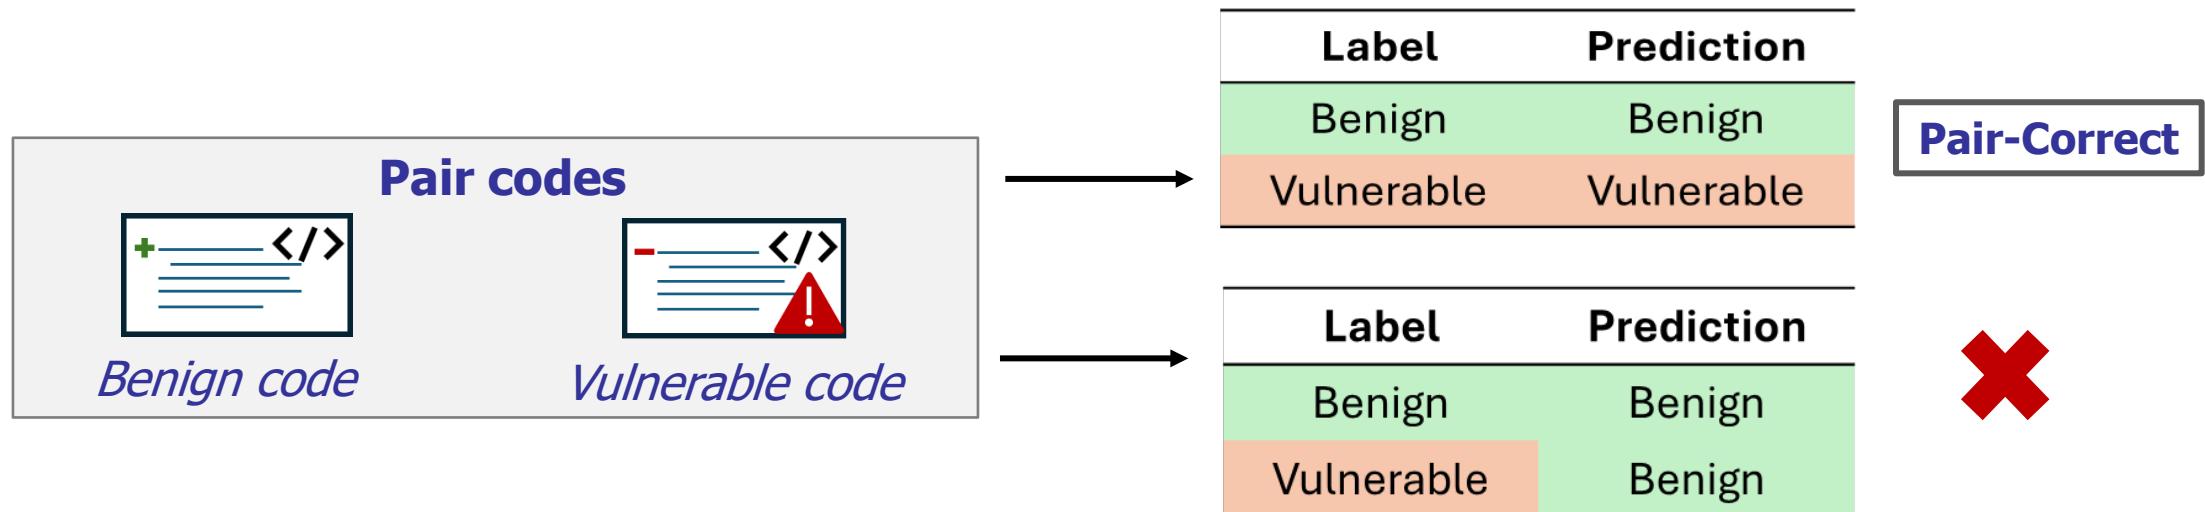
	Code Author
vulnerability	: Out-of-Bounds Read
reason	: The concern regarding out-of-bounds reads is valid. [...] I propose implementing a bounds checking [...].
vulnerability	: Type Handling
reason	: [...] any unsupported type will be caught in the default case, which reports an error. This is a standard practice [...].

	Moderator
researcher summary	: [...] could result in information disclosure [...]
author summary	: The author acknowledges the concern [...].

Review Board		
vulnerability	: Out-of-Bounds Read	decision : valid
severity	: high	action : fix immediately
reason	: [...] function does not perform bounds checking [...].	
vulnerability	: Type Handling	decision : partially valid
severity	: medium	action : monitor
reason	: [...] potential for new tensor types w/o proper handling [...].	

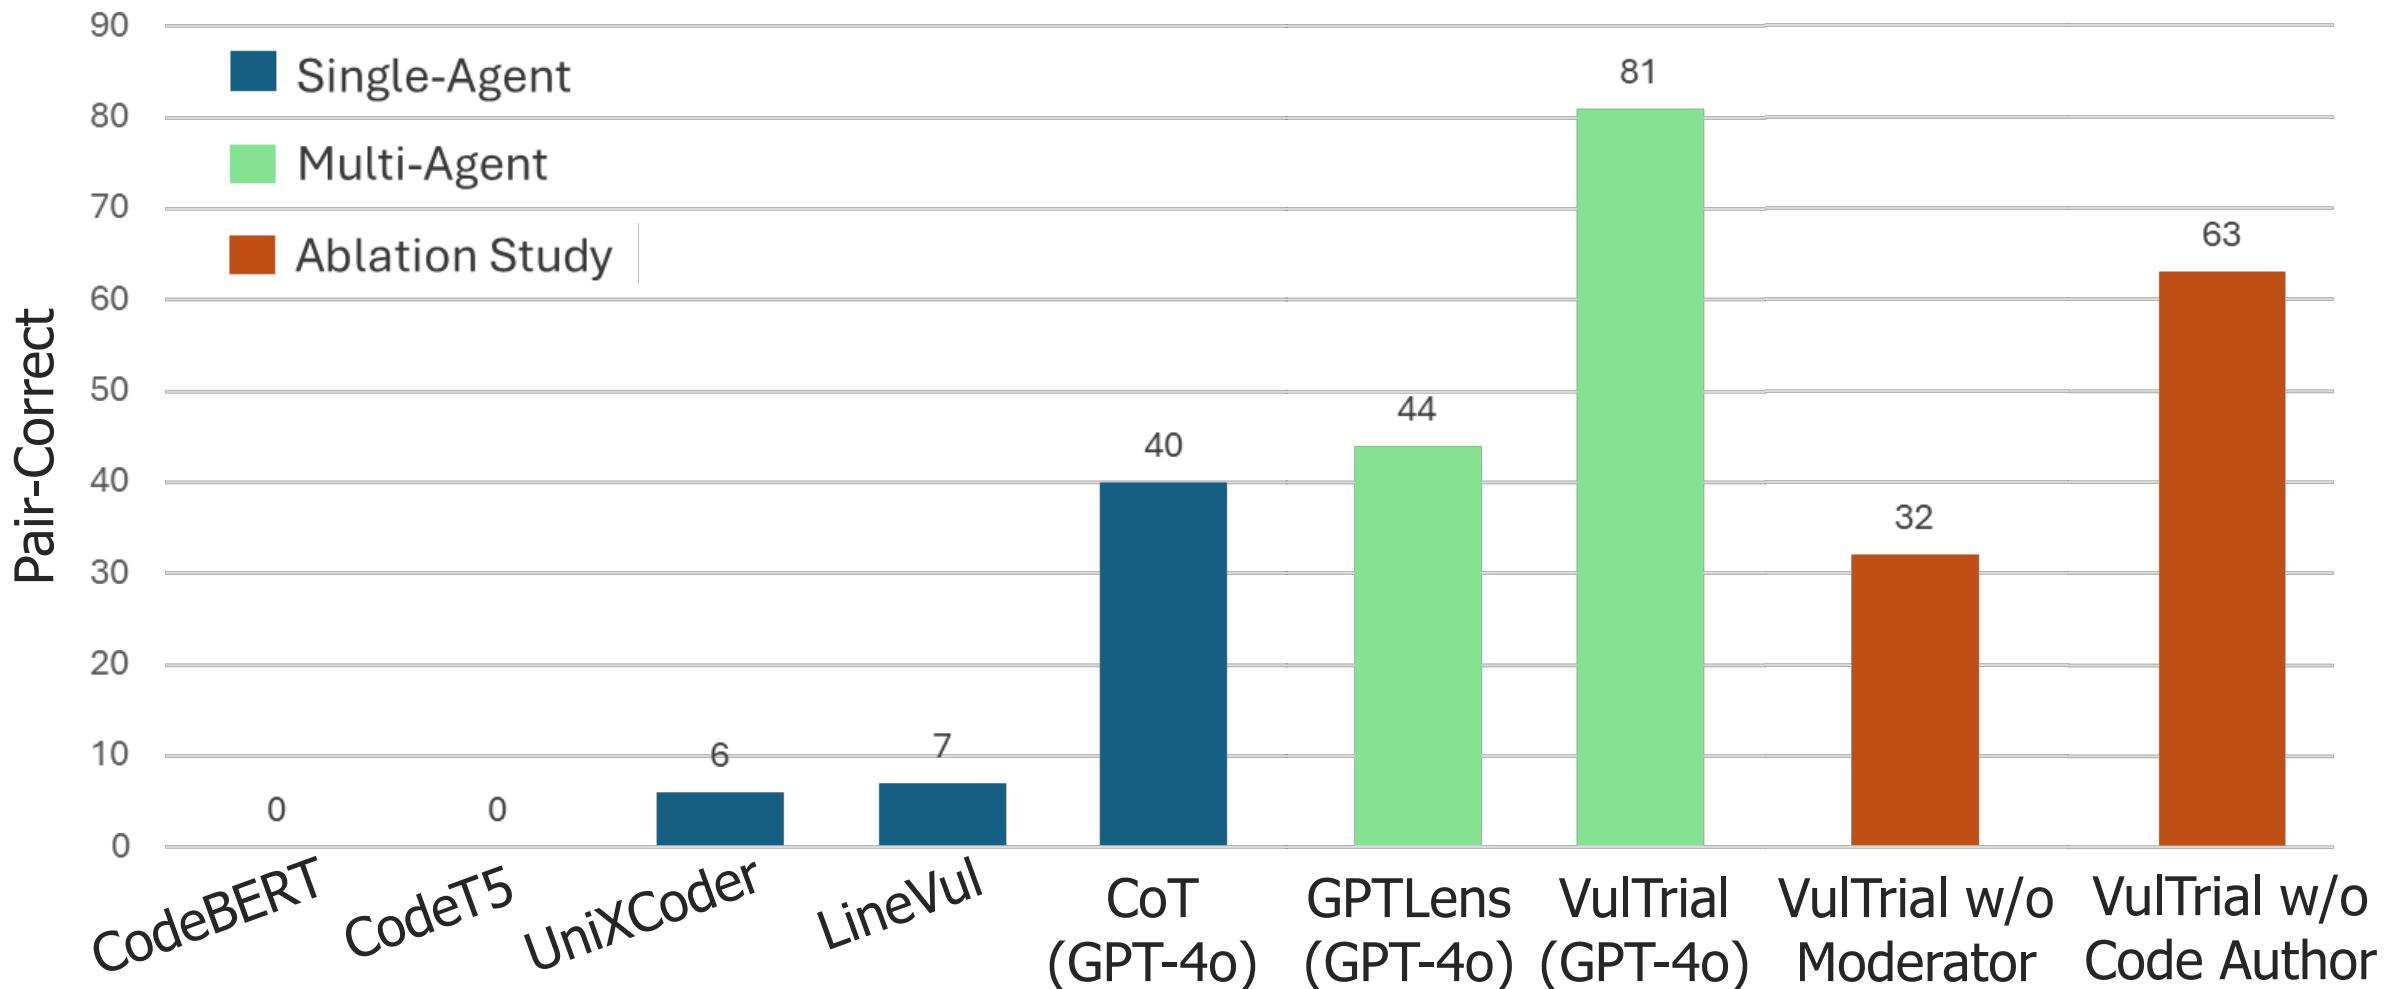
Evaluation Settings

- Evaluation on PrimeVul Pair [1] data.



[1] Yangruibo Ding, Yanjun Fu, Omniyyah Ibrahim, Chawin Sitawarin, Xinyun Chen, Basel Alomair, David A. Wagner, Baishakhi Ray, Yizheng Chen:
Vulnerability Detection with Code Language Models: How Far are We? ICSE 2025: 1729-1741

Evaluation Results



Multi-agent methods **outperform** the single-agent methods

Each agent in VulTrial is **necessary** to achieve the best performance

VulTrial in the Wild

User Interaction	Privileges Required	Attack Vector	Attack Complexity
Required	Low	Network	Low

CVE-2025-49000

Disclosure Date: June 03, 2025 · (Last updated August 14, 2025) ▾

CVE-2025-49000 CVSS v3 Base Score: 3.5

Easy to weaponize Authenticated Requires user interaction

Description

InvenTree is an Open Source Inventory Management System. Prior to version 0.17.13, the skip field in the built-in `label-sheet` plugin lacks an upper bound, so a large value forces the server to allocate an enormous Python `list`. This lets any authenticated label-printing user trigger a denial-of-service via memory exhaustion. The issue is fixed in versions 0.17.13 and higher. No workaround is available aside from upgrading to the patched version.

Related Paper @ ICSME 2025

SAEL: Leveraging Large Language Models with Adaptive Mixture-of-Experts for Smart Contract Vulnerability Detection

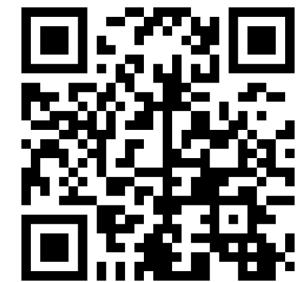
Lei Yu^{†‡1}, Shiqi Cheng^{†1}, Zhirong Huang^{†‡}, Jingyuan Zhang^{†‡}, Chenjie Shen^{†‡},
Junyi Lu^{†‡}, Li Yang^{†*}, Fengjun Zhang^{†§*}, Jiajia Ma[†]

[†]Institute of Software, Chinese Academy of Sciences, Beijing, China

[‡]University of Chinese Academy of Sciences, Beijing, China

[§]State Key Laboratory of Computer Science, Institute of Software, Chinese Academy of Sciences, Beijing, China

{yulei2022, chengshiqi, huangzhirong2022, zhangjingyuan2023, lujunyi2022}@iscas.ac.cn,
shenchenjie22@mails.ucas.ac.cn, {yangli2017, fengjun, majiajia}@iscas.ac.cn



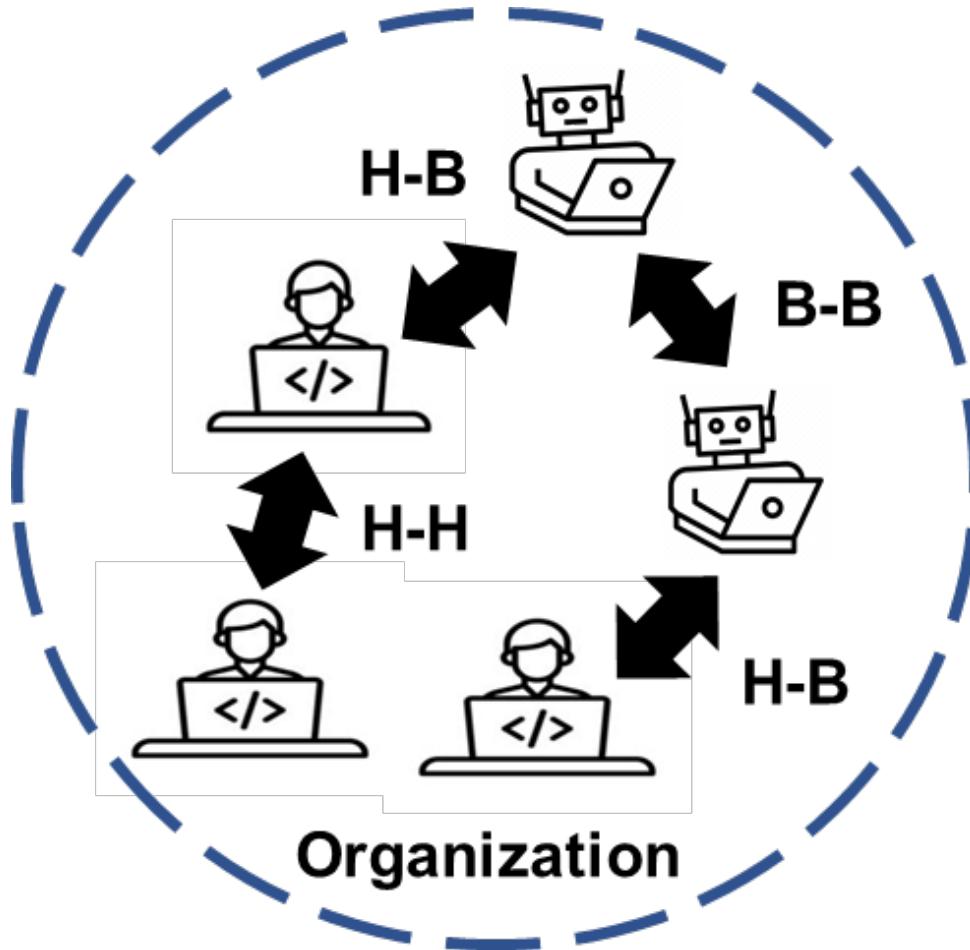
Session 12 - Security 1

Thu 11 Sep 2025 15:45 - 16:00

Case Room 260-057

Towards Full Multi-Agent Solution: Humans (H) and Bots (B)

N:M??



Road Ahead



Greening LLM4SE

LLM4SE solutions are large, slow, and not green, introducing much cost, latency, and carbon footprint

TOSEM 2025

Efficient and Green Large Language Models for Software Engineering: Vision and the Road Ahead

JIEKE SHI, ZHOU YANG, and DAVID LO, Singapore Management University, Singapore



Greening LLM4SE: Three Strategies



Stop



Simplify



Shrink

Greening LLM4SE

ISSTA 2024

When to Stop? Towards Efficient Code Generation in LLMs with Excess Token Prevention

Lianghong Guo

Sun Yat-sen University
Zhuhai, China
guolh8@mail2.sysu.edu.cn

Wanjun Zhong

Sun Yat-sen University
Guangzhou, China
zhongwj25@mail2.sysu.edu.com

Ruikai Zhang

Huawei Cloud Computing
Technologies Co., Ltd.
Shenzhen, China

Yanlin Wang*

Sun Yat-sen University
Zhuhai, China
wangyulin36@mail.sysu.edu.cn

Hongyu Zhang

Chongqing University
Chongqing, China
hyzhang@cqu.edu.cn

Yuchi Ma

Huawei Cloud Computing
Technologies Co., Ltd.
Shenzhen, China

Ensheng Shi

Xi'an Jiaotong University
Xi'an, China
s1530129650@stu.xjtu.edu.cn

Jiachi Chen

Sun Yat-sen University
Zhuhai, China
chenjch86@mail.sysu.edu.cn

Zibin Zheng

Sun Yat-sen University
Zhuhai, China
zhzibin@mail.sysu.edu.cn



Stop



Improve code generation speed by up to
4.5x by terminating inference early

Greening LLM4SE



Simplify

ISSTA 2024

AI Coders Are Among Us: Rethinking Programming Language Grammar Towards Efficient Code Generation

Zhensu Sun
Singapore Management University
Singapore
zssun@smu.edu.sg

Xiaoning Du*
Monash University
Australia
xiaoning.du@monash.edu

Zhou Yang
Singapore Management University
Singapore
zyang@smu.edu.sg

Li Li
Beihang University
China
lilicoding@ieee.org

David Lo
Singapore Management University
Singapore
davidlo@smu.edu.sg



Won ACM SIGSOFT Distinguished Paper Award

First work to propose a programming language grammar for AI agents

Greening LLM4SE



ESEC/FSE 2023

Shrink

Towards Greener Yet Powerful Code Generation via Quantization: An Empirical Study

Xiaokai Wei

xiaokaiw@amazon.com
AWS AI Labs
USA

Wasi Ahmad

wuahmad@amazon.com
AWS AI Labs
USA

Sujan Kumar Gonugondla

gsujan@amazon.com
AWS AI Labs
USA

Baishakhi Ray

rabaisha@amazon.com
AWS AI Labs
USA

Shiqi Wang

wshiqi@amazon.com
AWS AI Labs
USA

Haifeng Qian

qianhf@amazon.com
AWS AI Labs
USA



Reduce size by 3x and latency by 50% by
quantizing parameter into int8

Related Paper @ ICSME 2025

Is Quantization a Deal-breaker? Empirical Insights from Large Code Models

Saima Afrin

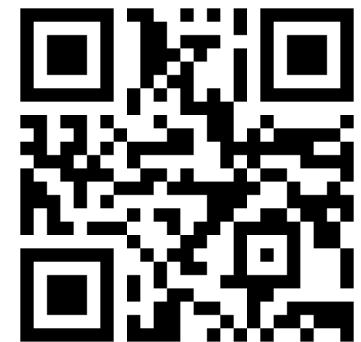
*Department of Computer Science
William & Mary
Williamsburg, VA, USA
safrin@wm.edu*

Bowen Xu

*Department of Computer Science
North Carolina State University
Raleigh, NC, USA
bxu22@ncsu.edu*

Antonio Mastropaoletti

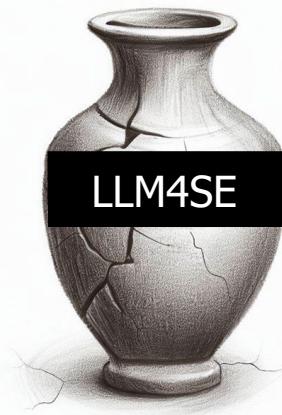
*Department of Computer Science
William & Mary
Williamsburg, VA, USA
amastropaoletti@wm.edu*



Session 2 - Quality Assurance 1
Wed 10 Sep 2025 10:45 - 11:00
Case Room 260-057

Many Open Problems wrt. **Non-Functional** Properties

LLM4SE



***Robustness, Security, Privacy, Explainability, Efficiency, and Usability* of Large Language Models for Code**

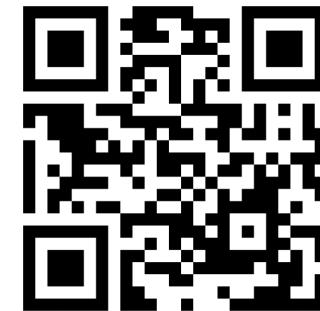
ZHOU YANG, Singapore Management University, Singapore

ZHENSU SUN, Singapore Management University, Singapore

TERRY ZHUO YUE, Singapore Management University, Singapore

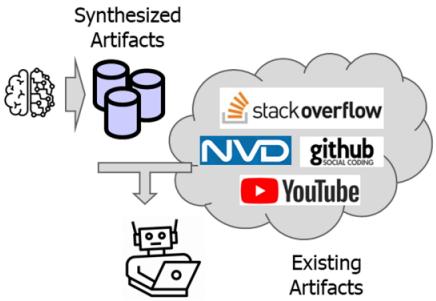
PREMKUMAR DEVANBU, Department of Computer Science, UC Davis, USA

DAVID LO, Singapore Management University, Singapore



Data-Centric LLM4SE: Summary

"How Can We Systematically Engineer Software Artifacts to Build Better LLM4SE Bots?"



Beyond

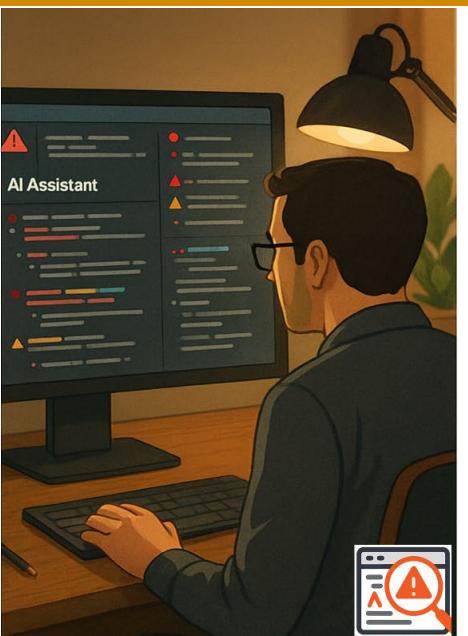


Prompt
Engineering

- We want to **elevate the data**:
 - enriching it, restructuring it, contrasting it, and linking it in ways that *promote deeper reasoning*
- And **utilizing them effectively** to train LLMs to *do better for SE tasks*



Engineering **structured reasoning traces** help LLM reasoning for code generation



Critique

Engineering **contrastive reasoning trace pairs** help LLM reasoning for vulnerability detection

88



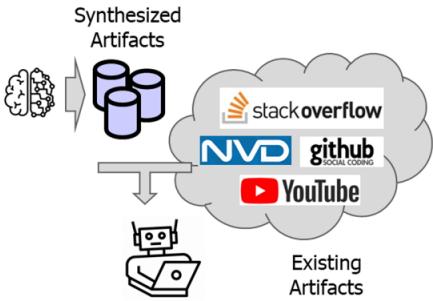
Cure

Engineering **holistic representation from diverse inputs and sources** help LLM reasoning for repair

59

Data-Centric LLM4SE: Summary

"How Can We Systematically Engineer Software Artifacts to Build Better LLM4SE Bots?"



Beyond



Prompt
Engineering



- We w
- er
it,
d
- And i
LLMs

Road Ahead



Multi-Agent LLM4SE

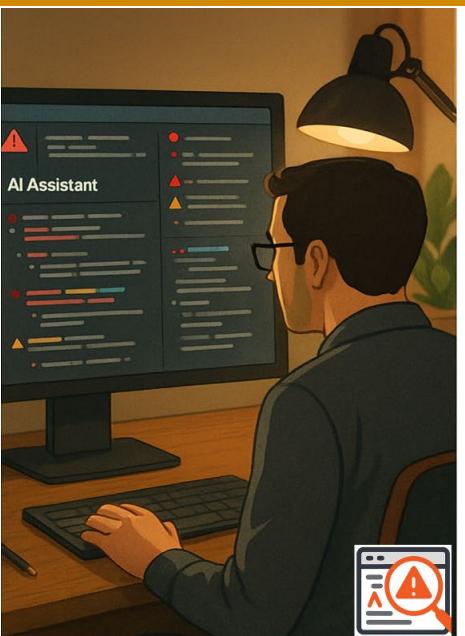
Real-world problems demand "synergistic collaboration" and multi-agent AI transforms isolated AI into coordinated experts.

LLM-Based Multi-Agent Systems for Software Engineering: Literature Review, Vision and the Road Ahead

JUNDA HE, Singapore Management University, Singapore
CHRISTOPH TREUDE, Singapore Management University, Singapore
DAVID LO, Singapore Management University, Singapore



TOSEM 2025



Critique

Engineers
trace pairs help LLM reasoning
for vulnerability detection

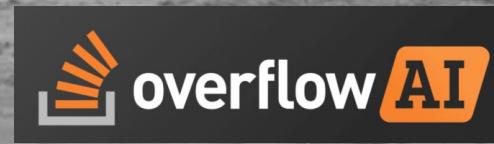


Cure

Engineering **structured reasoning**

races help LLM reasoning
for code generation

Engineering **holistic** representation
from **diverse inputs and sources**
help LLM reasoning for repair



Acknowledgements



ICSME
2025

Auckland, New Zealand



OUNIVERSITY
Chair
Professorship Fund





A background image of a person wearing a futuristic, metallic-looking suit, possibly a spacesuit or a high-tech protective suit. The person is shown from the side, looking down at something in their hands. The background is a light grey with faint, diagonal lines of colorful text or code.

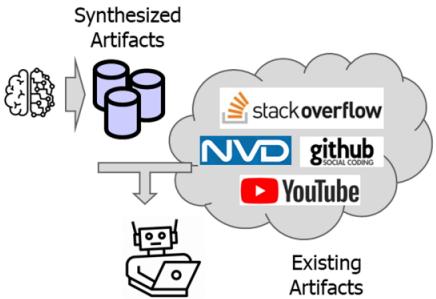
Thank you!

Questions? Comments? Advice?

davidlo@smu.edu.sg

Data-Centric LLM4SE: Summary

"How Can We Systematically Engineer Software Artifacts to Build Better LLM4SE Bots?"



Beyond

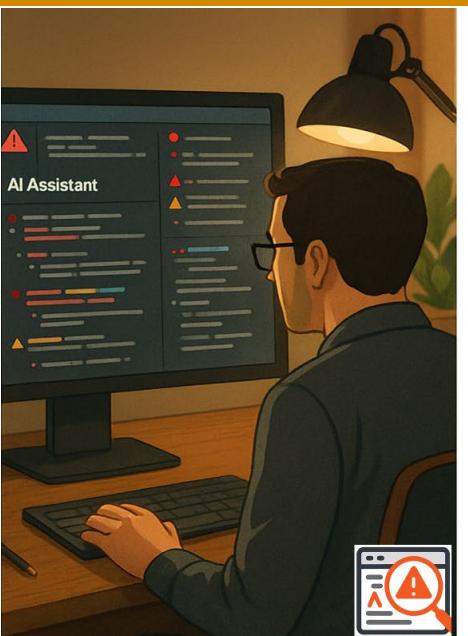


Prompt
Engineering

- We want to **elevate the data**:
 - enriching it, restructuring it, contrasting it, and linking it in ways that *promote deeper reasoning*
- And **utilizing them effectively** to train LLMs to *do better for SE tasks*



Engineering **structured reasoning traces** help LLM reasoning for code generation



Critique

Engineering **contrastive reasoning trace pairs** help LLM reasoning for vulnerability detection

88



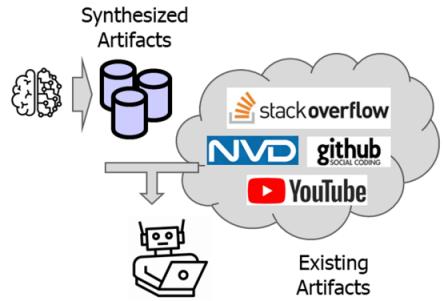
Cure

Engineering **holistic representation from diverse inputs and sources** help LLM reasoning for repair

59

Data-Centric LLM4SE: Summary

"How Can We Systematically Engineer Software Artifacts to Build Better LLM4SE Bots?"



Beyond



Prompt
Engineering



Engineering **structured reasoning**

races help LLM reasoning
for code generation

- We w
- er
it,
d
- And i
LLMs

Road Ahead



Multi-Agent LLM4SE

Real-world problems demand "synergistic collaboration" and multi-agent AI transforms isolated AI into coordinated experts.

LLM-Based Multi-Agent Systems for Software Engineering: Literature Review, Vision and the Road Ahead

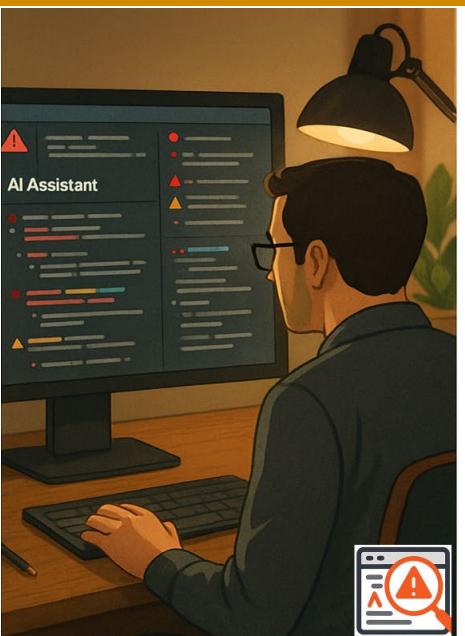
JUNDA HE, Singapore Management University, Singapore
CHRISTOPH TREUDE, Singapore Management University, Singapore
DAVID LO, Singapore Management University, Singapore



TOSEM 2025

Engine

trace pairs help LLM reasoning
for vulnerability detection



Critique



Cure

Engineering **holistic** representation
from **diverse inputs and sources**
help LLM reasoning for repair